Formally Verified Electronic Voting Scheme : A Case Study

Mukesh Tiwari

A thesis submitted for the degree of
Doctor of Philosophy at
The Australian National University

January 2020

Except where otherwise indicated, this thesis is my own original work.

Mukesh Tiwari
14 January 2020

To my grandparents who — despite being poor and uneducated —
understood the value of education.

Acknowledgments

Going through this PhD has truly been a remarkable experience. It has taught me many valuable lessons, but the most important one is to not get intimidated by a hard problem. During this journey, I was fortunate to have guidance and support from many people.

First and foremost, this thesis could not have been possible without the support of my supervisor, Dirk Pattinson. I really admire his abilities and intuition to make sure that I stay clear from many dead ends. Moreover, I thank him for patiently answering my stupid questions happily and guiding me towards the answer by asking the right questions. My only wish to be a researcher like him and incorporate more of his qualities, but I am less optimistic about my chance.

I also want to thank my dog, Turbo, whom I can talk for hours, and he never judged me. He would patiently listen to my proofs and ideas about electronic voting with occasionally barking at me if he was bored of listening. His listening helped me in developing my ideas more better. Turbo truly made my PhD a breeze, and I never felt any pressure of PhD when I was with him. Thank you Turbo, and you are the best dog.

I want to thank Optus, Australian Mobile Service Provider, for unlimited calling hours scheme to India. Because of this scheme, I was able to talk to my family, specifically my mom, everyday for hours. She is from the generation who has seen the mobile and internet revolution in their late age and had hard time in coping up with technological advancement.

Some of the great friends who made this journey possible are:

- Caitlin D'Abrera: She is one the smartest person I know of. She has a super power to break down any problem at finer details to have a better understanding. Moreover, her curiosity to understand the *git* tool forced me to learn it properly, so that I can answer her all the queries. Always be curious Caitlin.

- Milad Ketabi Ghale Ali: Milad was a senior PhD student at logic group

when I started my PhD at logic group, but we have not had any inter-
action until I was in my 8th month of PhD. But once we come to know
each other, it mostly turned out to be discussing Immanuel Kant, Lud-
wig Wittgenstein, David Hume, and Rumi. One of his favourite thing
during any discussion was to not let anyone speak, hide the argument or
not mention it all, and later bring it to end the discussion. In the begin-
ning, it was frustrating, but later we (me and Caitline) learnt to enjoyed
it because it was always a fruitful discussion and chance to know some
good philosophical theories.

- Ali Cheraghian: One of the most funniest person who can lighten up
  any heavy atmosphere with his stupid jokes. One of the most sought
  up skills of Ali is to make everyone excited about disclosing something
  that we did not know, but he would withhold it at the very last moment
  driving everyone crazy.

- Jim De Groot: The funny Dutch who is really very good at making jokes
  in all the situations. I also admire to be a Category theorist like him, but
  I do not see it happening in near future. He is too much into exploring
  nature and because of him, I got a chance to explore the nature around
  Canberra.

- Ian Shillito: The philosopher who did not mind my philosophy jokes.
  Sometimes, I had very profound discussion about the philosophy itself
  with him.

- Saeed Alhamlan: Saeed was a business student when I started my PhD.
  I came to know him through Milad. I travelled to Sydney with Saeed
  in my new year Holidays, and it was a memorable trip. He is great at
  making travelling plans and a excellent travelling companion.

Many thanks to my co-supervisors Rajeev Goré and Michael Norrish.
Their valuable feedback during my PhD monitoring helped me a lot in shap-
ing my presentation skills. Also, I attended programming language reading
group every Wednesday with Michael, and as usual, he was always flawless
in explaining ideas.

I want to thank Thomas Haines, my co-author, for teaching me all the bits
pieces of cryptography, specifically zero-knowledge-proofs. I really enjoyed
working with him and hope to produce many more papers in the upcoming
future.

One of the best experience of this PhD was travelling to Princeton to participate in DeepSpec summer and attending the lectures of Coq stalwarts like Andrew Appel, Benjamin Pierce, and Adam Chlipala. I have learned plenty of tricks about the Coq by watching their lectures on the YouTube. Now that I understand dependent types enough, but it was not always the case. In order to dissect the dependent types, I was watching Benjamin Pierce's Coq lecture on the YouTube. During his lecture he said "proof objects are just data structures", and it was precisely my aha moment in the path of demystifying the dependent types. Unfortunately, there were some sad moments as well. I could not attend Marktoberdorf Summer School because it was next to impossible for me to get the visa of Germany, but I hope things would be more easier in the upcoming future.

I want to thank the Australian National University for providing me resources and scholarship to attend conferences and explore the academic world. In addition, many thanks to Inger Mewburn and her team at ANU for organizing the thesis boot camp which was very helpful in understanding the thesis writing. If, during this thesis, I am not able to convey something in a meaningful way, then the fault is mine because I could not incorporate their teachings properly during the thesis boot camp.

No thesis can be complete in Australia without mentioning the Australian beers and wines. Australia has some of the best beers and wines, and they were very helpful at many occasions during this PhD journey.

Finally, one of the biggest joy of this PhD was to meet my girlfriend, Mina, who came as a visitor to our logic group. Being by your side for the last one and half years has been the best thing that ever happened to me, and I could not imagine life without you. I am also grateful to all the efforts you made to bear my tendencies to constantly talk about set-theory, formal-verification, and every thing except romantic talks. Thank you for being my side all the time Mina.

Abstract

Since the introduction of secret ballot by Victoria, Australia in 1855, paper (ballots) are widely used around the world to record the preferences of eligible voters. Paper ballots provide three important ingredient: correctness, privacy, and verifiability. However, the paper ballot election poses various other challenges, e.g. slow for large democracies like India, error prone for complex voting method like single transferable vote, and poses operational challenges for massive countries like Australia. In order to solve these problems and various others, many countries are adopting electronic voting. However, electronic voting has a whole new set of problems. In most cases, the software programs used to conduct the election has numerous problems, including, but no limited to, counting bugs, ballot identification, etc. Moreover, these software programs are treated as commercial in confidence and are not allowed to be inspected by general member of public. As a consequence, the result produced by these software programs can not be substantiated.

In this thesis, we address the three main concerns posed by electronic voting, i.e. correctness, privacy, and verifiability. We address the correctness concern by using theorem prover to implement the vote counting algorithm, privacy concern by using homomorphic encryption, and verifiability concern by generating a independently checkable scrutiny sheet (certificate). Our work has been carried out in Coq theorem prover.

Contents

List of Figures

List of Tables

Introduction

> The best weapon of a dictatorship is secrecy, but the best weapon of a democracy should be the weapon of openness.

> *Niels Bohr*

1.1   Problem Statement

A democracy can be described as a system where all eligible voters have equal rights to express their opinion(s) on different matters. One of the most important example of expressing the opinion is holding elections to elect the leader of country. During the elections, all eligible voters express their opinion on a paper, also known as ballot, in a manner, depending on the voting method, which reflects their true opinion. For example, if the voting method is *ranked voting (preferential voting)*, then the voters rank the candidates according to their preference, and if the method is *first past the post*, then each voter selects one candidate by marking against the candidate name in the ballot. Later, once the ballot cast finishes, a candidate is elected as a winner from the participating candidates by combing the choices of all the voters. The paper ballot method works great, except it is very time consuming, expensive, error prone, and not very inclusive for disabled voters such as visually impaired. In order to solve the various problems posed by paper ballot, many countries are adopting electronic voting as an alternative. Electronic voting is getting popular in many countries, and the reason for its popularity is cost-effective, faster result, high voter turn out, and accessible for disabled voters.

Undeniably, electronic voting has helped, for example, Australia to ease the logistic challenges of elections because of its massive land size and sparsely population and save millions of dollars. In addition, it has helped India, the second most populous country with 900 million eligible voter, to declare 2019 election with 67 percent voter turn out (roughly 600 million) in 2 days, and Estonia, a labour shortage country, has saved thousands of man hours, 11,000 working days, by using electronic voting [Est].

Despite all these benefits, electronic voting is an arduous effort because a minuscule possibility of going anything wrong in software or hardware could lead to a disaster, possibly inverting the results [Lewis et al.], [Halderman and Teague, 2015], [Aranha et al., 2019], [Feldman et al., 2007]. The nature of (electronic) data and ease of its manipulability/misinterpretation causes electronic voting many problems, which are not present in paper ballot elections, that makes it perfectly susceptible to delivering wrong and unverifiable result [Wolchok et al., 2010]. For instance, if a software program used in electronic voting for reading the ballots has byte order bug, or even if it depends on some other software which has byte order bug (the data is suppose to read from left to right, but software is reading right to left), then the interpretation of ballot would completely be different from what the voter had in mind. More often than not, these software programs are configured incorrectly [Kohno et al., 2004] and run at the top of (untrusted) operating system and hardware. Usually, operating systems have millions of lines code (Linux has 15 millions lines) which exposes a large attack surface and could be exploit, possibly by current government or foreign countries, for illegal gain. The worst, these software and hardware are commercial in confidence and treated as a black-box, and, most often, their source code or design is not open for public scrutiny [Australian Electoral Commission, 2013]. In addition, these software programs take a list of ballots and produce the result without producing any evidence about the correctness of result. As a consequence, from casting the ballot electronically and declaring winner based on cast (electronic) ballot, the whole process lacks basic assumptions of democracy such as transparency, genuineness, and verifiability.

In order to make the electronic voting process genuine and trustworthy, electronic voting research community has recognized some must have properties of electronic voting protocol [Küsters et al., 2011], [Benaloh and Tuinstra, 1994], [Delaune et al., 2010a], [Bernhard et al., 2017]:

- Correctness: The produced results are correct, and convincing to all leaving no ground for suspicion.

- Coercion-resistance: A voter can not cooperate with a coercer to prove anything about her choices.

- Eligibility: Only eligible voter can cast the ballot.

- Privacy: All the votes must be secret, and voter should not be able to convince anyone the value of her vote.

- End-to-end Verifiability: Any independent third party should be able to verify the final outcome of election based on cast ballots. It can be further divided into three sub-categories:

  - Cast-as-intended: Every voter can verify that their ballot was cast as intended.

  - Collected-as-cast: Every voter can verify that their ballot was collected as cast.

  - Tallied-as-cast: Everyone can verify final result on the basis of the collected ballots.

In this thesis, we focus on privacy, correctness, coercion-resistance, and tallied-as-cast, the third part of end-to-end verifiability, property of an election. Furthermore, we assume that the first two properties of end-to-end verifiability, *cast-as-intended* and *collected-as-cast*, hold for an election. *Cast-as-intended* is a verification method that is is used to audit the front end voting software, also known as voting client software, to make sure that it is not modifying the options of voters. In a nutshell, the cast-as-intended is assurance to a voter that front-end software is transparent and her vote is recorded according to her intent. *Cast-as-intended* is itself an active area of research in its own right; however, it is not the focus of this thesis. Similarly, *Collect-as-cast* is a notion related to the voters to make sure that the ballots appearing on the bulletin board are indeed the ballots that cast during the election. A consequence of *collected-as-cast* notion is that any attempt to change or delete the ballots from the bulletin board would be detected. This notion is indeed a crucial one and works as a bridge between cast-as-intended and tallied-as-cast notion. However, it is related to voters' behaviour; hence, the reason for assumption. Moreover, assuming these two notions, cast-as-intended and collected-as-cast, helped us in isolating the irrelevant details and paved a way to focus on more complex problem of counting, i.e. tallied-as-cast.

1.2   Research Motivation and Contribution

Given the potential advantages of electronic voting, we need to address the correctness, privacy, and verifiability concerns for its widespread adoption. This thesis sets out to address these concerns of electronic voting. The questions we asked ourself was:

1. Can we implement a vote counting protocol with the "guaranteed" correctness of the implementation and practical enough to count the real life election involving millions of ballots (Correctness)?

2. Can we produce the result by counting encrypted ballot without revealing its content, and at the same time, assuring everyone that the result produced is only based on "valid" ballots, and "invalid" ones have been discarded (Privacy and Coercion-resistance)?

3. Can we decouple the verifiability from implementation, i.e. generating enough evidence so that any independent auditor can ascertain the outcome of election without trusting the implementation of software used to conduct the election (Verifiability)?

In order to answer these questions, at first we need two things: (i) a voting protocol and (ii) a tool to implement the voting protocol and prove the correctness properties of the implementation. Our choice of voting protocol is the Schulze method [Schulze, 2011] and the tool is Coq [Bertot et al., 2004] theorem prover for implementing and proving the correctness of Schulze method. Even though Schulze method is not used in any democratic election to public office, it has many interesting properties and, at the same time, it is non-trivial. While no preferential voting scheme can guarantee all desirable properties that one would like to impose due to Arrow's theorem [Arrow, 1950a], the Schulze method offers a good compromise, with a number of important properties already established in Schulze's original paper. Amongst the various properties, Schulze method satisfies the *resolvability criterion*, i.e. elects a single winner under the assumption that number of voters are much larger than number of candidates (and in case of tie, when there are more than one winner, a random vote can be selected to declare winner. However, our formalization has not taken the randomness into account, so it can produce more than one winner). Coq is a theorem prover (proof assistant) based on *Calculus of Inductive Construction* [Coquand and Huet, 1988] [Coquand and Paulin, 1988]. The *Calculus of Inductive Construction* allows the "proof" terms and "computation" terms to live in the same universe (level), which leads to

a highly expressive type system. Moreover, during the proof development, it provides a step by step feedback to the user and possibility to automate the proofs by writing the custom tactics using the Ltac [Delahaye, 2000]. In addition, Coq proofs can be extracted into the Haskell, OCaml, and Scheme.

Now that we have a voting protocol (Schulze method) and a tool (Coq theorem prover), we demonstrate that it is possible to achieve correctness, privacy, coercion-resistance, and (tallied-as-cast) verifiability in electronic voting. We achieve the following:

- *Correctness* by formally specifying the Schulze method and prove its correctness properties inside the Coq theorem prover. Coq has a well developed extraction facility that we use to extract proofs into OCaml programs, and using these extracted OCaml programs, we have counted the ballots from election to produce the result.

- *Privacy and Coercion-resistance* by encryption. We use homomorphic encryption to compute the finally tally without decrypting any individual ballot.

- *Verifiability* by tabulating the relevant data of election (we call it scrutiny-sheet/certificate). Achieving verifiability in a plain-text ballot counting is fairly straight forward. To achieve verifiability in encrypted ballot counting, we augment the scrutiny sheet with zero-knowledge-proof for the each claim we make during the counting which can later be checked by any auditor.

Formally Verified Checker: Third party independent election audit based on scrutiny sheet data is a crucial step towards establishing the trust in the system. However, auditing scrutiny sheet of a election involving encrypted ballots is not straight forward in comparison to a election with plain-text ballots. In general, auditing scrutiny sheet of a election involving plain-text ballot simply requires the knowledge of basic arithmetic, e.g. addition, subtraction and multiplication, and virtually anyone can audit the election based on the data produced in the scrutiny sheet by using a calculator or by writing a simple program in her preferred language. However, encrypted ballot election scrutiny sheet involves various cryptographic concepts (homomorphic encryption, zero knowledge proof, commitment scheme, etc.) which are accessible to a very few voters, mainly cryptographers, so auditing it requires deep understanding of cryptographic principals. To ease this situation, we have developed a formally verified certificate checker as a proof of concept

for automating the auditing an election conducted on encrypted ballots. Having said that, our certificate generated by encrypted ballots is very complex, and formalizing all the cryptographic primitives involved would be fairly time consuming, so we have developed a proof of concept formally verified certificate checker for International Association of Cryptographic Research 2018 election scrutiny sheet, relatively simple than ours.

Properties of Schulze Method: We have proved couple of the properties, Condercet winner, and Reversal symmetry, of the Schulze method inside Coq theorem prover (ongoing work).

1.3 Cryptographic Blackbox

Since the beginning of this project, our primary goal was to achieve privacy (using encryption) and verifiability (using zero knowledge proof) in electronic voting using the cryptographic primitives (but not the verification of primitives itself). To achieve this goal, we have taken the axiomatic approach and assumed the existence of cryptographic primitives inside Coq. Moreover, we assume the axioms about their correctness behaviour, e.g. decryption is left inverse of encryption. These primitives, in general, provide functionality of generating a random permutation, encrypting a plain-text, decrypting a cipher-text, producing commitment of a value, constructing a zero-knowledge-proof, and verifying a zero-knowledge-proof. Later, in extracted OCaml code from Coq code, these functions are instantiated with Unicrypt [Locher and Haenni, 2014] function. [1]

1.4 Publication

The chapters, or some part of it, of this thesis are based on the following papers:

1. Pattinson, D. and Tiwari, M., 2017. Schulze Voting as Evidence carry-

---

[1]Formalizing the whole cryptographic stack used in our project would be very time consuming (probably a PhD itself), but it would be worth trying. Although, we have formalized the (El-Gamal) encryption, and decryption inside Coq, but we still are very far from achieving the goal of fully verified cryptographic stack. We leave the formalisation of cryptographic primitives for future work (work in progress).

ing computation. In Proc. ITP 2017, vol. 10499 of Lecture Notes in Computer Science, 410–426. Springer.

2. Lyria Bennett Moses, Rajeev Goré, Ron Levy, Dirk Pattinson, Mukesh Tiwari. No More Excuses: Automated Synthesis of Practical and Verifiable Vote-Counting Programs for Complex Voting Schemes. E-VOTE-ID 2017: 66-83

3. Milad K. Ghale, Rajeev Goré, Dirk Pattinson, Mukesh Tiwari. Modular Formalisation and Verification of STV Algorithms. E-Vote-ID 2018: 51-66

4. Thomas Haines, Dirk Pattinson, and Mukesh Tiwari. 2019. Verifiable Homomorphic Tallying for the Schulze Vote Counting Scheme. 11th International Conference Verified Software: Theories, Tools, and Experiments. VSTTE 2019 (to appear)

5. Thomas Haines, Rajeev Goré, and Mukesh Tiwari. 2019. Verified Verifiers for Verifying Elections. In Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security (CCS '19).

Part of chapter 2 is based on *No More Excuses: Automated Synthesis of Practical and Verifiable Vote-Counting Programs for Complex Voting Schemes*, chapter 4 is based on *Schulze Voting as Evidence Carrying Computation*, chapter 5 is based on *Verifiable Homomorphic Tallying for the Schulze Vote Counting Scheme*, and part of chapter 6 is based on *Verified Verifiers for Verifying Elections*.

1.5   Related Work

There is extensive work that addresses the different issues related of electronic voting protocols in symbolic model, but there are very few, to the best of my knowledge, that have used theorem provers to implement the voting protocol (counting algorithm) and verify its correctness properties. [Kremer and Ryan, 2005], and [Delaune et al., 2010b] have used pi-calculus, a formalism for describing concurrent systems, to model and analyse various properties, such as fairness, eligibility, vote-privacy, receipt-freeness and coercion-resistant, of the protocol FOO developed by [Fujioka et al., 1993]. [Backes et al., 2008] presented a general technique to model remote electronic voting protocol and automatically verifying its security properties using pi-calculus. [Cortier and Smyth, 2011] have used pi-calculus to analyse the ballot secrecy of [Helios,

2016]. [Cortier and Wiedling, 2012] have used pi-calculus to ascertain the properties of Norwegian electronic voting protocol. [Bruni et al., 2017] have used Tamarin to prove receipt-freeness and vote-privacy of Selene voting protocol [Ryan et al., 2016]. Most of these work differs from ours in the sense that their primary focus is verification of security protocol in Dolev-Yao model or complexity-theoretic model, whereas our work is more focused on verified implementation and verifiability aspect of vote counting.

The closest to our work is [Cochran and Kiniry, 2010] [DeYoung and Schürmann, 2012], [Pattinson and Schürmann, 2015], [Pattinson and Verity, 2016], [Verity and Pattinson, 2017], and [Ghale et al., 2017]. [Cochran and Kiniry, 2010] have used Business Object Notation (BON) and Java Modelling Language (JML) to formally specify their Java implementation of Irish Proportional Representation by Single Transferable Vote (PR-STV) method. They relied on Extended Static Checking to validate the correctness of their implementation. Upon further investigation [Cochran and Kiniry, 2013], they improved it by writing formal specification of candidate, ballot, and ballot box datatypes using the Alloy model checker [Jackson, 2002]. However, they themselves pointed out that:

> Note that this automated consistency checking is not the same as providing a full interactive proof of a soundness theorem in a higher-order logical framework. Such formalization is an interesting and useful exercise, but we did not do it for this case study. Instead, checking the dozens of theorem stipulated in law text is more akin to the kind of validation that we are advocating in this work. It gives us high confidence, but not a proof, that the mechanical formalization is sound and complete.

[DeYoung and Schürmann, 2012] used linear logic [Girard, 1987] to model the different entity in electronic voting as a resource. The use of linear logic makes it very natural to capture the different entities in electronic voting, depending on their usage, by means of modality e.g. a voter can cast only one vote, but she might need to show her photo id to multiple times at counting booth. [Pattinson and Schürmann, 2015] treated the process of vote counting from the perspective of mathematical proof. They used (mathematical) proof theory to model the counting. [Ghale et al., 2017] have formalized single transferable vote in Coq and extracted Haskell code from the formalization. The extracted Haskell code produces the result and a certificate for a given set of input ballots. This certificate can be used by any third party to verify or audit the outcome of election result. However, none of these work considers

privacy and coercion resistance as a key issue in electronic voting, and their method simply works for plaintext ballots which are susceptible to italian attack [Otten, 2003] [Benaloh et al., 2009].

1.6   Outline of the Chapters

- Chapter 2 provides an overview of electronic voting around the world, problems in general, and rationale for formal verification of election voting software.

- Chapter 3 provides the overview of concept of Coq theorem prover and cryptographic primitives.

- Chapter 4 describes Schulze method, its formal specification, proof of correctness, experimental result, and scrutiny sheet.

- Chapter 5 describes verifiable homomorphic tally for the Schulze method, its realization in theorem prover, experimental result, instructions to audit the scrutiny sheet.

- Chapter 6 focuses on the notion of software independence, and sketches the details for the formalization of cryptographic concepts involved in the certificate generated by encrypted ballot.

- Chapter 7 puts forward the idea of machine checked properties of electronic voting scheme and describes couple of the properties, condercet winner and reversal symmetry, of Schulze method.

- Chapter 8 concludes the thesis, and some possible direction of future work.

1.7   Trivia

Before 1856, Victoria and NSW held their elections to elect its democratic representative in pub where it was legal for candidates to offer beer to voters to influence their decision!

A NORTH BOURKE CANDIDATE ADDRESSING HIS CONSTITUENCY ON A "NATIONAL" QUESTION.

Figure 1.1: Election held in 1855 in Victoria, Australia was conducted in pub!

Background

> People shouldn't be afraid of
> their government. Governments
> should be afraid of their people.
>
> *Alan Moore, V for Vendetta*

Counting ballots by hand is tedious, error prone, slow, and costly process. For example, the Senate election conducted in Western Australia in September 2013 was declared void by the high court because of the loss of 1370 votes. It was re-conducted in April 2014 at the cost of 20 Million AUD with additional delay in results [Aus]. Before introduction to electronic voting machines in India, it used to take months to declare the result. As a consequence, many countries are now adopting electronic voting to alleviate the problems introduced by hand counting. The world can be divided into five broad categories according to the usage of electronic voting [Evo] 2.1: i) No electronic voting (Grey Area), ii) Discussion and/or voting technology pilots (Yellow Area), iii) Discussion and concrete plans for Internet voting (Orange Area), iv) Ballot scanners, Electronic Voting Machines, and Internet Voting (Green and Dark Green), v) Withdrawn voting technology because of public concern (Red Area)

Chapter Outline: In the section 2.1, we discuss the major concern in electronic voting, software/hardware bug used in the election process, by high lighting the state of electronic voting in Australia, Germany, India, and Netherlands. In the following section 2.2, we give two anecdotes that how formal verification helped in achieving the correctness and eliminating bug in *CompCert*, and *Athelon* with the emphasis that we should formally verify the software programs used in electronic voting to alleviate the concerns of

Figure 2.1: World map of Electronic Voting

software/hardware bug. Section 2.3 emphasizes that formal verification is not enough to establish the trust and puts forward the concept of scrutiny sheet (2.3.1), which can be used independently to attest the result of election, to achieve verifiability. Finally, we summarize in the section 2.4 by emphasizing that formal verification and verifiability both are needed to ensure the trust in electronic voting.

2.1   Electronic Voting

Electronic voting is projected as a step towards the future with many benefits, such as increased voter turnout, faster result, accessible to everyone including challenged voters, and reduced carbon footprint (for each national election, India saves about 10,000 tonnes of the ballot paper by using electronic voting

machines). There is no doubt that electronic voting has many advantages over paper ballot, but it is certainly not flawless. Electronic voting makes the process faster, but it has its own layer of added complexities which creates trust issues amongst voters. For this reason, some countries who were the early adopters were also the early abandoner, e.g. Germany, and The Netherlands (countries in red color in the world map 2.1).

Germany: In 2005 German election, two voters filed a case in the German Constitutional Court (Bundesverfassungsgericht) because their appeal to scrutinize the election was not heeded by the Committee. They argued that using electronic voting machines to conduct the election was unconstitutional. Furthermore, they added that these machines could be hacked, hence results of the 2005 election could not be trusted on the grounds of public examinability of elections according to German Constitution (Basic Law for the Federal Republic of Germany) [Ger]. The Court noted that, under the constitution, elections are required to be public in nature [Ger]:

> The principle of the public nature of elections requires that all essential steps in the elections are subject to public examinability unless other constitutional interests justify an exception. Particular significance attaches here to the monitoring of the election act and to the ascertainment of the election result.

In its verdict, the court did not rule out or prevent the usage of electronic voting machines, but suggested to make the process more transparent and trustworthy [Ger]:

> The legislature is not prevented from using electronic voting machines in the elections if the constitutionally required possibility of a reliable correctness check is ensured. In particular, voting machines are conceivable in which the votes are recorded elsewhere in addition to electronic storage. This is for instance possible with electronic voting machines which print out a visible paper report of the vote cast for the respective voter, in addition to electronic recording of the vote, which can be checked prior to the final ballot and is then collected to facilitate subsequent checking. Monitoring that is independent of the electronic vote record also remains possible when systems are deployed in which the voter marks a voting slip and the election decision is recorded simultaneously, or subsequently by electronic means in order to evaluate these by electronic means at the end of the election day.

The Netherlands: The Netherlands was among a few countries who adopted electronic voting in the early nineties (1990), but it did not go very well in the long run and was abolished in 2008 [Jacobs and Pieters, 2009]. The reason for abolishing the electronic voting was that the voting machines used in elections were susceptible to many attacks, and the results of elections conducted using these machines were not publicly verifiable. Besides, a Dutch public foundation, Wij vertrouwen stemcomputers niet (We do not trust voting computers), demonstrated that e-voting machines used in election leaks enough information to guess the choice of a voter from distance of 20 to 30 meters [Net].

Germany and The Netherlands are some of the rare cases where electronic voting was withdrawn because it was not able to replicate the same trust environment as created by paper ballot systems whereas Australia, and India continued with electronic voting despite having the concerns expressed by researchers about the security of system.

India: India, one of the largest democracies in world, uses electronic voting machines (also known as EVMs) for national and state level elections despite the fact that many political parties have raised security concern against it. Moreover, it has already been shown in [Wolchok et al., 2010] that it is possible to manipulate the election results. In their attack, they replaced the parts of electronic voting machine with malicious look alike components. These components were capable of receiving instruction over wireless communication. As a result, any malicious attacker can control these components from nearby vicinity by sending instructions over wireless channel by using a simple hand-held device and can manipulate the results in their favour [1]. India is mainly criticised for keeping the design of electronic voting machines a closely guard secretes (security by obscurity). However, it is not impossible to get access of these machines as shown by [Wolchok et al., 2010]. The worst part, the design of these machines were never audited by any independent third party.

Australia: In March, 2015 state election of New South Wales, Australia, a Internet voting system, iVote, was used and 280,000 votes were cast through it. NSW Election commissioner claimed that it was:

> It's fully encrypted and safeguarded, it can't be tampered with, and for the first time people can actually after they've voted go into the system and check to see how they voted just to make sure

---

[1]https://indiaevm.org/

everything was as they intended [NSW].

The voting on iVote opened on Monday March 16 and continued until March 28. On 22 March, two security researchers, Vanessa Teague and J. Alex Halderman, announced that iVote has critical security bug, and they demonstrated that it was good enough to steal any ballot. From their paper [Halderman and Teague, 2015]:

> While the election was going on, we performed an independent, uninvited security analysis of public portions of the iVote system. We discovered critical security flaws that would allow a network-based attacker to perform downgrade-to-export attacks, defeat TLS, and inject malicious code into browsers during voting. We showed that an attacker could exploit these flaws to violate ballot privacy and steal votes. We also identified several methods by which an attacker could defeat the verification mechanisms built into the iVote design.

Basically, New South Wales ran a online election for 6 days on buggy software which was susceptible to many attacks with a possible outcome of tampered ballot without anyone noticing it.

We do not have to think very hard to figure out the reasons for these debacles. There are various factors for these debacles, but one of the most common denominator among all these debacles, which contributed significantly, is the software/hardware used in the election process were buggy. But, this begs the question: why various governments were (Germany, Netherlands)/are (Australia, India) using buggy software in the first place for conducting elections? In general, no entity related to government or electoral commission develops the software programs for electronic voting, and predominantly it is outsourced to the companies having experience in electronic voting software development. Most of these companies do not follow the rigours methodology leading to buggy software programs production. Moreover, these software programs are closely guarded secrets and their source code in not open for public scrutiny because of commercial interests of companies [Australian Electoral Commission, 2013] involved in the process. Overall, the whole process lacks transparency, which violates the fundamental principals of democracy, i.e. openness.

The process of turning a vague idea into a concrete software, also known as software development process, involves requirement gathering, software

design, implementation, testing, and maintenance. During this whole process of software development, software testing is the only mechanism for quality assurance. However, testing is not enough for instilling the confidence in software that it is bug free as stated by Edsger W. Dijkstra [Dijkstra, 1972]:

> Program testing can be used to show the presence of bugs, but never to show their absence!

Additionally, most of these companies have poor testing practices which exacerbate the software quality problem.

In the next section, given the mission critical importance of electronic voting software we will discuss that software testing in not sufficient for these software, and we should prove the correctness of these software by using formal verification techniques [Beckert et al., 2014]. Furthermore, we will argue that having a formal verification software development methodology would alleviate the bug problem with few case studies as a supporting evidence. The success of these case studies should be a good motivation for us to adopt formal method for electronic voting software development.

2.2 Correctness: Formal Method Approach

Formal verification has been successfully applied in many areas ranging from verified C compilers CompCert [Leroy, 2006], verified ML compiler CakeML [Kumar et al., 2014], Vellvm: Verifying the LLVM [Zhao and Zdancewic, 2012], verified cryptography Fiat-crypto [Erbsen et al., 2019], verified operating system CertiKOS [Gu et al., 2011] and SeL4 [Klein et al., 2009], verified theorem prover Milwa [Myreen and Davis, 2014], verified crash resistant file system FSCQ [Chen et al., 2015], verified distributed system Verdi [Wilcox et al., 2015], mechanisation of Four Color Theorem [Gonthier, 2008], Fundamental Theorem of Algebra [Geuvers et al., 2002], and Kepler Conjecture [Hales et al.]. One thing I would like to emphasize is that none of these are toy project, and it has taken years to develop and verify them. Also, some of these products are used commercially e.g CompCert is used by AIRBUS, and MTU [2] and Fiat-crypto is used in Google's BoringSSL library for elliptic-curve arithmetic [3]. Based on the cost and efforts of these project, the very

---

[2]https://www.absint.com/compcert/
[3]https://deepspec.org/entry/Project/Cryptography

basic question to ponder about using formal method to develop software: does it achieve the goal, produces bug free software?

We would give two anecdote to answer this question. One of the most basic way to break the software is generating random tests and throwing it to the software under consideration [Miller et al., 1990]. [Yang et al., 2011] developed random C program generator and used these programs to test various compilers. In three years of its usage, they have found 325 unknown bugs in various compiler including GCC [4] and LLVM [5]; however, they could not find any bug in verified component of CompCert. In their own words [Yang et al., 2011]:

> The striking thing about our CompCert results is that the middle-end bugs we found in all other compilers are absent. As of early 2011, the under-development version of CompCert is the only compiler we have tested for which Csmith cannot find wrong-code errors. This is not for lack of trying: we have devoted about six CPU-years to the task. The apparent unbreakability of CompCert supports a strong argument that developing compiler optimizations within a proof framework, where safety checks are explicit and machine-checked, has tangible benefits for compiler users.

Formal verification is not only helpful in proving the correctness, but sometimes, it helps in uncovering the bugs in design of software. ACL2, a Lisp based theorem prover, helped AMD to uncover a floating point bug in Athelon processor which has survived 80 million floating point test cases! In the paper, Milestones from the Pure Lisp theorem prover to ACL2 [Moore, 2019], Moore, one of the developer of ACL2, writes:

> When AMD developed their translator from their register-transfer language (in which designs are expressed) to ACL2 functions they ran 80 million floating point test cases through the ACL2 model of Athlon's FMUL and their own RTL simulator. However, the subsequent proof attempt exposed bugs not covered by the test suite. These bugs were fixed before the Athlon was fabricated.

There are numerous instances where formal verification was very useful, and it caught the lurking bugs in design in early stage which could never have

---

[4]https://embed.cs.utah.edu/csmith/gcc-bugs.html
[5]https://embed.cs.utah.edu/csmith/llvm-bugs.html

been found by testing. For electronic voting software used in democratic election, where we can not afford to lose a single ballot or miscalculation or any undefined behaviour, should be developed using formal method techniques. In order to ascertain that the formal verification of voting software has been carried out diligently, one therefore needs to

1. read, understand and validate the formal specification: is it error free, and does it indeed reflect the intended functionality?

2. scrutinize the formal correctness proof: has the verification been carried out with due diligence, is the proof complete or does it rely on other assumptions?

The above mentioned requirements can be met by publishing or open sourcing both the specification and the correctness proof so that the specification can be analysed, and the proof can be replayed (inside the tool used for verification) by any independent third party. Both need a considerable amount of expertise, but it can be carried out by (ideally more than one group of) domain experts.

2.3   Verifiability: Trust in Electronic Voting

Formal verification is useful in producing the bug free code, but we solely can not establish the trust in the system based on argument of formal verification. The reason is that how would a voter:

• ascertain that it was indeed the verified program that was executed in order to obtain the claimed results?

• ensure that the computing equipment on which the (verified) program is executed has not been tampered with or is otherwise compromised?

Formal verification is certainly a necessary thing in electronic voting because it helps in achieving the correctness of software, but it is not sufficient because it has no mechanism to provide verifiability. Combing both *verification* of the computer software that counts votes, and *verifiability* of individual counts are critical for building trust in an election process. These two facts, *verification* and *verifiability*, can also be viewed as two sides of a coin from the perspective

of two major stake holders of a democracy: i) electoral commission/government, and, ii) voters/participants. Using formally verified software to count the ballots would increase the confidence of electoral commission that they has announced and published the correct result, and the published result always verifies would increase the confidence of voter in the system.

Given the mission-critical importance of correctness of vote-counting, both for the legal integrity of the process and for building public trust, it is crucial to replace the currently used black-box software for vote-counting with a counterpart that is both verified and produces evidence which can later be used to certify the outcome of election [Bernhard et al., 2017] [Rivest, 2008].

2.3.1 Scrutiny Sheet

A scrutiny sheet is the tabulation of relevant data to verify the result of election. The idea of requiring that computations provide not only results, but also a witness attesting to the correctness of the computation is not new and has been put forward in [Sullivan and Masson, 1990] [McConnell et al., 2011] [Arkoudas and Rinard, 2005], and, in the context of electronic voting by [Schürmann, 2009] [Pattinson and Schürmann, 2015]. In general, the idea of computation is that a computable function f, it takes a input x and produces output y (figure 2.2); however, in case of certified computation, the computable function f on the given input x, not only produces the output y, but it also produces a witness w for the fact that $f(x) = y$ (figure 2.3).

Figure 2.2: Function f computing y on input x

Below is a certificate which has been generated by a program which computes the greatest common divisor of two numbers. The program has produced the certificate on the concrete input 34 and 21:

```
gcd 34 21
```

Figure 2.3: Function f computing y and producing witness w on input x

```
- - - - - - - - - - - - -
gcd 21 13
- - - - - - - - - - - -
        .
- - - - - - - - - - - -
gcd 1 0
- - - - - - - - - - - -
     1
```

In order to verify the correctness of the computation of greatest common divisor, we need to make sure that one of the rules of Euclidean algorithm, given below, is applicable at each line of the certificate:

1. Rule-zero: $\forall x$, gcd x 0 = x

2. Rule-inductive: $\forall x\ y$, gcd $x\ y$ = gcd y (mod $x\ y$)

The same certificate (augmented with rules and variables instantiated) from certificate-checker perspective.

```
gcd 34 21 = gcd 21 13 Rule-inductive: x := 34, y := 21, mod 34 21 = 13
------------------------------------------------------------------
gcd 21 13  = gcd 13 8 Rule-inductive: x := 21, y := 13, mod 21 13 = 8
------------------------------------------------------------------
     .
------------------------------------------------------------------
gcd 1 0 = 1   Rule-zero: x := 0
```

Unfortunately, the example we gave, greatest common divisor, is very simple, and it is not very helpful example to put the usefulness of certificate

in perspective. However, this approach, generating a certificate to certify the computation later, is very useful in context of complicated and unverified programs. One such example is algorithmic library LEDA [Mehlhorn and Näher, 1995] ("Library of Efficient Data types and Algorithms") written in C++. Checkers are integral part of LEDA which can later be invoked by user to certify that the result produced by unverified code is correct. Initially, the checkers came with library were unverified, and Kurt Mehlhorn defended this decision by admitting that [Alkassar et al., 2014]:

> Checkers are simple programs with little algorithmic complexity. Hence, one may assume that their implementations are correct.

Later, the checkers [Alkassar et al., 2014] were verified by using VCC [Cohen et al., 2009], and Isabelle/Hol [Nipkow et al., 2002b]. One advantages of this approach is that it is easier to formally verify the checker than the algorithm itself, because checkers are very simple in nature, and this approach scales very well.

One may ask the question that can we follow the same approach for vote counting, i.e. unverified counting code, and verified checker? The answer is: it depends. If the verified checker validates the result, i.e. it returns true on the certificate generated by unverified vote counting program then everything is fine from every stakeholders' perspective. However, what about the situation when the verified checker returns false, i.e. invalidates the result. This kind of situation must be dealt carefully by electoral commission, and commission should inspect everything carefully including the vote counting software and various other thing involved in the process. To eliminate this kind of problematic situation, we propose: i) a formally verified vote counting software which produces the result with evidence (certificate), and ii) formally verified certificate checker. The advantage of this approach is that the result produced is always correct (modulo specification), a confidence building measure for electoral commission. Furthermore, the verified checker would always return true on the evidence (certificate) produced by verified vote counting program, confidence booster for voters into the deployed system.

2.4   Summary

We reiterate that to maximise trust, reliability and auditability of electronic vote counting, we need both approaches, verification and verifiability, to be combined. To ensure verifiability, we advocate that vote-counting programs do not only compute a final result, but additionally produce an independently verifiable certificate that attests to the correctness of the computation, together with a formal verification that valid certificates indeed imply the correct determination of winners. Given a certificate-producing vote-counting program, external parties or stakeholders can then satisfy themselves to the correctness of the count by checking the certificate.

In the next chapter, I will briefly discussion about Coq theorem prover and basic cryptographic primitives which would enable us later in counting encrypted ballots (without revealing any information about voter's choice).

Theorem Prover and Cryptography

> All our knowledge begins with
> the senses, proceeds then to the
> understanding, and ends with
> reason. There is nothing higher
> than reason.
>
> *Immanuel Kant*

A proof assistant or theorem prover is a computer program which assists users in development of mathematical proofs. Basically, the idea of developing mathematical proofs using computer goes back to Automath (automating mathematics) [de Bruijn, 1983] and LCF [Milner, 1972]. The Automath project (1967 until the early 80's) was initiative of De Bruijn, and the aim of the project was to develop a language for expressing mathematical theories which can be verified by aid of computer. Moreover, the Automath was first practical project to exploit the Curry-Howard isomorphism (proofs-as-programs and formulas-as-types). DeBruijn was likely unaware of this correspondence, and he almost re-invented it. The Automath project can be seen as the precursor of proof assistants NuPrl [Constable et al., 1986] and Coq [Bertot et al., 2004]. Some other notable proof assistants are Nqthm/ACl2 [Kaufmann and Strother Moore, 1996], PVS [Owre et al., 1992], HOL (a family of tools derived from LCF theorem prover) [Slind and Norrish, 2008] [Harrison, 1996] [Nipkow et al., 2002a], Agda [Norell, 2009], and Lean [de Moura et al., 2015].

Chapter overview: This chapter is overview of Coq theorem prover and Cryptographic primitives. In the section 3.1, we will give a brief overview of theoretical foundation, calculus of construction and calculus of inductive construction, of Coq. In the section 3.1.2, we will discuss about `Type` and `Prop` which is very crucial from program extraction point of view. Our goal during

this course of formalization was not only proving the correctness of Schulze method, but extracting a executable OCaml/Haskell code to count ballots. In the section 3.1.3, we will focus on dependent types and how it leads to correct by construction paradigm by designing a type safe printf function. Section 3.1.4 focuses on Coq specification language `Gallina` with a example showing that why writing proofs using `Gallina` is difficult and cumbersome, and how it can be eased by using tactics. Finally, in the section 3.1.5, we will take philosophical route to justify that why should we trust in Coq proofs even though they do not appear anywhere near to a proof written by a human. In section 3.2, we give some historical context and modern day usage of Cryptography. In the following section 3.2.1, we describe the Group which is the underlying algebraic structure for Diffie-Hellman construction (3.2.2). In the next two sections, we describe the ElGamal encryption (3.2.3) and Homomorphic Encryption (3.2.4). In addition, we show the both homomorphic property, multiplicative and additive, of ElGamal encryption. We explain the concept of Zero-Knowledge-Proof and Zero-Knowledge-Proof of knowledge in section (3.2.5). In the next two sections, we discuss Sigma protocol (3.2.6), an efficient way to achieve zero knowledge proof, and Commitment schemes (3.2.7), a cryptographic protocol to force two mutually distrusting parties to behave honestly with the explanation of Pedersen commitment scheme based on discrete logarithm. Finally, we give a brief summary pointing to the resources for theorem proving and cryptography.

3.1   Coq: Interactive Proof Assistant

Coq is an interactive proof assistant (theorem prover) based on theory of Calculus of Inductive Construction [Paulin-Mohring, 1993] which itself is an augmentation of Calculus of Construction [Coquand and Huet, 1988]. The underlying theory of CoC and CIC is (typed) lambda calculus, so before we describe the CoC and CIC syntax and its typing judgement, we will take a brief detour to explain different variants of lambda calculus starting from untyped lambda calculus and moving up in the ladder by adding various abstractions. Later, we will show that these all variant, including CoC, can be abstracted into one framework by using pure type system [Berardi, 1988] [Barendregt, 1992]. In addition, pure type system can be extended with three rules, inductive data type, pattern matching, and recursion to accommodate Calculus of Inductive Construction.

Lambda calculus was invented by Alonzo Church in the 1930s, and his

motive was to use lambda calculus as a foundation for formal mathematics, specifically the notation of computable function by means of an algorithm. It is a simplest programming language having just three constructs, i.e. variable, application, and abstraction, and the abstract syntax tree of lambda calculus is:

```
T = V (* Variables *)
 | λ V. T (* Abstraction *)
 | T T (* Application *)
```

Using these three rules, we can construct the lambda terms corresponding to various mathematical notions. For example, we can represent *True* as $\lambda x.\lambda y.x$, *False* as $\lambda x.\lambda y.y$, *Zero* as $\lambda f.\lambda x.x$, *One* as $\lambda f.\lambda x.fx$, etc. However, there is nothing which is stopping us to construct lambda terms which has no meaning, e.g. applying a variable x with itself learning to a lambda term xx. To avoid these kind of terms, we extend the lambda calculus with another abstraction called *types*. Moreover, we add *typing judgement* (rule) which dictates which terms is well-typed and which one is not. This new lambda calculus augmented with *types* is known as *Simple Typed Lambda Calculus*, represented as $\lambda^{\rightarrow}$. The abstract syntax tree for *Simple Typed Lambda Calculus* is:

```
𝒯 = 𝒱 (* Type Variable *)
 | 𝒯 → 𝒯 (* Arrow Type *)
```

```
T = V (* Variables *)
 | λ V : 𝒯. T (* Abstraction *)
 | T T (* Application *)
```

The typing judgement is a relation between *types* and *terms* in some abstract typing context Γ. The Γ is a set/list of typing assumption of the form $x : A$, meaning x is of type A. Moreover, $\Gamma \vdash x : A$ means that the term x has type A in the context Γ. The typing judgement of *Simple Typed Lambda Calculus* has three rules, *Var, Abstraction*, and *Application*, to ensure that the terms are well-typed:

- Var:
$$\frac{x : A \ \in \ \Gamma}{\Gamma \ \vdash \ x : A}$$

- Abstraction:

$$\frac{\Gamma, x : A \ \vdash \ e : B}{\Gamma \ \vdash \ (\lambda\,x : A.\,e\,) : A \to B}$$

- Application:

$$\frac{\Gamma \ \vdash \ f : A \to B \quad \Gamma \ \vdash \ x : A}{\Gamma \ \vdash \ f\,x : B}$$

Now these three typing judgement rules rule out the term *xx* because it is not well-typed term, and this can be inferred from *Application* rule. The *Application* rule states that for *f x* to be a well typed term of some type *B*, the *f* has to have a type $A \to B$ for some type *A* and *x* has to have the type *A*. Following the *Application rule*, for *xx* to be well typed, the *x* has to have a arrow type $A \to B$ and type *A* in some typing context Γ. However, it is not possible that $A = A \to B$ leading to rejection of term *xx*.

Simple typed lambda calculus is great for many practical purposes except it is verbose. Consider a function which takes a input and simply returns it, also known as identity function. Assuming that we two base type, *nat* for the type of natural numbers and *bool* for type of boolean values, as member of type variable set $\mathcal{V}$. We can represent a identity function on boolean value as $\lambda x : bool.x$, and on natural number as $\lambda x : nat, x$. In general, we would have one identity function per type. We can abstract these types in to type variable, but we need to type these type variables as well to keep everything well typed. Consequently, abstracting the *types* over type variable, which itself is of sort *kinds* and represented as $*$, leads to *Second Order Lambda Calculus* ($\lambda 2$), and now the identity function over different types can be abstracted into a single function: $\lambda \alpha : \star.\lambda x : \alpha.x$. There are various other variants or abstractions of typed lambda calculus, which we would not discuss here, that can be categorized into:

- Terms depending on terms ($\lambda^{\to}$)

- Terms depending on types ($\lambda 2$)

- Types depending on types ($\lambda \underline{\omega}$)

- Types depending on terms (λP)

All these variation of lambda calculus can be captured into a unified framework known as *Pure Type System* or *Generalized Type System*. The *PTS* is group of type system that allows the dependencies between types and terms.

Unlike the simple typed lambda calculus ($\lambda^{\to}$) where terms and types live in two disjoint world, *PTS* blurs this distinction between types and terms. The abstract syntax of pure type system:

```
T = V (* variable *)
  | C (* constant *)
  | T T (* application *)
  | λ V : T. T (* abstraction*)
  | ∏ V : T. T (* dependent function type *)
```

The pure type system parametrized by a specification, i.e. set of sorts S, axioms A, and rules R, such that:

- S is a subset of C, i.e. $S \subseteq C$.

- A is the axioms of form $c : s$ where $c \in C$ and $s \in S$, i.e. $A \subseteq C \times S$.

- R is the set of rules of form (s_1, s_2, s_3) such that s_1, s_2, and $s_3 \in S$, i.e. $R \subseteq S \times S \times S$.

The typing judgement for *PTS* in typing context Γ is defined by following rules (s ranges of S, and x ranges over V with usual notion of variable capture avoidance):

- Axiom:
$$\frac{c : s \ \in A}{\Gamma \vdash c : s}$$

- Start:
$$\frac{\Gamma \vdash A : s}{\Gamma, x : A \vdash x : A}$$

- Weakening:
$$\frac{\Gamma \vdash A : B \quad \Gamma \vdash C : s}{\Gamma, x : C \vdash A : B}$$

- Product:
$$\frac{\Gamma \vdash A : s_1 \quad \Gamma, x : A \vdash B : s_2}{\Gamma \vdash (\prod x : A. B) : s_3}$$

- Application:
$$\frac{\Gamma \vdash F : (\prod x : A. B) \quad \Gamma \vdash a : A}{\Gamma \vdash Fa : B \, [x := a]}$$

- Abstraction:

$$\frac{\Gamma, \ x : A \ \vdash \ b : B \quad \Gamma \ \vdash \ (\prod x : A. \ B) : s}{\Gamma \ \vdash \ (\lambda x : A. \ b) : (\prod x : A. \ B)}$$

- Conversion:

$$\frac{\Gamma \ \vdash \ A : B \quad \Gamma \ \vdash \ B' : s \quad B =_\beta B'}{\Gamma \ \vdash \ A : B'}$$

3.1.1   Calculus of Construction/Inductive Construction

The Calculus of Construction is a higher order natural deduction style proof system for constructive proofs where every proof a typed λ-abstractions. Using the *Pure Type System* syntax, it can be expressed as:

$S =$
$$\{Prop\} \cup \{Type_i \mid i \in \mathbb{N}\}$$

$A =$
$$\{Prop : Type_0\} \cup \{Type_i : Type_{i+1} \mid i \in \mathbb{N}\}$$

$R =$
$$\left\{ \begin{array}{l} (Prop, Type_i, Type_i) \ i \in \mathbb{N} \\ (s, Prop, Prop) \ s \in S \\ (Type_i, Type_j, Type_{max(i,j)}) \end{array} \right\}$$

The sort *Prop* captures type of expression which represents logical proposition, while the sort *Type* captures the computational content. Calculus of Construction is powerful enough to encode inductive definitions [Pfenning and Paulin-Mohring, 1989], but one of the main drawback is efficiency of computation of function over these encoded inductive definitions, and some other properties could not be proven [Geuvers, 2001]. In order to solve these problems, [Paulin-Mohring, 1993] introduced Inductive definitions, pattern matching, and fixpoint in the Calculus of Construction to make the data structure representation more efficient. Below is the (incomplete) syntax of Calculus of Inductive Construction:

```
T = ... (* Pure Type System *)
| Ind { V : T := V : T}.V (* inductive definition)
| case T of V => T (* pattern matching *)
```

| fix$_n$ { V : T := T } (* recursion *)

We would like to point that the current underlying theory of Coq has been extended with Co-Inductive types [Giménez, 1995]; however, the discussion of Co-Inductive types is not very relevant for this thesis.

3.1.2   Type vs. Prop: Code Extraction

Every term in Coq has a type, and proofs are constructive in nature, so we can extract a computational program from a suitable proof. Moreover, every proof in Coq has two parts: (i) computational part, and (ii) logical proposition. The type of logical propositions are *Prop*, while the type of computational parts are *Type*. During the extraction process [Letouzey, 2008], every term of sort Prop is removed and no longer exists in extracted code, and only the terms of sort Type are translated into target language (OCaml/Haskell/Scheme). Because of this, Coq in general does not allow the case analysis on the terms (logical objects) of sort Prop when the goal is in not in Prop, but in certain cases it can be achieved (we call this special case reification and explained next).

3.1.2.1   Reification

Sometimes it is very natural to express certain properties/definitions in the Prop than the Type. Moreover, the definitions/terms in the Prop are very intuitive for human inspection and understanding. The only problem is that the terms of the sort Prop do not carry any computational content but only the proof part. However, we can escape this situation if the term of sort Prop is decidable and its domain is finite (in this case Prop is equivalent to bool, but this is not the case in general in *Constructive Logic* compared to *Higher Order Logic* with the law of excluded middle). In the case of decidable predicate in Prop over a finite domain, we can extract the witness constructively by using a program of linear search that tries the decidable predicate on every element of its (finite) domain. Below is the linear search reification code which produces a Type level witness, *existsT*, from a Prop level witness, *exists*, by iterating through all the elements of finite type *A* (the finiteness of *A* is captured by the list *l*).

```
Require Import Coq.Lists.List.
Import ListNotations.

(* type level existential quantifier *)
Notation "'existsT' x .. y , p" :=
  (sigT (fun x => .. (sigT (fun y => p)) ..))
    (at level 200, x binder, right associativity,
     format "'[' 'existsT'  '/  ' x  ..  y ,  '/  ' p ']'")
  : type_scope.

(* the following shows that a decidable (or boolean valued)
   predicate on a finite list
   can always be reified in terms of strong existence *)
Theorem reify {A: Type} (P: A -> bool) : forall (l: list A),
    (exists x, In x l /\ P x = true) -> existsT x, P x = true.
Proof.
  refine (
     fix Fn l :=
       match l with
       | [] => fun H => _
       | h :: tl => fun H => _
       end).
  contradict H. intro.
  destruct H as [x [H1 H2]].
  firstorder.

  assert (Hbiv: {P h  = true} + {P h <> true}).
  decide equality.
  destruct Hbiv as [Htrue | Hfalse].
  exists h. assumption.
  specialize (Fn tl). apply Fn.
  destruct H as [x [H1 H2]].
  destruct H1. subst.
  firstorder. exists x.
  firstorder.
Defined.
```

We have used many standard tricks like this to make our formalization more accessible for human inspection. For example, we have two definition, one in Prop and other in Type, of winner, loser and path. The rationale behind two definitions is that that Prop definitions are very natural and easy

to understand compared to their Type counter part. Furthermore, we have shown that they are equivalent to each other, and used the definitions in Type for computation. Also, there is a nice Coq library ConstructiveEpsilon[1] which uses the similar trick as ours; however, we have not used this library in our formalization.

3.1.3   Correct by Construction: Type Safe Printf

One of the highly sought feature of Coq is dependent type, a type which is parametrised by value. The expressiveness of dependent type make it possible to express specification at type level, and these specifications enables larger set of logical errors to eliminated at compile time.

Using the expressiveness of dependent type, we construct a type-safe version of printf [Pierce, 2004]. Our goal is to generate compiler error when the given format string and the type of corresponding input values do not match, e.g. printf "%d %s" "hello Coq" 42 should be compiler error because %d is a directive for integer value, but the type of input, "hello Coq", is string. In addition, type-safe printf should print the input when the format string is aligned with type of input, e.g. printf "%s %d" "hello Coq" 42 should print the string "hello Coq 42" because the first directive of format string, %s, and type of input, "hello Coq", are aligned. Similarly, the second directive of format string, %d, is also aligned with the type of input, 42.

The high level idea is to split the printf arguments into two parts: i) format string, and ii) values to be printed. For example, printf "%s %d" "hello Coq" 42 would be split into "%s %d", and "hello Coq" 42. Based on the format string, we design two functions: i) a type level function, and ii) a value level function. The type level function would take format string and returns a variadic function type, e.g. on format string "%s %d", it would return a function type with signature `string -> Z -> string`. The value level function, whose type signature is constructed by the type level function, would take the values to printed as input. If the type of values to be printed is aligned with the type constructed by the type level function then we proceed to print the string, otherwise we generate compiler error.

First, we defined a abstract syntax tree, *format*, to make it explicit the characters we are interested in format string. Additionally, the *format_string* function takes the format string and returns the abstract syntax tree, and

---

[1]https://coq.github.io/doc/master/stdlib/Coq.Logic.ConstructiveEpsilon.html

the type level, *interp_format*, takes the abstract syntax tree and returns the function type corresponding to format string.

```
(* abstract syntax tree *)
Inductive format :=
| Fend : format
| Fint : format -> format
| Fstring : format -> format
| Fother : ascii -> format -> format.

(* turn the format string into abstract syntax tree *)
Fixpoint format_string (inp : string) : format :=
  match inp with
  | EmptyString => Fend
  | String ("%"%char) (String ("d"%char) rest) =>
        Fint (format_string rest)
  | String ("%"%char) (String ("s"%char) rest) =>
      Fstring (format_string rest)
  | String c rest => Fother c (format_string rest)
  end.



(* construct the type level function from abstract syntax tree *)
Fixpoint interp_format (f : format) : Type :=
  match f with
  | Fint f => Z  -> interp_format f
  | Fstring f => string -> interp_format f
  | Fother c f => interp_format f
  | Fend => string
  end.
```

The *interp_format* function returns a function type (`Z -> string -> string -> string`) on the (abstract syntax tree of) format string "%d %s %s"

```
Eval compute in interp_format (format_string "%d %s %s").
(* = Z -> string -> string -> string
     : Type *)
```

Now, we construct a value level function, *interp_value*, whose type is constructed by type level function, and it will take the values to be printed as

input. The type of values to print should match the type constructed by type level function for successful type checking otherwise it will be type error.

```
(* value level function whose type is constructed on fly by interp_format
    function *)
Fixpoint interp_value (f : format) (acc : string) : interp_format f :=
  match f with
  | Fint f' => fun i => interp_value f' (acc ++ of_Z i)
  | Fstring f' => fun i => interp_value f' (acc ++ i)
  | Fother c f' => interp_value f' (acc ++ String c EmptyString)
  | Fend => acc
  end.
```

Finally, we define the printf function, and evaluate it on two inputs: i) printf "%d %s" "hello Coq" 42, and ii) printf "%d %s" 42 "hello Coq".

```
Definition printf s := interp_value (format_string s) "".

Eval compute in  printf "\%d \%s" "hello Coq"%string 42.
(* Error: The term ""hello Coq"%string" has type "string"
while it is expected to have type "Z". *)

Eval compute in  printf "\%d \%s" 42 "hello Coq"%string.
(*  "\0b101010 \hello Coq"%string. The number
42 is printed in binary *)
```

The first input, printf "%d %s" "hello Coq" 42, is type error because printf "%d %s" returns a value level function whose type is Z -> string -> string, but the type of first argument, "hello Coq", is string which does not unifies with Z, while second one is successfully printed as string.

3.1.4 Gallina: The Specification Language

The example, type safe printf function, I gave in previous section was encoded in Coq's specification language Gallina. Gallina is a highly expressive specification language for development of mathematical theories and proving the theorems about these theories; however, writing proofs in Gallina is very

tedious and cumbersome. Furthermore, It is not suitable for large proof development. In order to ease the proof development, Coq also provides tactics. The user interacting with Coq theorem prover applies these tactics to build the Gallina term which otherwise would be very laborious.

We have written two proofs that addition on natural number is commutative. First proof, *addition_commutative_gallina*, is written using Gallina, while the second proof, *addition_commutative_tactics*, is written using the tactics. In general, we write programs directly in Gallina and use tactics to prove properties about the programs. However, there is no fixed set of rules, and tactics can be used to write programs with dependent types (which we have done during this formalization).

```
(* proof written using Gallina *)
Lemma addition_commutative_gallina :
    forall (n m : nat), n + m = m + n.
refine
      (fix Fn (n : nat) : forall m : nat, n + m = m + n :=
        match n as n0
              return (forall m : nat, n0 + m = m + n0) with
        | 0 => fun m : nat =>
            eq_ind_r (fun n0 : nat => m = n0)
                                eq_refl (Nat.add_0_r m)
        | S n' =>
            fun m : nat =>
              eq_ind_r (fun n0 : nat => S n0 = m + S n')
                        (eq_ind_r (fun n0 : nat => S (m + n') = n0)
                        eq_refl (Nat.add_succ_r m n')) (Fn n' m)
        end).
Qed.

(* proof written using tactics *)
Lemma addition_commutative_tactics :
  forall (n m : nat), n + m = m + n.
  induction n; intro m; simpl; try omega.
Qed.
```

3.1.5 Trusting Coq proofs

In general, Coq proofs are nowhere similar to a mathematical proof written by trained mathematician. Also, these proofs are verbose and fairly long, so a very fundamental question is: why should we accept or believe in a proof written in Coq [Pollack, 1998]? Generally, the answer of accepting or trusting Coq proos is two fold: i) is the logic (CIC) sound?, and ii) is the implementation correct? The logic has already been reviewed by many peers and proved correct using some meta-logic, therefore the answer of our question about trusting Coq proof hinges on the implementation. Coq implementation (written in OCaml) has two parts, the type checker (small kernel), and tactic language to build the proofs. We lay our trust in type checker, because it is a small kernel and can be manually inspected. Furthermore, if there is a bug in tactic language, which often is the case, then build proof would not pass the type checker. Also, we can use the publicly available proof checkers written by experts and inspected by many others. In addition, to increase the confidence, there have been efforts to certify type checker [Appel et al., 2003] [Barras, 1996], verifying meta theory of one proof system in other [Anand and Rahli, 2014], self certificate of theorem prover [Harrison, 2006]. However, no system can prove its own consistency (Gödel's second incompleteness theorem), therefore trusting human judgement is inevitable.

3.2 Cryptography

The word cryptography comes from the two Greek words: *kryptós*, meaning *hidden*, and *gráfein* meaning *to write*. As a matter of fact, in the past, hidden writing (cryptography), using the symbol replacement, has been used to conceal the message. For example, the earliest known usage of cryptography (symbol replacement) goes back to ancient Egyptian (Khnumhotep II, 1500 BCE); however, the purpose of replacing one symbol by other was not to protect any sensitive information but to enhance the linguistic appeal. The first known usage of cryptography to conceal the sensitive information goes back Mesopotamians (1500 BCE) where they used it to hide the formula for pottery glaze. Fast forward, around 100 BCE, Julius Caesar wrote a letter to Marcus Cicero using a method, now known as *Caesar cipher*, which would shift each character in letter by 3 position right with wrapping around, i.e. X would wrap A, Y would wrap to B, and Z would wrap to C. Decryption was 3 character left shift. Using the tools of modern mathematics, encryption

and decryption in *Caesar cipher* is modular addition and modular subtraction (modulo 26), respectively. Overall, cryptography is art and science of making thing unintelligible from everyone, except the intended recipient.

The modern day cryptography originated in 1970 with two ingenious ideas, *Data Encryption Standard (DES)*, and *Diffie-Hellman Algorithm*. Data Encryption Standard, developed at IBM in 1970, is a symmetric key encryption algorithm which uses the same key for encryption and decryption. Since its inception, Data Encryption Standard amassed a bad reputation because of *National Security Agency (NSA)* involvement; however, it had a practicality issue, key management (sort of chicken egg problem). If two parties wanted to communicate securely over insecure channel using Data Encryption Standard, then they needed to agree on a common key. In order to agree on common key, they needed a secure channel where they can securely communicate the key. The solution to this problem came from *Diffie-Hellman* key exchange where two parties can exchange the key securely over insecure channel. Moreover, the advent of *Diffie-Hellman* key exchange started the whole new area of public key cryptography where encryption and decryption key are different. Although *Diffie-Hellman* key exchange suffers from man-in-the-middle (MITM) attack if used for keys exchange in its naivety, e.g. Logjam [Adrian et al., 2015], nonetheless, it is a building block for many other algorithm, e.g. ElGamal Encryption [ElGamal, 1985].

In this thesis, we are mostly concern about public key cryptography and will no longer discuss or explain symmetric key cryptography, hence any mention of cryptography hereafter should be assumed as public key cryptography. The basic working principals of modern day cryptography is based on mathematical principal than vanilla symbol replacement. In addition, it is no longer just used to achieving confidentiality, but various other things, e.g. integrity, authentication, non-repudiation protocol, digital signature, digital cash, etc. These mathematical principal involves various algebraic structures and algorithm to manipulate the object from these structures. For example, the underlying mathematical principal of *Diffie-Hellman* algorithm is hardness of computing discrete logarithm in finite field.

Now we describe the workings of `Diffie-Hellman` [Diffie and Hellman, 2006] algorithm, because all the constructions we have used are based on `Diffie-Hellman` construction. Before we describe the algorithm, we briefly sketch the algebraic structure Group because it is underlying algebraic structure of `Diffie-Hellman` construction (typically, the underlying structure is multiplicative group of a finite field). Also, note that our definition is influenced theorem-provers/type-theory because we have written the type signa-

ture of group operator $*$ and inverse operator *inv*.

3.2.1   Group

A group is a set G, with a binary operator $* : G \to G \to G$, identity element e, and inverse operator $inv : G \to G$ such that the following laws hold:

- `Associativity`: $\forall$ a b c $\in G$, $a * (b * c) = (a * b) * c$

- `Closure`: $\forall$ a b $\in G$, $a * b \in G$

- `Inverse Element`: $\forall$ a $\in G$ $\exists$ $a^{-1} \in G$, such that $a * a^{-1} = a^{-1} * a = e$. a^{-1} is called inverse of a (*inv a*).

- `Identity`: $\forall$ a $\in G$, $a * e = e * a = a$

Furthermore, if a group is commutative, i.e. $\forall$ a b $\in G$, $a * b = b * a$, then we call it abelian group (in honour of Niels Henrik Abel). In addition, if a *group* is *cyclic group* if it can be generated by a single element, also known as generator of group and denoted as g, by repeatedly applying the group operator $*$ to itself. Moreover, a group is *finite cyclic group* if it is cyclic and the cardinality of the underlying set (carrier set) G is finite. The cardinality is also known as order of group.

3.2.2   Diffie-Hellman Construction

Now we explain `Diffie-Hellman` construction. The construction can be divided into two steps:

1. The two communicating parties, say Alice and Bob, agree with shared public parameters which are finite cyclic group G of order p (p is a large prime) and generator element g.

2. After agreeing with public parameters, Alice and Bob initiates the key exchange protocol (assuming that Alice goes first):

   (a) Alice selects a random number a, where $1 < a < p$, computes g^a ($g * g * g... * g$ a times), and shares g^a to Bob.

(b) Similarly, Bob selects are random number b, where $1 < b < p$, computes g^b, and shares g^b with Alice.

(c) Finally, Alice computes the key $(g^b)^a$, and Bob computes the key $(g^a)^b$. A basic algebraic simplification on Alice's key and Bob's key would show that they both have the common key g^{ab}.

During the whole process, Eve, the adversary, would have g^a and g^b, but she can not compute the g^{ab} from these two values assuming that discrete logarithm is hard to compute. There are, off course, other attacks exists, e.g. denial of man in the middle attack, Logjam, etc. The security property of `Diffie-Hellman` construction is formalized using complexity theoretic notion given below (we would not go into the details of complexity theoretic notions):

DL - Discrete Logarithm problem: An instance of *DL* problem states that given a finite cyclic group G, a generator of g of G, and an element y, finding an element $x \in G$ such that $g^x = y$.

DH - Diffie-Hellman problem: An instance of *DH* problem states that given a finite cyclic group G, a generator of g of G, elements g^a and g^b, finding the element g^{ab}.

DDH - Decision Diffie-Hellman Problem: An instance of *DDH* problem states that given a finite cyclic group G, a generator of g of G, elements g^a, g^b, and g^c, determining if $c = a * b$.

3.2.3 El-Gamal Encryption Scheme

In 1985, Tahir El-Gamal [ElGamal, 1985] proposed a new encryption system which was based on Diffie-Hellman algorithm. *El-Gamal* turned the interactive Diffie-Hellman algorithm into a non-interactive, no need for any active second party, by introducing a randomness. The *ElGamal* scheme has three phases:

1. **Key Generation:** The user, say Alice, first chooses a finite-cyclic group G of order p (p is a large prime) and a group group generator g. She randomly selects a an element x from $\{1, \ldots, p-1\}$ as a private key, computes her public key $h = g^x$. Subsequently, she publishes the <G, g, p, h> and keeps x private.

2. **Encryption:** If any party, say Bob, wants to send a encrypted message *m* to Alice, then he would randomly select an element r, where 1 < r < p, computes $c_1 := g^r$ and $c_2 := m * h^r$, and send the pair (c_1, c_2) to Alice.

3. **Decryption:** Upon receiving any pair (c_1, c_2), Alice would compute $c_2 * c_1^{-x}$. A basic simplification of $c_2 * c_1^{-x}$ shows that it recovers the plaintext message. The simplification proceeds by replacing the c_2 with $m * h^r$ and c_1 with g^r in $c_2 * c_1^{-x}$. This substitution leads to $m * h^r * g^{-rx}$ which upon further simplification by replacing the *h* with g^x leads to $m * g^{xr} * g^{-rx}$. Using the same base rule, the term $m * g^{xr} * g^{-rx}$ can be written as $m * g^{xr-rx}$. Since $xr = rx$, so we can replace $m * g^{xr-rx}$ with $m * g^0$. The term $g^0 = e$ (the identity of group *G*) and using the right identity group law, we can replace $m * e$ by *m*.

3.2.4 Homomorphic Encryption

Homomorphic encryption is a encryption scheme which allows us to perform useful operation on encrypted data without decrypting the data. It was first posed by Rivest, Adleman and Dertouzos in [Rivest et al., 1978]:

> Consider a small loan company which uses a commercial time-sharing service to store its records. The loan company's "data bank" obviously contains sensitive information which should be kept private. On the other hand, suppose that the information protection techniques employed by the time sharing service are not considered adequate by the loan company. In particular, the systems programmers would presumably have access to the sensitive information. The loan company therefore decides to encrypt all of its data kept in the data bank and to maintain a policy of only decrypting data at the home office – data will never be decrypted by the time-shared computer.

A encryption scheme is homomorphic if for any two plaintext *x* and *y*:

> $Enc_{pk}(x) \otimes Enc_{pk}(y) = Enc_{pk}(x \oplus y)$ where *Enc* is encryption function, *pk* is the public key, $\otimes$ is operation on ciphertext, and $\oplus$ is operation on plaintext.

These two operators $\otimes$ and $\oplus$ are very specific. If a cryptosystem that supports an arbitrary function f on ciphertext, then it is called fully homomorphic cryptosystem:

$$f(Enc_{pk}(m_1), Enc_{pk}(m_2), ..., Enc_{[}pk(m_k) = Enc_{pk}(f(m_1, m_2, ..., m_k))$$

The first fully homomorphic encryption system was proposed by Craig Gentry [Gentry, 2009]; however, in this thesis we are mostly concern with partially homomorphic encryption (either additive or multiplicative, but not both), specifically additive ElGamal, so we are not going to present the details overview of Craig Gentry fully homomorphic construction. From now on, we would be using the term homomorphic encryption for partially homomorphic encryption.

Now, keeping in mind that homomorphic encryption enables us to perform useful operation on encrypted data, we will see what kind of homomorphic property is exhibited by the ElGamal method discussed in the previous section. Given a public infrastructure $<G, p, g, h>$ for ElGamal scheme, we encrypt two message m_1 and m_2 by taking two random numbers r_1, r_2 from the group:

$$Enc(m_1, r_1) := (g^{r_1}, m_1 * h^{r_1})$$

$$Enc(m_2, r_2) := (g^{r_2}, m_2 * h^{r_2})$$

If we mutiply these two cipher together pairwise, we get $(g^{r_1+r_2}, m_1 * m_2 * h^{r_1+r_2})$. After decrypting this combined ciphertext, we will get $m_1 * m_2$. In the scheme, $\otimes$ is multiplication and $\oplus$ is also multiplication. Furthermore, if our end goal is to achieve multiplication on a bunch of plaintext, then rather than decrypting the corresponding ciphertext individually and multiplying them, we could simply multiply all the ciphertext together and decrypt the final result. The advantage of this scheme is that it does not leak the individual values which, sometimes, is a very crucial property in many application, specifically election voting. In electronic voting protocols, we do not want to reveal the choices of a individual voter, but it is acceptable to reveal the final tally. However, this scheme is not suitable for electronic voting schemes because it is multiplicative. Almost, to the best of my knowledge, all the electronic voting scheme calculate the finally tally by adding the individual choices of all voters, so the requirement is achieve the addition on

plaintext. There are many additive homomorphic encryption schemes, e.g. Benaloh cryptosystem, Paillier cryptosystem, etc. In addition, we can modify the ElGamal encryption scheme to make additive. In additive case, it works as:

$$Enc(m_1, r_1) := (g^{r_1}, g^{m_1} * h^{r_1})$$

$$Enc(m_2, r_2) := (g^{r_2}, g^{m_2} * h^{r_2})$$

Multiplying these two ciphers pairwise would give us, $(g^{r_1+r_2}, g^{m_1+m_2} * h^{r_1+r_2})$ which would decrypt as $g^{m_1+m_2}$. In this case, is multiplication and $\oplus$ is addition. We can calculate the value of $m_1 + m_2$ by using linear search algorithm, or more efficient one Pohlig–Hellman algorithm. However, the downside of this scheme is that if the values of $m_1 + m_2 + \cdots + m_n$ (assuming n values) is very large, then calculating it from $g^{m_1+m_2+\cdots+m_n}$ is not very practical [Cramer et al., 1997].

3.2.5 Zero Knowledge Proof

In conventional mathematics, a proof of mathematical statement is collection of basic axioms combined according to rules of the system. For example, we want to prove that in for any group (G, *), for any two elements x y $\in$ G, we have:

$$(x * y)^{-1} = y^{-1} * x^{-1}$$

Proof: we assume arbitrary x, y. We show that $(x * y)^{-1}$ and $y^{-1} * x^{-1}$ are inverse of each other by combining them together using the group operator $*$ and using the group laws lead to the identity of the group G.

$$
\begin{aligned}
(x * y) * (y^{-1} * x^{-1}) &= x * y * y^{-1} * x^{-1} (associativity) \\
&= x * (y * y^{-1}) * x^{-1} (associativity) \\
&= x * e * x^{-1} (inverses) \\
&= x * x^{-1} (identity) \\
&= e (inverse)
\end{aligned}
$$

$$(3.1)$$

Similarly, we can prove that $((y^{-1} \cdot x^{-1}) \cdot (x \cdot y) = e$. We can also formalize it inside theorem prover and prove it more formally (below is a proof in Coq theorem prover).

```
Lemma inv_distr : forall a b, inv (f a b) = f (inv b) (inv a).
Proof.
  intros a b. symmetry.
  apply inv_uniq_l.
  rewrite <- assoc.
  rewrite  (assoc (inv b) (inv a) a).
  rewrite (inv_l a).
  rewrite (assoc (inv b) e b).
  rewrite (id_l b).
  rewrite (inv_l b). auto.
Qed.
```

If a verifier wants to verify the correctness of our proof, then she would simply check that if the group rules are applied correctly. Moreover, these proofs are static in nature, i.e. once the prover has produced the proof, then the content of proof is not going to change over time, and there would not be any interaction between prover and verifier if verifier wants to verify the proof. In addition, the verifier not only learned that the statement is true, but she also learned the content of proof (gained some knowledge).

In contrast, zero-knowledge-proof, first introduced by Goldwasser, Micali, and Rackoff [Goldwasser et al., 1985] , is probabilistic proof that involves the explicit notion of a interaction between the prover and verifier. In addition, the goal of the prover is to convince the verifier about the validity of some statement without revealing any information, i.e. the only thing verifier would learn is that if statement is true or false without any other information. More formally, zero-knowledge-proof for a language $L \in \{0,1\}^*$ (generally NP) is a interactive proof between a (computationally unbounded) prover P and a (polynomial time) verifier V. Furthermore, the goal of P is to convince V that x $\in$ L such that:

Completeness: If x $\in$ L then the honest prover P would convince the honest verifier V to accept the claim with overwhelming probability. If P can always convince (probability 1) the V that x $\in$ L, then the proof system has perfect completeness.

Soundness: If x $\notin$ L then dishonest prover P^* can not convince the honest

verifier V to accept the claim (with some small probability error known as soundness error)

Zero knowledge: A malicious verifier V^* would gain no additional information by interacting with a honest prover P other than x $\in$ L. More formally, for every (polynomial time) program V^* there exists a (polynomial time) program S, also known as simulator, which can produce the transcript of protocol by itself without interactive with anyone. Moreover, the transcript produced by simulator S is indistinguishable from real transcript produced by interaction between

3.2.5.1   Zero Knowledge Proof of Knowledge

Sometimes, the fact that $x \in L$ is completely trivial. For example, for any given finite group G of order p (p is prime), a random element h from the group G, and generator g of the group G, a prover claims that there is a x such that $g^x = h$. This is trivial because we know that there always exists such x (discrete logarithm problem); however, the challenge is to show that the prover knows the witness x. Formally, zero-knowledge-proof of knowledge is defines as: let $R = (x, w) \subset LxW$ is a binary relation such that $x \in L$ is common string between prover P and verifier V and $w \in W$, also known as witness, is private to the prover P. Moreover, the goal of prover P is to convince verfier V that $(x, w) \in R$ in zero-knowledge.

3.2.6   Sigma Protocol

Sigma protocols are efficient way to achieve zero-knowledge-proof of a knowledge. Sigma protocol is a three step communication between a prover P and a verifier V where goal of the prover is to convince the verifier that she knows witness w for common input x such that $(x, w) \in R$.

1. P sends a message r

2. V sends a random string e

3. P replies with z

Based on public inputs (x, r, e, z), the verifier V decides to accept or reject the proof. A protocol is said to be sigma protocol for a relation R if:

Completeness: when prover and verifier follow the protocol for public input x and witness w then verifier accepts the proof

Special Soundness: For a given pubic input x, if prover can produce two accepting transcript (a, e, z) and (a, e', z') (e and e' are disjoint), then there exists a efficient program, extractor, which can extract the witness w.

Honest Verifier Zero Knowledge: For a given public input x and random input e, there is a simulator which outputs an accepting transcript (a, e, z) which is indistinguishable from a proof generated by a prover interacting with honest verifier.

A concrete example of sigma protocol is Schnorr protocol [Cramer et al., 1994]. In this example, the goal of a prover P is to prove the knowledge of discrete log in a Group of order p (p is prime) to a verifier V. Furthermore, g is the generator of group G, x is the public input and w is private input with relation $x = g^w$. The protocol follows:

- Prover P randomly selects an element r from $[0 \cdots q)$, computes $a = g^r$ and sends a to verifier V

- Verifier V randomly selects an element c from $[0 \cdots q)$ and sends it to P

- Prover P sends $z = r + c * w$ to V. V checks $g^z = a * x^c$

For the protocol described above, all three properties, completeness, special soundness, and honest verifier zero knowledge, hold.

- Completeness holds with probability 1. Simplifying the expression g^z shows that it is equal to $a * x^c$. Replacing the z by $r + c * w$ in expression g^z, we get g^{r+c*w}. Using addition rule of power, g^{r+c*w} can be simplified as $g^r * g^{c*w}$. First first step of protocol, $a = g^r$, so we can replace the $g^r * g^{c*w}$ by $a * g^{c*w}$. From the group infrastructure, we have $x = g^w$, so we can write x at place of g^w, therefore, $a * g^{c*w}$ transforms into $a * x^c$.

- Special soundness holds. For any two given response, $z_1 = r + w * c_1$ and $z_2 = r + w * c_2$, we can find the witness w by $(z_2 - z_1)/(c_2 - c_1)$.

- Honest Verifier Zero Knowledge also holds. Simulator can always produce a transcript $(g^z x^{-c}, c, z)$ by randomly choosing c (the random choice c is the reason for special honest verifier zero knowledge), and z.

Fiat-Shamir Transform: In practice, the *Fiat-Shamir* transform is used to turn a Sigma protocol into a non-interactive proof. As a consequence, there is no longer any interaction with verifier. A *Fiat-Shamir* transform to sigma protocol is:

- Prover P randomly selects an element r from $[0 \cdots q)$, computes $a = g^r$

- Prover P computes $c = H(a||x)$ where H is a hash function

- Prover P computer $z := r + c * w$

Finally, P publishes the transcript (a, c, z) for anyone to verify her claim. Subsequently, any one who is verifying the claim has to check two things: (i) $c := H(a||x)$, and (ii) $g^z = a * x^c$.

3.2.7 Commitment Schemes

Commitment schemes are cryptographic primitives equivalent to real life sealed lock-box. Once the lock-box is locked and sealed, the content inside it can not be changed without breaking the lock and seal. In general, commitment primitives are backbone of any cryptographic protocol between two parties, communicating over internet, to force them to follow the protocol honestly, even they would have a huge gain from deviating from protocol. For example, in order to save some time before a match, Indian cricket team captain, living in Delhi, and Australian cricket team captain, living in Canberra, decide to toss a coin in advance over the Internet, using a mobile application called toss-app, for a upcoming series of one-day matches [2]. Assuming the workings of toss-app is naive, i.e. one captain is going to post his call in the chat box, and the other other captain is going to toss the coin at their end and post the outcome in chat box. Furthermore, the decision is taken based on the messages posted by the two captains. In this scenario, we are assuming that the captain, who is tossing the coin, is honest and posting the outcome of the coin honestly in the chat box, which could or could not be the case. The question is can we devise some scheme which would force the both parties to behave honestly? The answer is yes, we can devise such scheme. We would use the sealed lock-box concept, albeit the digital one. Moreover, the first captain would put his call in digitally sealed lock-box and post it in the

---

[2]In a cricket match, which is very popular sport in India and Australia, both captains meet in the ground and toss a coin to decide who would have the first call.

chat box. Because it is sealed and locked, the other captain would have no idea what is the content inside it. Furthermore, it is impossible to break the lock box, so it is fruitless and waste of time for the other captain to even try. The other captain will toss the coin and post the outcome in a separate digital sealed lock box. Now that we two digital sealed lock box which can only be opened by the respective owners, they would move for the next phase of coin tossing called revealed phase. In the revealed phase, they both would open their sealed locked box to show that what they have locked, and the decision would be taken accordingly [3].

Formally, a commitment scheme is three step protocol between a sender *S* and a receiver *R*:

1. Commit phase: sender *S* commits a value *m* by generating a random number *r* and using some algorithm *C*, which takes the message and random *r*. Moreover, the committed value produced by the commitment algorithm *C*, $c = C(m, r)$, is shared with receiver *R*.

2. Reveal phase: In the reveal phase, the sender reveals the message *m* and randomness *r* which are subsequently used by receiver to verify the result, i.e. the receiver computes $c' = C(m, r)$ and matches it again the given *c* in the commit phase of protocol.

Security Properties: Commitment schemes have to have two properties: hiding and binding. Hiding property ensures that the receiver can not recover or recompute the original message *m* from the given commitment *c*, i.e. it forces the receiver to behave honestly in the protocol. Furthermore, binding property ensures that it is impossible for sender to come up with another message *m'* which is different from *m* but produces the same commitment *c*, i.e. it forces the sender to behave honestly in the protocol.

Pedersen commitment: Finally, we give a brief overview of a Pedersen commitment scheme which is based on discrete log. The protocol as follows assuming the public parameter available to sender and receiver, i.e. the set up has been conducted to generate the the public parameter, and both parties have these values. These values include a prime *p*, *y* a randomly chosen element from Z_p^*, and *g* a randomly chosen generator from Z_p^*.

- Commit phase: The sender generates a random *r* from Z_p^*, computes commitment $c = g^r * y^m$ and sent the commitment to receiver

---

[3]Story influenced by Manuel Blum's coin flipping by telephone

- Verification phase: In verification phase, the sender reveals the original message m and the randomness r. Finally, the receiver computes $g^r * y^m$. If the computed value matches with the commitment received in commitment phase, then she accepts it otherwise reject it.

3.3 Summary

In this chapter, we gave a brief summary of Coq theorem prover and cryptographic notation needed to understand the further chapters. By no means, these descriptions were exhaustive. For a detailed treatment of Coq theorem prover, [Bertot et al., 2004] [Chlipala, 2013] can be referred, and for cryptography, [Menezes et al., 1996] [Schneier, 1995] [Paar and Pelzl, 2009] can be referred. In the next chapter, we will discuss the Schulze method, and the machinery for its formalization.

Schulze Method : Evidence Carrying Computation

The negligence of a few could easily send a ship to the bottom, but if it has the wholehearted co-operation of all on board it can be safely brought to part.

Sardar Vallabhbhai Patel

4.1 Introduction

Correctness and verifiability/evidence are two main pillar of any democratic election. In case of paper ballot election, correctness and verifiability of counting is achieved by public scrutiny because each step is carefully observed by general member of public, and agents from different political parties. For example, casting ballot at booth is carefully observed by polling agents and counting ballots is observed by the scrutineers (Figure 4.1) appointed by different political parties. Given that electronic voting is relatively young, in this chapter we investigate how to achieve the correctness and verifiability similar to paper ballot election.

Chapter overview: In this chapter, we explain the Schulze method in section 4.2, and its formal specification in the section 4.3. The corner stone of our formalization is a correct by construction dependent inductive data type that represents all correct executions (4.3.1) with the formal proof of

Figure 4.1: Scrutineers, in green jacket, observing the ballot counting

that every Schulze election have winners (4.3.2). Every inhabitant of this dependent inductive data type not only produces a final result, but also all the intermediate steps which lead to the notion of evidence or scrutiny sheet (section 4.4). In the section 4.5, we discuss the optimization techniques to overcome the deficiencies in extracted Haskell code from Coq formalization. Based on these optimizations, the extracted Haskell code was able to count millions of ballots in few minutes. Finally, we conclude the chapter in the section 4.6 with the achievements and drawbacks of our work on the scale of *Correctness*, *Privacy*, and *Verifiability*.

4.2   Schulze Algorithm

The Schulze Method [Schulze, 2011] is a vote counting scheme that elects a single winner, based on preferential votes. The method itself rests on the relative *margins* between two candidates, i.e. the number of voters that prefer one candidate over another. The margin induces an ordering between candidates, where a candidate c is more preferred than d, if more voters pre-

fer c over d than vice versa. One can construct simple examples (see e.g. [Rivest and Shen, 2010]) where this order does not have a maximal element (a so-called *Condorcet Winner*). Schulze's observation is that this ordering *can* be made transitive by considering sequences of candidates (called *paths*). Given candidates c and d, a *path* between c and d is a sequence of candidates $p = (c, c_1, \ldots, c_n, d)$ that joins c and d, and the *strength* of a path is the minimal margin between adjacent nodes. This induces the *generalised margin* between candidates c and d as the strength of the strongest path that joins c and d. A candidate c then wins a Schulze count if the generalised margin between c and any other candidate d is at least as large as the generalised margin between d and c. More concretely:

- Consider an election with a set of m candidates $C = \{c1, \ldots, cm\}$, and a set of n votes $P = \{b1, \ldots, bn\}$. A vote is represented as function $b : C \to \mathbb{N}$ that assigns natural number (the preference) to each candidate. We recover a strict linear preorder $<_b$ on candidates by setting $c <_b d$ if $b(c) > b(d)$.

- Given a set of ballots P and candidate set C, we construct graph G based on the margin function $m : C \times C \to \mathbb{Z}$. Given two candidates $c, d \in C$, the *margin* of c over d is the number of voters that prefer c over d, minus the number of voters that prefer d over c. In symbols:

$$m(c,d) = \sharp\{b \in P \mid c >_b d\} - \sharp\{b \in P \mid d >_b c\}$$

  where $\sharp$ denotes cardinality and $>_b$ is the strict (preference) ordering given by the ballot $b \in P$.

- A directed *path* in the graph, G, from candidate c to candidate d is a sequence $p \equiv c_0, \ldots, c_{n+1}$ of candidates with $c_0 = c$ and $c_{n+1} = d$ ($n \geq 0$), and the *strength*, st, of path, p, is the minimum margin of adjacent nodes, i.e.

$$st(c_0, \ldots, c_{n+1}) = \min\{m(c_i, c_{i+1}) \mid 0 \leq i \leq n\}.$$

- For candidates c and d, let $M(c,d)$ denote the maximum strength, or generalized margin of a path from c to d i.e.

$$M(c,d) = \max\{st(p) : \text{p is path from c to d in G}\}$$

- The winning set is, always non empty and formally proved in 4.3.2,

defined as

$$W = \{c \in C : \forall d \in C \setminus \{c\}, M(c,d) \geq M(d,c)\}$$

The Schulze method stipulates that a candidate $c \in C$ is a *winner* of the election with margin function m if, for all other candidates $d \in C$, there exists a number $k \in Z$ such that

- there is a path p from c to d with strength $st(p) \geq k$

- all paths q from d to c have strength $st(q) \leq k$.

Informally speaking, we can say that candidate c *beats* candidate d if there's a path p from c to d which stronger than any path from d to c. Using this terminology, a candidate c is a winner if c cannot be beaten by any (other) candidate.

4.2.1 An Example

For some given set of ballots (the actual set of ballots are not very important because we want to demonstrate the Condercet Paradox), we have computed the margin function m such that $m(A,B) = 3$, $m(B,A) = $ -3, $m(A,C) = $ -1, $m(C,A) = 1$, $m(B,C) = 5$, and $m(C,B) = $ -5. We have drawn the graph below (Figure 4.2), and it shows that collective can be cyclic, even if the preferences of individual voters are not cyclic. This phenomena is known as Condercet paradox and first observed by french philosopher Marquis de Condorcet in late 18th century [1].

Figure 4.2: Margin Function/Matrix (Graph Interpretation)

---

[1]https://gallica.bnf.fr/ark:/12148/bpt6k417181

The main idea of the method is to resolve cycles by considering *transitive preferences* or a generalised notion of margin. Figure 4.3 is the graph interpretation of generalised margin, M, after running the Schulze method on margin function m (the word margin function is used interchangeably with margin matrix). In order to compute $M(A, B)$, we first compute all the path from candidate A to B. Well, we have just two path from A to B, a direct path between them and a intermediate path via candidate C. Now that we two paths, we compute the path strength st for each path, $st(A, B) = min\{m(A, B)\}$ and $st(A, C, B) = min\{m(A, C), m(C, B)\}$. Simply these expression:

$$st(A, B) = \min\{m(A, B)\}$$
$$= \min\{3\}$$
$$= 3$$

$$st(A, C, B) = \min\{m(A, C), m(C, B)\}$$
$$= \min\{-1, -5\}$$
$$= -5$$

Once we have the path strength for every path between A and B, we compute generalize margin $M(A, B) = \max\{st(A, B), st(A, C, B)\}$, and simplification leads to 3; hence the arrow going A to B has strength 3.

$$M(A, B) = \max\{st(A, B), st(A, C, B)\}$$
$$= \max\{3, -5\}$$
$$= 3$$

Similarly, we can compute other values as well.

$$st(B, A) = \min\{m(B, A)\}$$
$$= \min\{-3\}$$
$$= -3$$

$$st(B, C, A) = \min\{m(B, C), m(C, A)\}$$
$$= \min\{5, 1\}$$
$$= 1$$

$$M(B, A) = \max\{st(B, A), st(B, C, A)\}$$
$$= \max\{-3, 1\}$$
$$= 1$$

$$st(A, C) = \min\{m(A, C)\}$$
$$= \min\{-1\}$$
$$= -1$$

$$st(A, B, C) = \min\{m(A, B), m(B, C)\}$$
$$= \min\{3, 5\}$$
$$= 3$$

$$M(A, C) = \max\{st(A, C), st(A, B, C)\}$$
$$= \max\{-1, 3\}$$
$$= 3$$

$$st(C, A) = \min\{m(C, A)\}$$
$$= \min\{1\}$$
$$= 1$$

$$st(C, B, A) = \min\{m(C, B), m(B, A)\}$$
$$= \min\{-5, -3\}$$
$$= -5$$

$$M(C, A) = \max\{st(C, A), st(C, B, A)\}$$
$$= \max\{1, -5\}$$
$$= 1$$

$$st(C, B) = \min\{m(C, B)\}$$
$$= \min\{-5\}$$
$$= -5$$

$$st(C, A, B) = \min\{m(C, A), m(A, B)\}$$
$$= \min\{1, 3\}$$
$$= 1$$

$$M(C, B) = \max\{st(C, B), st(C, A, B)\}$$
$$= \max\{-5, 1\}$$
$$= 1$$

$$st(B, C) = \min\{m(B, C)\}$$
$$= \min\{5\}$$
$$= 5$$

$$st(B, A, C) = \min\{m(B, A), m(A, C)\}$$
$$= \min\{-3, -1\}$$
$$= -3$$

$$M(B, C) = \max\{st(B, C), st(B, A, C)\}$$
$$= \max\{5, -3\}$$
$$= 5$$

Now we have computed all the values of generalize margin M, we can interpret it as a graph show below. It is clear from the graph that candidate A is winner, as she beats B with strength 3 (reverse path from B to A is weaker, i.e. strength 1) and C with strength 3 (reverse path from C to A is weaker, i.e. strength 1).

Figure 4.3: Generalised Margin (Graph Interpretation)

4.3 Formal Specification

We start our Coq formalization assuming finite and non-empty set of candidates. Also, we assume decidable equality on candidates. For our purposes, the easiest way of stipulating that a type be finite is to require existence of a list containing all inhabitants of this type [Firsov and Uustalu, 2015].

```
Variable cand : Type.
Variable cand_all : list cand.
```

```
Hypothesis cand_fin : forall c: cand, In c cand_all.
Hypothesis dec_cand : forall n m : cand, {n = m} + {n <> m}.
Hypothesis cand_in : cand_all <> [].
```

For the specification of winners of Schulze elections, we take the margin function as given for the moment (and later construct it from the incoming ballots). In Coq, this is conveniently expressed as a variable:

```
Variable marg : cand -> cand -> Z.
```

We formalise the notion of path and strength of a path by means of a single (but ternary) inductive proposition that asserts the existence of a path of strength $\geq k$ between two candidates, for $k \in Z$. The notion of winning candidate is that it beats every other candidate, i.e. all the paths from the winner to other candidates are at least as strong as the reverse path. Dually, the notion of loser is that there is a candidate who beats the loser, i.e. the path from the candidate to the loser is stronger than the reverse path.

```
(* prop-level path *)
Inductive Path (k: Z) : cand -> cand -> Prop :=
  | unit c d : marg c d >= k -> Path k c d
  | cons  c d e : marg c d >= k ->
        Path k d e -> Path k c e.

(* winning condition of Schulze Voting *)
Definition wins_prop (c: cand) :=
  forall d: cand, exists k: Z,
    Path k c d /\ (forall l, Path l d c -> l <= k).

(* dually, the notion of not winning: *)
Definition loses_prop (c : cand) :=
  exists k: Z, exists  d: cand,
    Path k d c /\ (forall l, Path l c d -> l < k).
```

We reflect the fact that the above are *propositions* in the name of the definitions, in anticipation of type-level definitions of these notions later. The reason for having a *Prop* level definition is that it is very easy and intuitive for human to inspect the definitions, and ascertain the correctness of formalization. As we

discussed in the **Type vs. Prop** (section 3.1.2), the main reason for having an equivalent type-level versions of the above is that purely propositional information is discarded during program extraction, unlike the type-level notions of winning and losing that represent evidence of the correctness of the determination of winners. Our goal is to not only compute winners and losers according to the definition above, but also to provide independently verifiable evidence, a scrutiny sheet or certificate, of the correctness of our computation. The propositional definitions of winning and losing above serve as a reference to calibrate their type level counterparts, and we demonstrate the equivalence between propositional and type-level conditions in the next section.

One of the fundamental question about the declaring some one as a winner or loser is that how can we know that, say, a candidate c in fact wins a Schulze election, and that, say, d is not a winner? One possible answer is simply re-run an independent implementation of the method (usually hoping that results would be confirmed). But what happens if results diverge?

One major aspect of our work is that we can answer this question by not only computing the set of winners, but in fact presenting *evidence* for the fact that a particular candidate does or does not win. This is a re-emphasis on *Correctness*, and convincing to all, specifically to losers, leaving no ground for speculation. As we stated earlier that in the context of electronic vote counting, this is known as a *scrutiny sheet, or certification*: a tabulation of all relevant data that allows us to verify the election outcome. Again drawing on an already computed margin function, to demonstrate that a candidate c wins, we need to exhibit an integer k for all competitors d, together with

- evidence for the existence of a path from c to d with strength $\geq k$

- evidence for the non-existence of a path from d to c that is stronger than k

The first item is straight forward, as a path itself is evidence for the existence of a path, and the notion of path is inductively defined. For the second item, we need to produce evidence of membership in the *complement* of an inductively defined set.

Mathematically, given $k \in Z$ and a margin function $m : C \times C \to Z$, the pairs $(c, d) \in C \times C$ for which there exists a path of strength $\geq k$ that joins both are precisely the elements of the least fixpoint $LFP(V_k)$ of the monotone

operator $V_k : Pow(C \times C) \rightarrow Pow(C \times C)$, defined by

$$V_k(R) = \{(c,e) \in C \mid m(c,e) \geq k \text{ or } (m(c,d) \geq k \text{ and } (d,e) \in R \text{ for some } d \in C)\}$$

It is easy to see that this operator is indeed monotone, and that the least fixpoint exists, e.g. using Kleene's theorem [Stoltenberg-Hansen et al., 1994]. To show that there is *no* path between d and c of strength $> k$, we therefore need to establish that $(d,c) \notin LFP(V_{k+1})$.

By duality between least and greatest fixpoints, we have that

$$(c,d) \in C \times C \setminus LFP(V_{k+1}) \iff (c,d) \in GFP(W_{k+1})$$

where for arbitrary k, $W_k : Pow(C \times C) \rightarrow Pow(C \times C)$ is the operator dual to V_k, i.e.

$$W_k(R) = C \times C \setminus (V_k(C \times C \setminus R))$$

and $GFP(W_k)$ is the greatest fixpoint of W_k. As a consequence, to demonstrate that there is *no* path of strength $> k$ between candidates d and c, we need to demonstrate that $(d,c) \in GFP(W_{k+1})$. By the Knaster-Tarski fixpoint theorem [Tarski, 1955], this greatest fixpoint is the supremum of all W_{k+1}-coclosed sets, that is, sets $R \subseteq C \times C$ for which $R \subseteq W_{k+1}(R)$. That is, to demonstrate that $(d,c) \in GFP(W_{k+1})$, we need to exhibit a W_{k+1}-coclosed set R with $(d,c) \in R$. If we unfold the definitions, we have:

$$W_k(R) = \{(c,e) \in C^2 \mid m(c,e) < k \text{ and } (m(c,d) < k \text{ or } (d,e) \in R \text{ for all } d \in C)\}$$

so that given *any* fixpoint R of W_k and $(c,e) \in W$, we know that (i) the margin between c and e is $< k$ so that there's no path of length 1 between c and e, and (ii) for any choice of midpoint d, either the margin between c and d is $< k$ (so that $c, d, \ldots$ cannot be the start of a path of strength $\geq k$) or we don't have a path between d and e of strength $\geq k$. We use the following terminology:

Definition 1 *Let $R \subseteq C \times C$ be a subset and $k \in \mathbb{Z}$. Then R is W_k-coclosed, or simply k-coclosed, if $R \subseteq W_k(R)$.*

Mathematically, the operator W_k acts on subsets of $C \times C$ that we think of as predicates. In Coq, we formalise these predicates as boolean valued functions and obtain the following definitions where we isolate the function `marg_lt` (that determines whether the margin between two candidates is less than a given integer) for clarity:

```
Definition marg_lt (k : Z) (p : (cand * cand)) :=
  Zlt_bool (marg (fst p) (snd p)) k.

Definition W (k : Z) (p: cand * cand -> bool)
  (x: cand * cand) := andb  (marg_lt k x)
  (forallb (fun m => orb (marg_lt k (fst x, m))
                          (p (m, snd x))) cand_all).
```

In order to formulate type-level definitions, we need to promote the notion of path from a Coq proposition to a proper type, and formulate the notion of *k*-coclosed predicate.

```
Definition coclosed (k : Z) (f : (cand * cand) -> bool) :=
  forall x, f x = true -> W k f x = true.

Inductive PathT (k: Z) : cand -> cand -> Type :=
| unitT c d :  marg c d >= k -> PathT k c d
| consT c d e : marg c d >= k ->
        PathT k d e -> PathT k c e.
```

Now, we have following type-level definition of winning (and dually, nowinning) for Schulze counting. As we see that these definition not only produces the result, but they also produce witness, e.g. the `wins_type` definition states that if a candidate, say *c*, is the winner, then for each individual candidates participating in election, it produces two witness: (i) a path from itself to the beating candidate of certain strength, say *k*, and (ii) a *k*+1 co-closed set. These witness are basic building blocks of the scrutiny sheet we produce after election.

```
Definition wins_type c :=
  forall d : cand, existsT (k : Z),
  ((PathT k c d) * (existsT (f : (cand * cand) -> bool),
   f (d, c) = true /\ coclosed (k + 1) f))%type.

Definition loses_type (c : cand) :=
  existsT (k : Z) (d : cand),
  ((PathT k d c) * (existsT (f : (cand * cand) -> bool),
    f (c, d) = true /\ coclosed k f ))%type.
```

We have two definitions of winning, `wins_prop` which is easier for a human to inspect; on the other hand, `wins_type` which is useful for the machine. We close the gap by formally establishing that type level winning and prop level winning (dually, not winning) are in fact equivalent.

```
Lemma wins_type_prop :
    forall c, wins_type c -> wins_prop c.

Lemma wins_prop_type :
    forall c, wins_prop c -> wins_type c.

Lemma loses_type_prop :
    forall c, loses_type c -> loses_prop c.

Lemma loses_prop_type :
    forall c, loses_prop c -> loses_type c.
```

The different nature of the two propositions does not allow us to claim an equivalence between both notions, as Coq defines bi-implication only on propositions.

The proof of the first statement, *wins_type_prop*, is completely straight forward, as the type, *win_type*, carries all the information needed to establish the propositional winning, *wins_prop*. However, for the second statement *wins_prop_type*, Coq does not allow the case analysis or induct on a term of sort `Prop` when the sort of goal is not in `Prop`. We follow the techniques that we have described in `Type vs Prop` section (3.1.2.1). To prove the second statement, we first introduced an intermediate lemma based on the *iterated margin function* $M_k : C \times C \to Z$. Intuitively, $M_k(c, d)$ is the strength of the strongest path between c and d of length $\leq k + 1$. Formally, $M_0(c, d) = m(c, d)$ and

$$M_{i+1}(c, d) = \max\{M_i(c, d), \max\{\min\{m(c, e), M_i(e, d) \mid e \in C\}\}\}$$

for $i \geq 0$. It is intuitively clear (and we establish this fact formally) that the iterated margin function stabilises at the n-th iteration (where n is the number of candidates), as paths with repeated nodes don't contribute to maximising the strength of a path. This proof loosely follows the evident pen-and-paper proof given for example in [Carr'e, 1971] that is based on cutting out segments of paths between repeated nodes and so reaches a fixed point.

```
Lemma iterated_marg_fp: forall (c d : cand) (n : nat),
    M n c d <= M (length cand_all) c d.
```

That is, the *generalised margin,* i.e. the strength of the strongest (possibly infinite) path between two candidates is effectively computable.

This allows us to relate the propositional winning conditions to the iterated margin function and showing that a candidate *c* is winning implies that the generalised margin between this candidate and any other candidate *d* is at least as large as the generalised margin between *d* and *c*.

```
Lemma wins_prop_iterated_marg (c : cand) : wins_prop c ->
    forall d, M (length cand_all) d c <=
    M (length cand_all) c d.
```

This condition on iterated margins can in turn be used to establish the type-level winning condition, thus closing the loop to the type level winning condition.

```
Lemma iterated_marg_wins_type (c : cand) : (forall d,
    M (length cand_all) d c <= M (length cand_all) c d)
    -> wins_type c.
```

Similarly, we connect the propositional losing to type level losing via generalized margin. We show that candidate *c* is losing then there is a candidate *d* and generalized margin between candidate *d* and *c* is more that generalized margin between *c* and *d*. Using this fact, we can prove the type level losing condition.

```
Lemma loses_prop_iterated_marg (c : cand):
    loses_prop c ->
    (exists d, M (length cand_all) c d <
    M (length cand_all) d c).

Lemma iterated_marg_loses_type (c : cand) :
    (exists d, M (length cand_all) c d <
    M (length cand_all) d c)
    -> loses_type c.
```

The proof of lemma *iterated_marg_loses_type* is not straight forward because we are in a similar situation as we were in *wins_prop_type*. We can not eliminate `exists n, P n` in order to show `existsT n, P n`, because Coq would not allows to do case analysis on `exists n, P n` (a term of type `Prop`) since the goal, `existsT n, P n` (a term of type `Type`), is not in `Prop`. We again follow the technique described in `Type vs Prop` section (3.1.2.1). We do a linear search on list of candidates to find the witness constructively, and since, the list of candidates is finite we would eventually terminate and find one. This completes our loop of prop level loser to type level loser.

```
Corollary reify_opponent (c: cand):
  exists  d,
  M (length cand_all) c d < M (length cand_all) d c ->
  existsT d,
  M (length cand_all) c d < M (length cand_all) d c.
```

The crucial part of establishing the type-level winning conditions in the proof of the lemma above is the construction of a co-closed set. First note that `M (length cand_all)` is precisely the generalised margin function. Writing g for this function, we assume that $g(c,d) \geq g(d,c)$ for all candidates d, and given d, we need to construct a $k+1$-coclosed set S where $k = g(c,d)$. One option is to put $S = \{(x,y) \mid g(x,y) < k+1\}$. As every i-coclosed set is also j-coclosed for $i \leq j$, the set $S' = \{(x,y) \mid g(x,y) < g(d,c)+1\}$ is also $k+1$-coclosed and (in general) of smaller cardinality. We therefore witness the existence of a $k+1$-coclosed set with S' as this leads to certificates that are smaller in size and therefore easier to check.

We note that the difference between the type-level and the propositional definition of winning is in fact more than a mere reformulation. As remarked before, one difference is that purely propositional evidence is erased during program extraction so that using just the propositional definitions, we would obtain a determination of election winners, but no additional information that substantiates this (and that can be verified independently). The second difference is conceptual: it is easy to verify that a set is indeed coclosed as this just involves a finite (and small) amount of data, whereas the fact that *all* paths between two candidates don't exceed a certain strength is impossible to ascertain, given that there are infinitely many paths.

In summary, determining that a particular candidate wins an election based on the `wins_type` notion of winning, the extracted program will *additionally* deliver, for all other candidates,

- an integer k and a path of strength $\geq k$ from the winning candidate to the other candidate

- a co-closed set that witnesses that no path of strength $> k$ exists in the opposite direction.

It is precisely this additional data, which we call scrutiny sheet, (on top of merely declaring a set of election winners) that allows for scrutiny of the process, as it provides an orthogonal approach to verifying the correctness of the computation: both checking that the given path has a certain strength, and that a set is indeed co-closed, is easy to verify. We reflect more on this in Section 4.6, and present an example of a full scrutiny sheet in the next section, when we join the type-level winning condition with the construction of the margin function from the given ballots.

4.3.1 Vote Counting as Inductive Type

Up to now, we have described the specification of Schulze voting relative to a given margin function. We now describe the specification (and computation) of the margin function given a profile (set) of ballots. Our formalisation describes an individual *count* as a type with the interpretation that all inhabitants of this type are correct executions of the vote counting algorithm. In the original paper describing the Schulze method [Schulze, 2011], a ballot is a linear preorder over the set of candidates.

In practice, ballots are implemented by asking voters to put numerical preferences against the names of candidates as represented by the image on the right. The most natural representation of a ballot is therefore a function $b : C \to Nat$ that assigns a natural number (the preference) for each candidate, and we recover a strict linear preorder $<_b$ on candidates by setting $c <_b d$ if $b(c) > b(d)$.

As preferences are usually numbered beginning with 1, we interpret a preference of 0 as the voter failing to designate a preference for a candidate as this allows us to also accommodate incomplete ballots. This is clearly a design decision, and we could have formalised ballots as functions $b : C \to 1 + Nat$ (with 1 being the unit type) but it would add little to our analysis.

```
Definition ballot := cand -> nat.
```

**Rank all candidates
in order of preference**

4 Lando Calrissian

3 Boba Fett

1 Mace Windu

2 Poe Dameron

2 Maz Kanata

Figure 4.4: Ballot Representation

The count of an individual election is then parameterised by the list of ballots cast, and is represented as a dependent inductive type. More precisely, we have a type State that represents either an intermediate stages of constructing the margin function or the determination of the final election result:

```
Inductive State: Type :=
| partial: (list ballot * list ballot)  ->
  (cand -> cand -> Z) -> State
| winners: (cand -> bool) ->  State.
```

The interpretation of this type is that a state either consists of two lists of ballots and a margin function, representing

- the set of ballots counted so far, and the set of invalid ballots seen so far

- the margin function constructed so far

or, to signify that winners have been determined, a boolean function that determines the set of winners.

The type that formalises correct counting of votes according to the Schulze method is parameterised by the profile of ballots cast (that we formalise as

a list), and depends on the type State. That is to say, an inhabitant of the type Count n, for n of type State, represents a correct execution of the voting protocol up to reaching state n. This state generally represents intermediate stages of the construction of the margin function, with the exception of the final step where the election winners are being determined. The inductive type takes the following shape:

```
Inductive Count (bs : list ballot) : State -> Type :=
| ax us m : us = bs -> (forall c d, m c d = 0) ->
  Count bs (partial (us, []) m)          (* zero margin      *)
| cvalid u us m nm inbs :
  Count bs (partial (u :: us, inbs) m) ->
  (forall c, (u c > 0)%nat) ->           (* u is valid       *)
  (forall c d : cand,
    ((u c < u d) -> nm c d = m c d + 1) (* c preferred to d *) /\
    ((u c = u d) -> nm c d = m c d)     (* c, d rank equal  *) /\
    ((u c > u d) -> nm c d = m c d - 1))(* d preferred to c *) ->
    Count bs (partial (us, inbs) nm)
| cinvalid u us m inbs :
  Count bs (partial (u :: us, inbs) m) ->
  (exists c, (u c = 0)%nat)   (* u is invalid      *) ->
  Count bs (partial (us, u :: inbs) m)
| fin m inbs w
  (d : (forall c, (wins_type m c) + (loses_type m c))) :
  Count bs (partial ([], inbs) m)  (*no ballots left*) ->
  (forall c, w c = true <-> (exists x, d c = inl x)) ->
  (forall c, w c = false <-> (exists x, d c = inr x)) ->
  Count bs (winners w).
```

The intuition here is simple: the first constructor, ax, initiates the construction of the margin function, and we ensure that all ballots are uncounted, no ballot are invalid (yet), and the margin function is constantly zero. The second constructor, cvalid, updates the margin function according to a valid ballot (all candidates have preferences marked against their name), and removes the ballot from the list of uncounted ballots. The constructor cinvalid moves an invalid ballot to the list of invalid ballots, and the last constructor fin applies only if the margin function is completely constructed (no more uncounted ballots). In its arguments, w : cand -> bool is the function that determines election winners, and d is a function that delivers, for every candidate, type-level evidence of winning or losing, consistent with w. Given this, we can

conclude the count and declare w to be the set of winners (or more precisely, those candidates for which w evaluates to `true`).

Together with the equivalence of the propositional notions of winning or losing a Schulze count with their type-level counterparts, every inhabitant of the type `Count b (winners w)` then represents a correct count of ballots b leading to the boolean predicate `w : cand -> bool` that determines the winners of the election with initial set b of ballots.

The crucial aspect of our formalisation of executions of Schulze counting is that the transcript of the count is represented by a type that is *not* a proposition. As a consequence, extraction delivers a program that produces the (set of) election winner(s), *together* with the evidence recorded in the type to enable independent verification.

4.3.2 All Schulze Election Have Winners

The main theorem, the proof of which we describe in this section, is that all elections according to the Schulze method engender a boolean-valued function `w : cand -> bool` that determines precisely which candidates are winners of the election, together with type-level evidence of this. Note that a Schulze election can have more than one winner, the simplest (but not the only) example being when no ballots at all have been cast. The theorem that we establish (and later extract as a program) simply states that for every incoming set of ballots, there is a boolean function that determines the election winners, together with an inhabitant of the type `Count` that witnesses the correctness of the execution of the count.

```
Theorem schulze_winners: forall (bs : list ballot),
  existsT (w: cand -> bool) (p : Count bs (winners w)), True.
```

The first step in the proof is elementary: We show that for any given list of ballots we can reach a state of the count where there are no more uncounted ballots, i.e. the margin function has been fully constructed.

```
Lemma all_ballots_counted: forall (bs : list ballot),
    existsT i m, (Count bs (partial ([], i) m)).
```

The second step relies on the iterated margin function already discussed in Section 4.3. As $M_n(c, d)$ (for n being the number of candidates) is the strength of the strongest path between c and d, we construct a boolean function w such that $w(c) = \texttt{true}$ if and only if $M_n(c, d) \geq M_n(d, c)$ for all $d \in C$. We then construct the type-level evidence required in the constructor `fin` using the function (or proposition) `iterated_marg_wins_type` described earlier.

4.4 Scrutiny Sheet and Experimental Results

The crucial aspect of our formalisation is that the vote counting protocol itself is represented as a dependent inductive type that represents all (correct) partial executions of the protocol. A complete execution is can then be understood as a state of vote counting where election winners have been determined. Our main theorem, `schulze_winners`, then asserts that an inhabitant of this type exists, for all possible sets of incoming ballots. Crucially, every such inhabitant contains enough information to (independently) verify the correctness of the election result, and can be thought of as a *certificate* for the count. From a computational perspective, we view tallying not merely as a function that delivers a result, but instead as a function that delivers a result, *together* with evidence that allows us to verify correctness. In other words, we augment verified correctness of an algorithm with the means to verify each particular *execution*.

From the perspective of electronic voting, this means that we no longer need to trust the hardware and software (assuming the cast-as-intended and collected-as-cast verifiability) that was employed to obtain the election result, as the generated certificate can be verified independently. In the literature on electronic voting, this is known as (tallied-as-cast) *verifiability* and has been recognised as one of the cornerstones for building trust in election outcomes by electronic voting research community [Chaum, 2004] [Küsters et al., 2011], [Benaloh and Tuinstra, 1994], [Delaune et al., 2010a], [Bernhard et al., 2017].

Coq's extraction mechanism then allows us to turn our main theorem, `schulze_winners`, into a provably correct program. When extracting, all purely propositional information is erased and given a set of incoming ballots, the ensuing program produces an inhabitant of the (extracted) type `Count` that records the construction of the margin function, together with (type level) evidence of correctness of the determination of winners. That is,

we see the individual steps of the construction of the margin function (one step per ballot) and once all ballots are exhausted, the determination of winners, together with paths and co-closed sets. The following is the transcript of a Schulze election where we have added wrappers to pretty-print the information content. This is the (full) scrutiny sheet promised in Section 4.3.

```
V: [A3 B1 C2 D4,..], I: [],
M: [AB:0 AC:0 AD:0 BC:0 BD:0 CD:0]
----------------------------------------------------------
V: [A1 B0 C4 D3,..], I: [],
M: [AB:-1 AC:-1 AD:1 BC:1 BD:1 CD:1]
----------------------------------------------------------
V: [A3 B1 C2 D4,..], I: [A1 B0 C4 D3],
M: [AB:-1 AC:-1 AD:1 BC:1 BD:1 CD:1]
----------------------------------------------------------
                           . . .
----------------------------------------------------------
V: [A1 B3 C2 D4], I: [A1 B0 C4 D3],
M: [AB:2 AC:2 AD:8 BC:5 BD:8 CD:8]
----------------------------------------------------------
V: [], I: [A1 B0 C4 D3],
M: [AB:3 AC:3 AD:9 BC:4 BD:9 CD:9]
----------------------------------------------------------
winning: A
  for B: path A --> B of strenght 3, 4-coclosed set:
    [(B,A),(C,A),(C,B),(D,A),(D,B),(D,C)]
  for C: path A --> C of strenght 3, 4-coclosed set:
    [(B,A),(C,A),(C,B),(D,A),(D,B),(D,C)]
  for D: path A --> D of strenght 9, 10-coclosed set:
    [(D,A),(D,B),(D,C)]
losing: B
  exists A: path A --> B of strength 3, 3-coclosed set:
    [(A,A),(B,A),(B,B),(C,A),(C,B),(C,C),
     (D,A),(D,B),(D,C),(D,D)]
losing: C
  exists A: path A --> C of strength 3, 3-coclosed set:
    [(A,A),(B,A),(B,B),(C,A),(C,B),(C,C),
     (D,A),(D,B),(D,C),(D,D)]
losing: D
  exists A: path A --> D of strength 9, 9-coclosed set:
    [(A,A),(A,B),(A,C),(B,A),(B,B),(B,C),
```

```
(C,A),(C,B),(C,C),(D,A),(D,B),(D,C),(D,D)]
```

Here, we assume four candidates, A, B, C and D and a ballot of the form A3 B2 C4 D1 signifies that D is the most preferred candidate (the first preference), followed by B (second preference), A and C. In every line, we only display the first uncounted ballot (condensing the remainder of the ballots to an ellipsis), followed by votes that we have deemed to be invalid. We display the partially constructed margin function on the right. Note that the margin function satisfies $m(x,y) = -m(y,x)$ and $m(x,x) = 0$ so that the margins displayed allow us to reconstruct the entire margin function. In the construction of the margin function, we begin with the constant zero function, and going from one line to the next, the new margin function arises by updating according to the first ballot. This corresponds to the constructor cvalid and cinvalid being applied recursively: we see an invalid ballot being set aside in the step from the second to the third line, all other ballots are valid. Once the margin function is fully constructed (there are no more uncounted ballots), we display the evidence provided in the constructor fin: we present evidence of winning (losing) for all winning (losing) candidates. In order to actually verify the computed result, a third party observer would have to

1. Check the correctness of the individual steps of computing the margin function

2. For winners, verify that the claimed paths exist with the claimed strength, and check that the claimed sets are indeed coclosed.

Contrary to re-running a different implementation on the same ballots, our scrutiny sheet provides an *orthogonal* perspective on the data and how it was used to determine the election result.

We have evaluated our approach by extracting the entire Coq development into Haskell, with all types defined by Coq extracted as is, i.e. in particular using Coq's unary representation of natural numbers, and Haskell native integer representation. The results are displayed in Figure 4.5, and Figure 4.6 using a logarithmic scale.

Figure 4.5: Experimental Result (Coq Unary Natural Number, Slow)

4.5 Counting Millions of Ballots

The previous extracted Haskell code was very slow and was not practical for real life election involving millions of ballots. To scale it to real life election, we investigated the extracted Haskell code from Coq code. The most performance critical aspect of our code was the computation of margin function. Recall that the margin function is of type `cand -> cand -> Z` and that it depends on the *entire* set of ballots. Internally, it is represented by a closure [Landin, 1964] so that margins are re-computed with every call. The single largest efficiency improvement in our code was achieved by memoization, i.e. representing the margin function (in Coq) via list lookup. With this (and several smaller) optimisation, we can count millions of votes using verified code. However, this efficiency did not come for free, and we had to pay the cost in terms of (almost all) broken proofs. We had to redo all the proofs all over

Figure 4.6: Experimental Result (Haskell Native Integer, Slow)

again[2]. Below (Figure 4.7, Figure 4.8), we include our timing graphs, based on randomly generated ballots while keeping number of candidates constant i.e. 4 (The reason we kept it to 4 candidate to show the speed up compared to 4.5 and 4.6 . During the experimentation, we ran an election with 21 candidate, and we were able to count 2 million randomly generated ballots before running out of memory.)

In the Figure 4.7, we report timings (in seconds) for the computation of winners, whereas in the Figure 4.8, we include the time to additionally compute a universally verifiable certificate that attests to the correctness of the count. This is consistent with complexity of Schulze counting i.e. linear in no of ballots and cubic in candidates. The experiments were carried out on system equipped with intel core i7 processor and 16 GB of ram. We notice that the computation of the certicate adds comparatively little in computational cost.

---

[2]Redoing these proofs were trivial, but time consuming. I wished if there was a tool to automate this process

Figure 4.7: Computation of Winner (Without Certificate, Fast)

Our implementation requires that we store *all* ballots in main memory as we need to parse the entire list of ballots before making it available to our verified implementation so that the total number of ballots we can count is limited by main memory in practise. We can count real-world size elections (8 million ballot papers) on a standard, commodity desktop computer with 16 GB of main memory.

4.6 Discussion

In this chapter, we emphasize on correctness, and we take the approach that computation of winners in electronic voting (and in situations where correctness is key in general) should not only produce an end result, but an end result, *together* with a verifiable justification of the correctness of the computed result. We have exemplified this approach by providing a provably

Figure 4.8: Computation of Winner (With Certificate, Fast)

correct, and evidence-producing implementation of vote counting according to the Schulze method.

While the Schulze method is not difficult to implement, and indeed there are many freely available implementations on the Internet, comparing the results between different implementations can give some level of assurance for correctness only in case the results agree. If there is a discrepancy, a certificate for the correctness of the count allows to adjudicate between different implementations, as the certificate can be checked with relatively little computational effort.

From the perspective of computational complexity, checking a transcript for correctness is of the same complexity as computing the set of winners, as our certificates are cubic in size, so that certificate checking is not less complex than the actual computation. However, publishing an independently verifiable certificate that attests the individual steps of the computation helps to increase *trust* in the computed election outcome. Typically, the use of

technology in elections increases the amount of trust that we need to place both in technological artefacts, and in people. It raises questions that range from fundamental aspects, such as proper testing and/or verification of the software, to very practical questions, e.g. whether the correct version of the software has been run. On the contrast, publishing a certificate of the count dramatically reduces the amount of trust that we need to place into both people and technology: the ability to publish a verifiable justification of the correctness of the count allows a large number of individuals to scrutinise the count. While only moderate programming skills are required to check the validity of a certificate (the transcript of the count), even individuals without any programming background can at least spot-check the transcript: for the construction of the margin function, everything that is needed is to show that the respective margins change according to the counted ballot. For the correctness of determination of winners, it is easy to verify existence of paths of a given strength, and also whether certain sets are co-closed – even by hand! This dramatically increases the class of people that can scrutinise the correctness of the count, and so helps to establish a trust basis that is much wider as no trust in election officials and software artefacts is required.

Technically, we do not *implement* an algorithm that counts votes according to the Schulze method. Instead, we give a specification of the Schulze winning conditions (`wins_prop` in Section 4.3) in terms of an already computed margin function that (we hope) can immediately be seen to be correct, and then show that those winning conditions are equivalent to the existence of inhabitants of types that carry verifiable evidence (`wins_type`). We then join the (type level) winning conditions with an inductive type that details the construction of the margin function in an inductive type. Via propositions-as-types, a provably correct vote counting function is then equivalent the proposition that there exists an inhabitant of `Count` for every set of ballots. Coq's extraction mechanism then allows us to extract a Haskell program that produces election winners, together with verifiable certificates.

4.7   Summary

Our formalization achieves *Correctness*, *Practicality*, and (tallied-as-cast) *Verifiability*. The major problem in this formalization is *Privacy*. Our ballots are in plaintext and could easily be identified if the number of candidates participating in election in are large (Italian attack) [Otten, 2003]. In nutshell, the achieved and missed of this formalization:

- Achieved

    - Correctness: The implementation is formalized in Coq with emphasis on generating evidence to convince everyone about the outcome of election.

    - Practicality: The extracted code can count millions of ballots. Therefore, we can use it in any real life election.

    - Verifiability: The outcome of any election can be verified by any third party using the generated certificates. Certificates generated for plaintext ballot during the election are very simple. It requires basic math literacy to audit the certificate which would lead to increase in number of scrutineers.

- Missed

    - Privacy: There is no privacy because the ballots involved are simply plaintext, and potentially, it could lead coercion and vote-selling (coercion).

We remark that extracting Coq developments into a programming language itself is a non-verified process which could still introduce errors in our code. The most promising way to alleviate this is to independently implement (and verify) a certificate verifier, possibly in a language such as CakeML [Kumar et al., 2014] that is guaranteed to be correct to the machine level.

In the next chapter, we will try to solve privacy and coercion problem, using encryption, and to keep it verifiable, we will use zero-knowledge-proof. However, solution for privacy comes at a cost, e.g. a loss in the pool of scrutineers because auditing a certificate generated by counting encrypted ballot would requires intricate knowledge of cryptography.

Homomorphic Schulze Algorithm : Axiomatic Approach

Be melting snow. Wash yourself of yourself.

Rumi

5.1 Introduction

As we stated in the summary of last chapter that plaintext could lead to privacy problems, e.g. ballot identification (italian attack) Otten [2003]. In this chapter, we will try to achieve: privacy by using encryption, (tallied-as-cast) verifiability by using zero knowledge proof, and correctness of implementation by proving the correctness properties inside Coq theorem prover. One important point to note that we did not formalize any cryptographic primitive inside the Coq, but took an axiomatic approach, i.e. we assumed the existence of cryptographic primitive and correctness property about their behaviour (axiomatisation of cryptographic primitives). The reason we took the axiomatic approach because our goal was not to formalize cryptographic primitives, but use these primitives to conduct an election which has all three ingredient privacy, verifiability, and correctness. We then obtain, via program extraction, a provably correct implementation of vote counting, that we turn into executable code by providing implementations of the primitives based on a standard cryptographic library (Unicrypt). We conclude by presenting experimental results, and discuss trust the trust base, security and privacy as well as the applicability of our work to real-world scenarios.

Chapter Outline: In the section 5.2, we discuss the technique to achieve verifiable tally. In order to do so, we discuss why do we need our ballot to have a matrix representation and not a ranking function. Moreover, we discuss the concept of validity ballot, which comes naturally with matrix representation, steps of homomorphic counting, and cryptographic primitives needed to achieve all the required functionality. Section 5.3 takes a step forward and makes every concept from the previous section concrete using the Coq theorem prover. One important point in this section is our inductive data type *ECount* augmented with verification data in form of zero-knowledge-proof for various claims made during the counting (*ECount* is conceptually similar as *Count* (section 4.3.1), but in terms of data, *ECount* has state data related to counting and verification data in form of zero-knowledge proof related to claims made by us, while *Count* has just state data). In the section 5.4, we have demonstrated our main theorem which states that for any set of given encrypted ballots, winner can always be found. Apart from the main theorem, this section also incorporates the proof of correctness by stating the winners produced by encrypted ballots are same as the plain-text ballots (section 4.3.2) if encrypted ballots are aligned (decrypt) with plain-text ballot. Section 5.5 focuses on extraction, instantiating the cryptographic primitives with UniCrypt library, and experimental results. Finally, section 5.6 highlights our assumptions, scalability issues, the goals we achieved and the goals we missed.

Secure elections are a balancing act between integrity and privacy: achieving either is trivial but their combination is notoriously hard. One of the key challenges faced by both paper based and electronic elections is that results must substantiated with verifiable evidence of their correctness while retaining the secrecy of the individual ballot [Bernhard et al., 2017]. The combination of privacy and integrity can be realised using cryptographic techniques, where encrypted ballots (that the voters themselves cannot decrypt) are published on a bulletin board, and the votes are then processed, and the correctness of the final tally is substantiated, using homomorphic encryption [Hirt and Sako, 2000] and verifiable shuffling [Bayer and Groth, 2012]. (Separate techniques exist to prevent ballot box stuffing and to guarantee cast-as-intended.) Integrity can then be guaranteed by means of Zero Knowledge Proofs (ZKP), first studied by Goldwasser, Micali, and Rackoff [Goldwasser et al., 1985]. Informally, a ZKP is a probabilistic and interactive proof where one entity interacts with another such that the interaction provides no information other than that the statement being proved is true with overwhelming probability. Later results [Ben-Or et al., 1988; Goldreich et al., 1991] showed that all problems for which solutions can be efficiently verified have zero

knowledge proof (In practice, sigma protocol is used prove the knowledge of witness w for a public input x, and it is required to be zero-knowledge against the honest verifier).

5.2 Verifiable Homomorphic Tallying

The realisation of verifiable homomorphic tallying that we are about to describe follows the same two phases as the algorithm (Chapter 4.2): We first homomorphically compute the margin matrix from encrypted ballots, and then compute winners on the basis of the (decrypted) margin. Moreover, the computation also produces a verifiable certificate that leaks no information about choices in individual ballots other than the (final) margin matrix, which in turn leaks no information about individual ballots if the number of voters is large enough.

Format of Ballots. Recall that in preferential voting schemes, ballots are rank-ordered lists of candidates. For the Schulze Method, we require that all candidates are ranked, and two candidates may be given the same rank. That is, a ballot is most naturally represented as a function $b : C \to Nat$ that assigns a numerical rank to each candidate, and the computation of the margin amounts to computing the sum

$$m(x, y) = \sum_{b \in B} \begin{cases} +1 & b(x) > b(y) \\ 0 & b(x) = b(y) \\ -1 & b(x) < b(y) \end{cases}$$

where B is the multi-set of ballots, and each $b \in B$ is a ranking function $b : C \to Nat$ over a (finite) set C of candidates.

Ideally, we could have copied the same ballot structure in homomorphic Schulze method, but encrypting the choices, i.e. the ballot would have been represented as $b : C \to EncryptedNaturalNumber$ where $EncryptedNaturalNumber$ is the encrypted representation of a choice (natural number). However, we note that this representation of ballots is not well suited for homomorphic computation of the margin matrix as practically feasible homomorphic encryption schemes do not support comparison operators and case distinctions as used in the formula above (to the best of our knowledge).

We instead represent ballots as matrices $b(x, y)$ where $b(x, y) = +1$ if x

is preferred over y, $b(x, y) = -1$ if y is preferred over x and $b(x, y) = 0$ if x and y are equally preferred. The downside of this representation is that it takes $O(n^2)$ space to represent a ballot where n is the number of candidate participating in election.

While the advantage of the first representation is that each ranking function is necessarily a valid ranking and linear space ($O(n)$)) in number of candidates, n, the advantage of the matrix representation is that the computation of the margin matrix is simple, that is

$$m(c, d) = \sum_{b \in B} b(x, y)$$

where B is the multi-set of ballots (in matrix form), and can moreover be transferred to the encrypted setting in a straight forward way: if ballots are matrices $e(x, y)$ where $e(x, y)$ is the encryption of an integer in $\{-1, 0, 1\}$, then

$$encm = \bigoplus_{encb \in EncB} encb(x, y) \tag{5.1}$$

where $\oplus$ denotes homomorphic addition, $encb$ is an encrypted ballot in matrix form (i.e. decrypting $encb(x, y)$ indicates whether x is preferred over y), and $EncB$ is the multi-set of encrypted ballots. The disadvantage is that we need to verify that a matrix ballot is indeed valid, that is

- that the decryption of $encb(x, y)$ is indeed one of $1, 0$ or -1

- that $encb$ indeed corresponds to a ranking function.

Indeed, to achieve verifiability, we not only need *verify* that a ballot is valid, we also need to *evidence* its validity (or otherwise) in the certificate.

Validity of Ballots. By a plaintext (matrix) ballot we simply mean a function $b : C \times C \rightarrow Z$, where C is the (finite) set of candidates. A plaintext ballot $b(x, y)$ is *valid* if it is induced by a ranking function, i.e. there exists a function $f : C \rightarrow Nat$ such that $b(x, y) = 1$ if $f(x) < f(y)$, $b(x, y) = 0$ if $f(x) = f(y)$ and $b(x, y) = -1$ if $f(x) > f(y)$. A *ciphertext (matrix) ballot* is a function $encb : C \times C \rightarrow CT$ (where CT is a chosen set of ciphertexts), and it is valid if its decryption, i.e. the plaintext ballot $b(x, y) = dec(encb(x, y))$ is valid (where dec denotes decryption).

For a plaintext ballot, it is easy to decide whether it is valid (and should be counted) or not (and should be discarded). We use shuffles (ballot per-

mutations) to evidence the validity of encrypted ballots. One observes that a matrix ballot is valid if and only if it is valid after permuting both rows and columns with the same permutation. That is, $b(x, y)$ is valid if and only if $b'(x, y)$ is valid, where

$$b'(x, y) = b(\pi(x), \pi(y))$$

and $\pi : C \to C$ is a permutation of candidates. (Indeed, if f is a ranking function for b, then $f \circ \pi$ is a ranking function for b'). As a consequence, we can evidence the validity of a ciphertext ballot *encb* by

- publishing a shuffled version *encb'* of *encb*, that is shuffled by a secret permutation, together with evidence that *encb'* is indeed a shuffle of *encb*

- publishing the decryption b' of *encb'* together with evidence that b' is indeed the decryption of *encb'*.

We use zero-knowledge proofs in the style of [Terelius and Wikström, 2010] to evidence the correctness of the shuffle, and zero-knowledge proofs of honest decryption [Chaum and Pedersen, 1992] to evidence correctness of decryption. This achieves ballot secrecy as the (secret) permutation is never revealed.

In summary, the evidence of correct (homomorphic) counting starts with an encryption of the zero margin *encm*, and for each ciphertext ballot *encb* contains

1. a shuffle of *encb* together with a ZKP of correctness

2. decryption of the shuffle, together with a ZKP of correctness

3. the updated margin matrix, if the decrypted ballot was valid, and

4. the unchanged margin matrix, if the decrypted ballot is not valid.

Once all ballots have been processed in this way, the certificate determines winners and contains winners by

5. the fully constructed margin, together with its decryption and ZKP of honest decryption after counting all the ballots

6. publishes the winner(s), together with evidence to substantiate the claim

Cryptographic primitives. We require an additively homomorphic cryptosystem to compute the (encrypted) margin matrix according to Equation 5.1 (this implements Item 3 above). All other primitives fall into one of three categories. *Verification primitives* are used to syntactically define the type of valid certificates. For example, when publishing the decrypted margin matrix in Item 5 above, we require that the zero knowledge proof in fact evidences correct decryption. To guarantee this, we need a verification primitive that – given ciphertext, plaintext and zero knowledge proof – verifies whether the supplied proof indeed evidences that the given ciphertext corresponds to the given plaintext. In particular, verification primitives are always boolean valued functions. While verification primitives *define* valid certificates, *generation primitives* are used to *produce* valid certificates. In the example above, we need a decryption primitive (to decrypt the homomorphically computed margin) and a primitive to generate a zero knowledge proof (that witnesses correct decryption). Clearly verification and generation primitives have a close correlation, and we need to require, for example, that zero knowledge proofs obtained via a generation primitive has to pass muster using the corresponding verification primitive.

The three primitives described above (decryption, generation of a zero knowledge proof, and verification of this proof) already allow us to implement the entire protocol with exception of ballot shuffling (Item 1 above). Here, the situation is more complex. While existing mixing schemes (e.g. [Bayer and Groth, 2012]) permute an array of ciphertexts and produce a zero knowledge proof that evidences the correctness of the shuffle, our requirement dictates that every row and colum of the (matrix) ballot is shuffled with the *same* (secret) permutation. In other words, we need to retain the identity of the permutation to guarantee that each row and column of a ballot have been shuffled by the same permutation. We achieve this by committing to a permutation using Pedersen's commitment scheme [Pedersen, 1992]. In a nutshell, the Pedersen commitment scheme has the following properties.

- Hiding: the commitment reveals no information about the permutation

- Binding: no party can open the commitment in more than one way, i.e. the commitment is to one permutation only.

A combination of Pedersen's commitment scheme with a zero knowledge proof leads to a similar two step protocol, also known as commitment-consistent

proof of shuffle [Wikström, 2009].

- Commit to a secret permutation and publish the commitment (hiding).

- Use a zero knowledge proof to show that shuffling has used the same permutation which we committed to in previous step (binding).

This allows us to witness the validity (or otherwise) of a ballot by generating a permutation π which is used to shuffle every row and column of the ballot. We hide π by committing it using Pedersen's commitment scheme and record the commitment c_π in the certificate. However, for the binding step, rather than opening π we generate a zero knowledge proof, zkp_π, using π and c_π, which can be used to prove that c_π is indeed the commitment to some permutation used in the (commitment consistent) shuffling without being opened [Wikström, 2009]. We can now use the permutation that we have committed to for shuffling each row and column of a ballot, and evidence the correctness of the shuffle via a zero knowledge proof. To evidence validity (or otherwise) of a (single) ballot, we therefore:

1. generate a (secret) permutation and publish a commitment to this permutation, together with a zero knowledge proof that evidences commitment to a permutation

2. for each row of the ballot, publish a shuffle of the row with the permutation committed to, together with a zero knowledge proof that witnesses shuffle correctness

3. for each column of the row shuffled ballot, publish a shuffle of the column, also together with a zero knowledge proof of correctness

4. publish the decryption the ballot shuffled in this way, together with a zero knowledge proof that witnesses honest decryption

5. decide the validity of the ballot based on the decrypted shuffle.

The cryptographic primitives needed to implement this again fall into the same classes. To define validity of certificates, we need verification primitives

- to decide whether a zero knowledge proof evidences that a given commitment indeed commits to a permutation

- to decide whether a zero knowledge proof evidences the correctness of a shuffle relative to a given permutation commitment.

Dual to the above, to generate (valid) certificates, we need the ability to

- generate permutation commitments and accompanying zero knowledge proofs that evidence commitment to a permutation

- generate shuffles relative to a commitment, and zero knowledge proofs that evidence the correctness of shuffles.

Again, both need to be coherent in the sense that the zero knowledge proofs produced by the generation primitives need to pass validation. In summary, we require an additively homomorphic cryptosystem that implements the following:

Decryption Primitives. decryption of a ciphertext, creation and verification of honest decryption zero knowledge proofs.

Commitment Primitives. generating permutations, creation and verification of commitment zero knowledge proofs

Shuffling Primitives. commitment consistent shuffling, creation and verification of commitment consistent zero knowledge shuffle proofs

Witnessing of Winners. Once all ballots are counted, the computed margin is decrypted, and winners (together with evidence of winning) are computed using plaintext counting. We discuss this part only briefly, for sake of completeness, as it is identical to the previous (chapter 4). For each of the winners w and each candidate x we publish

- a natural number $k(w, x)$ and a path $w = x_0, \ldots, x_n = x$ of strength k

- a set $C(w, x)$ of pairs of candidates that is k-coclosed and contains (x, w)

where a set S is k-coclosed if for all $(x, z) \in C$ we have that $m(x, z) < k$ and either $m(x, y) < k$ or $(y, z) \in S$ for all candidates y. Informally, the first requirement ensures that there is no direct path (of length one) between a pair $(x, z) \in S$, and the second requirement ensures that for an element $(x, z) \in S$, there cannot be a path that connects x to an intermediate node y and then (transitively) to z that is of strength $\geq k$.

5.3   Formalization in a Theorem Prover

As we stated in the beginning of this chapter that the purpose of this work
was not to verify cryptographic primitives, but use them as a tool to construct
evidence which can be used to audit and verify the outcome during different
phase of election. Here, we treat them as abstract entities and assume axioms
about them inside Coq. In particular, we assume the existence of functions
that implement each of the primitives described in the previous section, and
postulate natural axioms that describe how the different primitives interact.
As a by-product, we obtain an axiomatisation of a cryptographic library that
we could, in a later step, verify the implementation of a cryptosystem against.
In particular, this allows us to not commit to any particular cryptosystem
in particular (although our development, and later instantiation, is geared
towards El Gamal [Gamal, 1984]).

The first part of our formalisation concerns the cryptographic primitives
that we collect in a separate module. Below is an example of the generation
/ verification primitives for decryption, together with coherence axioms.

```
Variable decrypt_message:
  Group -> Prikey -> ciphertext ->  plaintext.

Variable construct_zero_knowledge_decryption_proof:
  Group -> Prikey -> ciphertext -> DecZkp.

Axiom verify_zero_knowledge_decryption_proof:
  Group -> plaintext -> ciphertext -> DecZkp -> bool.

Axiom honest_decryption_from_zkp_proof: forall group c d zkp,
 verify_zero_knowledge_decryption_proof group d c zkp = true
 -> d = decrypt_message grp privatekey c.

Axiom verify_honest_decryption_zkp (group: Group):
  forall (pt : plaintext) (ct : ciphertext) (pk : Prikey),
  (pt = decrypt_message group pk ct) ->
  verify_zero_knowledge_decryption_proof group pt ct
  (construct_zero_knowledge_decryption_proof group pk ct)
  = true.
```

The difference between the keyword Variable and Axiom is purely syntactic,

and in our case, used as a convenience for extraction (The keyword `Variable` is used if want that variable of function to be lambda abstracted otherwise keyword `Axiom`). In the above, the first two functions, `decrypt_message` and `construct_zero _knowledge _decryption _proof` are *generation* primitives, whereas the function `verify_zero _knowledge _decryption _proof` is a *verification* primitive. We have two coherence axioms. The first says that if the verification of a zero knowledge proof of honest decryption succeeds, then the ciphertext indeed decrypts to the given plaintext. The second stipulates that generated zero knowledge proofs indeed verify.

For ballots, we assume a type `cand` of candidates, and represent plaintext and encrypted ballots as two-argument functions that take plaintext, and ciphertexts, as values.

```
Definition pballot := cand -> cand -> plaintext.
Definition eballot := cand -> cand -> ciphertext.
```

We now turn to the representation of certificates, and indeed to the definition of what it means to (a) count encrypted votes correctly according to the Schulze Method, and (b) produce a verifiable certificate of this fact. At a high level, we split the counting (and accordingly the certificate) into *states*. This gives rise to a (inductive dependent) type `ECount`, parameterised by the ballots being counted.

```
Inductive ECount (group : Group) (bs : list eballot) :
  EState -> Type
```

Given a list `bs` of ballots, `ECount bs` is a inductive dependent type. In this case, given a state of counting (i.e. an inhabitant `estate` of `EState`), the type level application `ECount bs estate` is the *type of evidence that proves that estate is a state of counting that has been reached according to the method*. The states itself are represented by the type `EState` where

- `epartial` represents a partial state of counting, consisting of the homomorphically computed margin so far, the list of uncounted ballots and the list of invalid ballots encountered so far

- `edecrypt` represents the final decrypted margin matrix, and

- `ewinners` is the final determination of winners.

This is readily translated to the following Coq code:

```
Inductive EState : Type :=
 | epartial : (list eballot * list eballot) ->
              (cand -> cand -> ciphertext) -> EState
 | edecrypt : (cand -> cand -> plaintext) -> EState
 | ewinners : (cand -> bool) -> EState.
```

The constructors of EState then allow us to move from one state to the next, under appropriate conditions that guarantee correctness of the count. The different states during the counting represented by *ECount* is tagged by five constructors:

- ecax: marks the beginning of counting

- ecvalid: take a ballot from cast-ballots pile, and the ballot is a valid ballot

- ecinvalid: take a ballot from cast-ballot pile, and the ballot is a invalid ballot

- ecdecrypt: decryption of fully constructed homomorphic margin from the cast-ballot

- ecfin: declaration of winner and loser based on the decrypted margin

Inductive type *ECount* with all the constructors filled with *state data*, *verification data*, and *correctness constraint*. The first constructor, ecax, bootstraps the count and ensures that

- all ballots are initially uncounted

- margin matrix is an encryption of the zero matrix

state data: here, the list of uncounted and invalid ballots, and the encrypted homomorphic margin

verification data: a zero knowledge proof that the encrypted homomorphic margin is indeed an encryption of the zero margin

correctness constraints: here, the constructor may only be applied if the list of uncounted ballots is equal to the list of ballots cast, and the fact that the zero knowledge proofs indeed verify that the initial margin matrix is identically zero.

The main difference between the correctness condition, and the verification data is that the former can be simply be inspected (here by comparing lists) whereas the latter requires additional data (here in the form of a zero knowledge proof).

```
Inductive ECount (grp : Group) (bs : list eballot) :
 EState -> Type :=
 | ecax (us : list eballot)
     (encm : cand -> cand -> ciphertext)
     (decm : cand -> cand -> plaintext)
     (zkpdec : cand -> cand -> DecZkp) :
     us = bs ->
     (forall c d : cand, decm c d = 0) ->
     (forall c d, verify_zero_knowledge_decryption_proof
         grp (decm c d) (encm c d) (zkpdec c d) = true) ->
   ECount grp bs (epartial (us, []) encm)
```

The constructor `ecvalid` represents the effect of counting a valid ballot. Here the crucial aspect is that validity needs to be evidenced. As before, we have:

state data: as before, the list of uncounted and invalid ballots, the homomorphic margin, but additionally evidence that the previous state has been obtained correctly

verification data: a commitment to a (secret) permutation, a row permutation of the ballot being counted, and a column permutation of this, and a decryption of the row- and column permuted ballot (all with accompanying zero knowledge proofs)

correctness constraints: all the zero knowledge proofs verify, the new margin is the homomorphic addition of the previous margin and the counted ballot, and the decrypted (shuffled) ballot is indeed valid.

```
| ecvalid (u : eballot) (v : eballot) (w : eballot)
  (b : pballot) (zkppermuv : cand -> ShuffleZkp)
  (zkppermvw : cand -> ShuffleZkp)
  (zkpdecw : cand -> cand -> DecZkp)
  (cpi : Commitment) (zkpcpi : PermZkp)
  (us : list eballot)
  (m nm : cand -> cand -> ciphertext)
  (inbs : list eballot) :
  ECount grp bs (epartial (u :: us, inbs) m) ->
  (* valid ballot *)
  matrix_ballot_valid b ->
  (* commitment proof *)
  verify_permutation_commitment grp
      (List.length cand_all) cpi zkpcpi = true ->
  (forall c, verify_row_permutation_ballot grp
      u v cpi zkppermuv c = true)   ->
  (forall c, verify_col_permutation_ballot grp
      v w cpi zkppermvw c = true)   ->
  (forall c d, verify_zero_knowledge_decryption_proof
      grp (b c d) (w c d) (zkpdecw c d) = true)   ->
  (forall c d, nm c d = homomorphic_addition
      grp (u c d) (m c d)) ->
  ECount grp bs (epartial (us, inbs) nm)
```

The constructor `ecinvalid` is very similar to `ecvalid`. We elide the description of the constructor that is applied when an invalid ballot is being encountered (the only difference is that the margin matrix is not being updated and ballot is moved to list of invalid ballots).

```
| ecinvalid (u : eballot) (v : eballot) (w : eballot)
  (b : pballot) (zkppermuv : cand -> ShuffleZkp)
  (zkppermvw : cand -> ShuffleZkp)
  (zkpdecw : cand -> cand -> DecZkp)
  (cpi : Commitment) (zkpcpi : PermZkp)
  (us : list eballot) (m : cand -> cand -> ciphertext)
  (inbs : list eballot) :
  ECount grp bs (epartial (u :: us, inbs) m) ->
  (* invalid ballot *)
  ~matrix_ballot_valid b ->
  (* commitment proof *)
```

```
    verify_permutation_commitment grp
       (List.length cand_all) cpi zkpcpi = true ->
    (forall c, verify_row_permutation_ballot grp
       u v cpi zkppermuv c = true) ->
    (forall c, verify_col_permutation_ballot grp
       v w cpi zkppermvw c = true) ->
    (forall c d, verify_zero_knowledge_decryption_proof
       grp (b c d) (w c d) (zkpdecw c d) = true) ->
    ECount grp bs (epartial (us, (u :: inbs)) m)
```

Counting finishes when there are no more uncounted ballots and this
state is marked by constructor *ecdecrypt*, in which case the next step is to
publish the decrypted margin matrix. Also here, we have

state data: the decrypted margin matrix, plus evidence that a state with no
more uncounted ballots has been obtained correctly

verification data: a zero knowledge proof that demonstrates honest decryp-
tion of the final margin matrix

correctness constraints: the given zero knowledge proof verifies, i.e. the
given decrypted margin is indeed the decryption of the (last) homo-
morphically computed margin matrix.

```
  | ecdecrypt inbs
      (encm : cand -> cand -> ciphertext)
      (decm : cand -> cand -> plaintext)
      (zkp : cand -> cand -> DecZkp) :
      ECount grp bs (epartial ([], inbs) encm) ->
      (forall c d, verify_zero_knowledge_decryption_proof
         grp (decm c d) (encm c d) (zkp c d) = true) ->
      ECount grp bs (edecrypt decm)
```

The last constructor, *ecfin*, finally declares the winners of the election, and
we have:

state data: a function cand -> bool that determines winners, plus evidence
of the fact that the decrypted final margin matrix has been obtained
correctly

verification data: paths and co-closed sets that evidence the correctness of the function above

correctness constraints: that ensure that the verification data verifies the winners given by the state data.

This last part is same as the previous chapter's scrutiny sheet (section 4.4).

```
| ecfin dm w
    (d : (forall c, (wins_type dm c) +
                    (loses_type dm c))) :
    ECount grp bs (edecrypt dm) ->
    (forall c, w c = true <-> (exists x, d c = inl x)) ->
    (forall c, w c = false <-> (exists x, d c = inr x)) ->
    ECount grp bs (ewinners w).
```

5.4 Correctness by Construction and Verification

In the previous section, we have presented a data type that *defines* the notion of a verifiably correct count of the Schulze Method, on the basis of encrypted ballots. To obtain an executable that in fact *produces* a verifiable (and provably correct) count, we can proceed in either of two ways:

1. implement a function that – give a list `bs` of ballots – produces a boolean function `w` (for winners) and an element of the type `ECount bs (winners w)`. This gives both the election winners (`w`) as well as evidence (the element of the `ECount` data type).

2. to prove that for every set `bs` of encrypted ballots, we have a boolean function `w` and an inhabitant of the type `ECount bs (winners w)`.

Under the proofs-as-programs interpretation of constructive type theory, both amount to the same. We chose the latter approach, and our main theorem formally states that all elections can be counted according to the Schulze Method (with encrypted ballots), i.e. a winner can always be found. Formally, our main theorem takes the following form:

```
Lemma encryption_schulze_winners (group : Group)
 (bs : list eballot) : existsT (f : cand -> bool),
 ECount group bs (ewinners f).
```

The proof proceeds by successively building an inhabitant of EState by homomorphically computing the margin matrix, then decrypting and determining the winners. Within the proof, we use both generation primitives (e.g. to construct zero knowledge proofs) and coherence axioms (to ensure that the zero knowledge proofs indeed verify).

The correctness of our entire approach stands or falls with the correct formalisation of the inductive data type ECount that is used to determine the winners of an election counted according to the Schulze Method. While one can argue that the data type itself is transparent enough to be its own specification, the cryptographic aspect makes things slightly more complex. For example, it appears to be credible that our mechanism for determining validity of a ballot is correct – however we have not given proof of this. Rather than scrutinising the details of the construction of this data type, we follow a different approach: we demonstrate that homomorphic counting always yields the same results as plaintext counting, where plaintext counting is already verified against its specification (chapter 4). This correspondence has two directions, and both assume that we are given two lists of ballots that are the encryption (resp. decryption) of one another.

The first theorem, plaintext_ schulze_to _homomorphic, reproduced below shows that every winner that can be determined using plaintext counting can also be evidenced on the basis of corresponding encrypted ballots. The converse of this is established by Theorem homomorphic _schulze _to_plaintext.

```
Lemma plaintext_schulze_to_homomorphic
 (group : Group) (bs : list ballot):
 forall (pbs : list pballot) (ebs : list eballot)
 (w : cand -> bool), (pbs = map (fun x => (fun c d =>
 decrypt_message group privatekey (x c d))) ebs) ->
 (mapping_ballot_pballot bs pbs) ->
 Count bs (winners w) -> ECount group ebs (ewinners w).

Lemma homomorphic_schulze_to_plaintext
 (group : Group) (bs : list ballot):
 forall (pbs : list pballot) (ebs : list eballot)
 (w : cand -> bool) (pbs = map (fun x => (fun c d =>
```

```
    decrypt_message group privatekey (x c d))) ebs) ->
    (mapping_ballot_pballot bs pbs) ->
    ECount grp ebs (ewinners w) -> Count bs (winners w).
```

The theorems above feature a third type of ballot that is the basis of plaintext counting, and is a simple ranking function of type `cand -> Nat`, and the two hypotheses on the three types of ballots ensure that the encrypted ballots (`ebs`) are in fact in alignment with the rank-ordered ballots (`bs`) that are used in plaintext counting. The proof, and indeed the formulation, relies on an inductive data type `Count` (section 4.3.1) that can best be thought of as a plaintext version of the inductive type `ECount` given here. Crucially, `Count` is verified against a formal specification of the Schulze Method. Both theorems are proven by induction on the definition of the respective data types, where the key step is to show that the (decrypted) final margins agree. The key ingredient here are the coherence axioms that stipulate that zero knowledge proofs that verify indeed evidence shuffle and/or honest decryption.

5.5   Extraction and Experiments

As already mentioned, we are using the Coq extraction mechanism[Letouzey, 2003] to extract programs from existence proofs[1]. In particular, we extract the proof of the Theorem `pschulze_winners`, given in Section 5.4 to a program that delivers not only provably correct counts, but also verifiable evidence. Give a set of encrypted ballots and a `Group` that forms the basis of cryptographic operations, we obtain a program that delivers not only a set of winners, but additionally independently verifiable evidence of the correctness of the count.

Indeed, the entire formulation of our data type, and the split into state data, verification data, and correctness constraints, has been geared towards extraction as a goal. Technically, the verification conditions are *propositions*, i.e. inhabitants of Type *Prop* in the terminology of Coq, and hence erased at extraction time. This corresponds to the fact that the assertions embodied in the correctness constraints can be verified with minimal computational overhead, given the state and the verification data. For example, it can simply be verified whether or not a zero knowledge proof indeed verifies honest decryption by running it through a verifier. On the other hand, the zero knowledge proof

---

[1]https://github.com/mukeshtiwari/EncryptionSchulze/tree/master/code/Workingcode

itself (which is part of the verification data) is crucially needed to be able to verify that a plaintext is the honest decryption of a ciphertext, and hence cannot be erased during extraction. Technically, this is realised by formulating both state and verification data at type level (rather than as propositions).

As we have explained in Section 5.3, the formal development does not pre-suppose any specific implementation of the cryptographic primitives, and we assume the existence of cryptographic infrastructure. From the perspective of extraction, this produces an executable with "holes", i.e. the cryptographic primitives need to be supplied to fill the holes and indeed be able to compile and execute the extracted program.

To fill this hole, we implement the cryptoraphic primitives with help of the UniCrypt library[Locher and Haenni, 2014]. UniCrypt is a freely available library, written in Java, that provides nearly all of the required functionality, with the exception of honest decryption zero knowledge proofs. We extract our proof development into OCaml and use Java/OCaml bindings [Aguillon] to make the UniCrypt functionality available to our OCaml program. Due to differences in the type structure between Java and OCaml, mainly in the context of sub-typing, this was done in the form of an OCaml wrapper around Java data structures. After instantiating the cryptographic primitives in the extracted OCaml code with wrapper code that calls UniCrypt we tested the executable on a three candidate elections between candidates A, B and C. The computation produces a tally sheet that is schematically given below: it is trace of computation which can be used as a checkable record to verify the outcome of election. We elide the cryptographic detail, e.g. the concrete representation of zero knowledge proofs. A certificate is be obtained from the type ECount where the head of the certificate corresponds to the base case of the inductive type, here ecax. Below, M is encrypted margin matrix, D is its decrypted equivalent, required to be identically zero, and Z represents a matrix of zero knwoledge proofs, each establishing that the XY-component of M is in fact an encryption of zero. All these matrices are indexed by candidates and we display these matrices by listing their entries prefixed by a pair of candidates, e.g. the ellipsis in AB(...) denotes the matrix entry at row A and column B.

```
M: AB(rel-marg-of-A-over-B-enc),  AC(rel-marg-of-A-over-C-enc),  ...
D: AB(0)                       ,  AC(0)                       ,  ...
Z: AB(zkp-for-rel-marg-A-B)    ,  AC(zkp-for-rel-marg-A-C)    ,  ...
```

Note that one can verify the fact that the initial encrypted margin is in fact the zero margin by just verifying the zero knowledge proofs. Successive entries in the certificate will generally be obtained by counting valid, and discarding invalid ballots. If a valid ballot is counted after the counting commences, the certificate would continue by exhibiting the state and verification data contained in the `ecvalid` constructor which can be displayed schematically as follows:

```
V: AB(ballot-entry-A-B) , AC(ballot-entry-A-C), ...
C: permutation-commitment
P: zkp-of-valid-permutation-commitment
R: AB(row-perm-A-B)     , AC(row-perm-A-C)    , ...
RP: A(zkp-of-perm-row-A), B(zkp-of-perm-row-B), ...
C: AB(col-perm-A-B),      AC(col-perm-A-C)    , ...
CP: A(zkp-of-perm-col-A), B(zkp-of-perm-col-B), ...
D: AB(dec-perm-bal-A-B) , AC(dec-perm-bal-A-C), ...
Z: AB(zkp-for-dec-A-B)  , AC(zkp-for-dec-A-C) , ...
M: AB(new-marg-A-B)     , AC(new-marg-A-C)    , ...
```

Here `V` is the list of ballots to be counted, where we only diplay the first element. We commit to a permutation and validate this commitment with a zero knowledge proof, here given in the second and third line, prefixed with `C` and `P`. The following two lines are a row permutation of the ballot `V`, together with a zero knowledge proof of correctness of shuffling (of each row) with respect to the permutation committed to by `C` above. The following two lines achieve the same for subsequently permuting the columns of the (row permuted) ballot. Finally, `D` is the decrypted permuted ballot, and `Z` a zero knowledge proof of honest decryption. We end with an updated homomorphic margin matrix `M`. Again, we note that the validity of the decrypted ballot can be checked easily, and validating zero knowledge proofs substantiate that the decrypted ballot is indeed a shuffle of the original one. Homomorphic addition can simply be re-computed.

The steps where invalid ballots are being detected is similar, with the exception of not updating the margin matrix. Once all ballots are counted, the only applicable constructor is `ecdecrypt`, the data content of which would continue a certificate schematically as follows:

```
V: []
M: AB(fin-marg-A-B), AC(fin-marg-A-C), ...
D: AB(dec-marg-A-B), AC(dec-marg-A-C), ...
```

```
Z: AB(zkp-dec-A-B) , AC(zkp-dec-A-C) , ...
```

Here the first line indicates that there are no more ballots to be counted, `M` is the final encrypted margin matrix, `D` is its decryption and `Z` is a matrix of zero knowledge proofs verifying the correctness of decryption.

The certificate would end with the determination of winners based on the encrypted margin, and would end with the content of the `ecfin` constructor

```
winning: A, <evidence that A wins against B and C>
losing:  B, <evidence that B loses against A and C>
losing:  C, <evidence that C loses against A and B>
```

where the notion of evidence for winning and losing is as in the plaintext version of the protocol (4).

We note that the schematic presentation of the certificate above is nothing but a representation of the data contained in the extracted type `ECount` that we have chosen to present schematically. Concrete certificates can be inspected with the accompanying proof development, and are obtained by simply implementing datatype to string conversion on the type `ECount`.

To demonstrate proof of concept, we have run our experiment on an Intel i7 2.6 GHz Linux desktop computer with 16GB of RAM for three candidates and randomly generated ballots (Figure 5.1). The largest amount of ballot we counted was 10,000 (not included in graph), with a runtime of 25 hours. A more detailed analysis reveals that the bottleneck are the bindings between OCaml and Java. More specifically, producing the cryptographic evidence using the UniCrypt Library for 10,000 ballots takes about 10 minutes, and the subsequent computation (which is the same as for the plaintext count) takes negligible time. This is consistent with the mechanism employed by the bindings: each function call from OCaml to Java is inherently memory bounded and creates an instance of the Java runtime, the conversion of OCaml data structures into Java data structures, computation by respective Java function producing result, converting the result back into OCaml data structure, and finally destroying the Java runtime instance when the function returns. While the proof of concept using OCaml/Java bindings falls short of being practically feasible, our timing analysis substantiates that feasibility can be achieved by eliminating the overhead of the bindings.

Figure 5.1: Experimental Result

5.6   Summary

The main contribution of our formalisation is that of independently verifiable *evidence* for a set of candidates to be the winners of an election counted according to the Schulze method. Our main claim is that our notion of evidence is both safeguarding the privacy of the individual ballot (as the count is based on encrypted ballots) and is verifiable at the same time (by means of zero knowledge proofs). To do this, we have axiomatised a set of cryptographic primitives to deal with encryption, decryption, correctness of shuffles and correctness of decryption. From formal and constructive proof of the fact that such evidence can always be obtained, we have then extracted executable code that is provably correct by construction and produces election winners together with evidence once implementations for the cryptographic primitives are supplied.

In a second step, we have supplied an implementation of these primitives, largely based on the UniCrypt Library. Our experiments have demonstrated that this approach is feasible, but quite clearly much work is still needed to improve efficiency.

Assumptions for Provable Correctness. While we claim that the end product embodies a high level of reliability, our approach necessarily leaves some gaps between the executable and the formal proofs. First and foremost, this is of course the implementation of the cryptographic primitives in an external (and unverified) library. We have minimised this gap by basing our implementation on a purpose-specific existing library (UniCrypt) to which we relegate most of the functionality.

Modelling Assumptions. In our modelling of the cryptographic primitives, in particular the zero knowledge proofs, we assumed properties which in reality only hold with very high probability. As a consequence our correctness assertions only hold to the level of probability that is guaranteed by zero knowledge proofs (sigma protocols).

Scalability. We have analysed the feasibility of the extracted code by counting an increasing number of ballots. While this demonstrates a proof of concept, our results show that the bindings used to couple the cryptographic layer with our code adds significant overhead compared to plaintext tallying (4). Given that both parts are practically efficient by themselves, scalability is merely the question of engineering a more efficient coupling.

In a nutshell, the achieved and missed part of this formalization:

- Achieved

  - Correctness: The implementation is formalized in Coq assuming the existence of cryptographic functions and axioms about their correctness behaviour. These primitives were used for constructing evidence, or certificate.

  - Privacy : We don't reveal the content of ballot at any phase of election counting. Therefore, there is no possibility of anyone knowing the choices of a voter other than the voter herself.

  - Verifiability: The outcome of any election can be verified by any third party because of the generated certificates. However, certificates are only accessible to someone having specialized knowledge of cryptography.

- Missed

  - Verifiability: The nature of certificates in this case is very complex, and they can only be scrutinize by someone having specialized knowledge of cryptography which decreases the pool of potential scrutinizers dramatically.

  - Correctness: We used external unverified libraries for cryptographic code. In general, these libraries could be buggy and may produce a wrong result. This, per say, is not a problem because it will be caught during the certificate checking by any independent party, but, it may create a atmosphere of distrust among voters about electronic voting.

Our formalization leaves some gaps which needed to be filled:

- A formally verified cryptographic library to fill the *correctness gap*.

- A formally verified checker to ease the auditing of election to fill the *Scrutineers* gap

Developing a formally verified library to fill the correctness gap would have taken more time, specifically the commitment consistent shuffle, so we chose to formally verify the certificate checker to ease the auditing of election to increase the number of scrutineers gap.

In the next chapter, we will focus on all the details needed to develop formally verified certificate checker for the certificate we produced in this chapter. However, due to time constraint and complex nature of our certificate, we could not formalize every cryptographic primitive needed to verify our certificate. Rather, we have developed a proof of concept formally verified certificate checker for International Association for Cryptologic Research 2018 election, a simpler scrutiny sheet than ours which does not involve any shuffle.

Scrutiny Sheet : Software Independence

> Somewhere inside of all of us is the power to change the world.
>
> *Roald Dahl*

6.1   Introduction

One of the major disadvantage of introducing cryptography to achieve privacy, encryption to make the content of ballot private, and verifiability, zero-knowledge-proof for verification of claims, makes the verification process cumbersome. As a consequence, the verification process (checking the scrutiny sheet) is only accessible to tiny fraction of representative population, cryptographers, results into a sharp decrease in number of scrutineers. While it is not very difficult to find cryptographers to verify the election, they are, off course, not the representative population of any democracy. In order to increase the number of scrutineers and subsequently confidence in electronic voting, we follow the route of providing a formally verified open-source reference certificate checker, which anyone can inspect and run on the election data. The rationale behind formally verifying the certificate checker is *correctness* and open sourcing is to gain the public trust via scrutiny or openness. For example, consider a scenario where we do not provide the reference checker, then how likely would it be for community/voters to develop the verified checker? Moreover, assuming that we publish one unverified certificate checker, what would happen if it returns false on a valid certificate because of its own bug?

Both situations, off course, would be a devastating situation, so not only we should provide a reference certificate checker, but it should be a formally verified one. Additionally, a formally verified reference certificate checker would open the gate for debate in case of someone's implementation for checking certificate diverges from the reference checker. In the case of diverging situation, there are two possibility, either the reference checker is verified using wrong assumptions, or the implementation itself is wrong. The first situation is certainly not very pleasant because it would deteriorate the public trust in the system, but nonetheless, it is always good to have openness in democracy to make it more strong.

At this point of time, the astute reader can criticize us that the source code of formally verified certificate checker should be inspected by someone having the expertise of cryptography and formal verification (logic) to see if the verification has been carried out diligently. We sincerely accept the criticism of the reader that this is indeed the case, but what makes this effort worthwhile is the increased ability of voters to verify the election by themselves by simply running the checker, which has been checked by a community having the expertise of cryptography and logic.

In this chapter, we will discuss all the concepts required to develop a verified certificate checker for the certificate we generated in the last chapter. We have already explained our certificate in the section 5.5, but intuitively, checking our certificate amounts to proving that the homomorphic margin has been computed correctly and zero-knowledge-proof for every claim is correct. Recall that our claims were that we have honestly decrypted every value, computed the finally tally correctly, and a ballot has been permuted by the same permutation whose commitment is published (commitment consistent shuffle [Wikström, 2009]). We will sketch the proof of honest decryption, Pedersen's commitment, and homomorphic computing of tally. As a consequence, translating these sketches into Coq proof is fairly trivial. The triviality comes from the fact that most of these proofs are basic algebraic manipulation over various algebraic structures, group, and field. Furthermore, we have formalized these algebraic manipulation in Coq. We have not formalized the shuffle algorithm [Wikström, 2009] because of its complexity and would take a fair amount of time (the efforts are ongoing). Finally, all of the concepts explained in this chapter would be sufficient to develop the formally verified certificate checker for IACR election.

Chapter Outline: In the section 6.2, we discuss the underlying algebraic structures needed for various cryptographic operations. Section 6.3 focuses on generalizing the Pedersen commitment scheme for a matrix. In the follow-

ing section 6.4, we discuss the details of sigma protocol, and its formalization in Coq. Moreover, we also discuss how to turn a probabilistic reasoning, which is the heart of sigma protocols, to deterministic reasoning without losing the meaning of sigma protocol. In section 6.4.1, we show how to make a concrete instance of sigma protocol by giving an example of discrete logarithm. In addition, we show the protocol needed for honest decryption (section 6.4.2). Section 6.5 sketches the homomorphically tally based on additive ElGamal scheme. Finally, we discuss the IACR scrutiny sheet in section 6.6, and we summarize the chapter in the section 6.7

6.2 Algebraic Structures: Building Blocks

The basic building blocks any cryptographic system are algebraic structures, specifically cyclic group of prime order, field and vector space. In general, we do not need vector space, and group and field are sufficient for most of the cryptographic purposes. However, vector space of a cyclic group of prime order over the field of integers modulo the same order is nicer to work because multiplying a group element by a field element can be abstracted over scalar multiplication of vector space.

Definition 2 (Group) *A group is a set G, with a binary operator* $* : G \to G \to G$, *identity element e, and inverse operator inv* $: G \to G$ *such that the following laws hold:*

- *Associativity*: $\forall$ a b c $\in$ G, $a * (b * c) = (a * b) * c$

- *Closure*: $\forall$ a b $\in$ G, $a * b \in G$

- *Inverse Element*: $\forall$ a $\in$ G $\exists$ $a^{-1} \in G$, such that $a * a^{-1} = a^{-1} * a = e$. a^{-1} is called inverse of a (*inv* a).

- *Identity*: $\forall$ a $\in$ G, $a * e = e * a = a$

If the group is commutative, i.e. $\forall$ a b $\in$ G, $a * b = b * a$, then we call it abelian group. We can represent the abelian group in Coq by using the record data type.

```
Record AbelGroup (G : Type)
    (Hdec : forall x y : G, {x = y} + {x <> y})
```

```
   (dot : G -> G -> G)  (inv : G -> G) (e : G) :=
  {
    dot_associativity : forall x y z,
      dot x (dot y z) = dot (dot x y) z;
    dot_left : forall x, dot x e = x;
    dot_right : forall x, dot e x = x;
    left_inverse : forall x, dot (inv x) x = e;
    right_inverse : forall x, dot x (inv x) = e;
    commutative : forall x y, dot x y = dot y x
  }.
```

Our Coq encoding is slight different from the definition of the group we gave above, mainly that *G : Type* "read it as A has the type Type" and *(Hdec : forall x y : G, x = y + x <> y)*. Because of the theoretical foundations of Coq, every term in Coq has to have a type; hence we need to explicitly state the type of term *G*, and we assume a decidable equality *Hdec* on the elements of *G*.

Definition 3 (Field) *A field is a set* $\mathbb{F}$*, with two binary operator* $+ : \mathbb{F} \rightarrow \mathbb{F} \rightarrow \mathbb{F}$*, and* $\cdot : \mathbb{F} \rightarrow \mathbb{F} \rightarrow \mathbb{F}$*, two identity element* 0 *and* 1*, and two unary operator* $- : \mathbb{F} \rightarrow \mathbb{F}, 1/ : \mathbb{F} \rightarrow \mathbb{F}$ *such that:*

- $(\mathbb{F}, +, 0, -)$ forms an abelian group.

- $(\mathbb{F} - \{0\}, \cdot, 1, 1/)$ forms an abelian group.

- $\cdot$ distributes over +.

Definition 4 (Vector Space) *A set V with two binary operation* $+ : V \rightarrow V \rightarrow V$ *and* $\cdot : \mathbb{F} \rightarrow V \rightarrow V$*, is a vector space over a field* $\mathbb{F}$ *if the following properties hold:*

- Closure under addition: $(V, +)$ forms an abelian group.

- Scalar distribute with respect to vector addition: $\forall r \in \mathbb{F}, u\ v \in V, r \cdot (u + v) = r \cdot u + r \cdot v$

- Scalar distribute with respect to field addition: $\forall a\ b \in \mathbb{F}, u \in V, (a +_{\mathbb{F}} b) \cdot u = a \cdot u + b \cdot v$

- Scalar multiplication is associative with respect to field multiplication: $\forall a\ b \in \mathbb{F}, u \in V, (a \cdot_{\mathbb{F}} b) \cdot u = a \cdot (b \cdot u)$

- Identity: $\forall a \in V, 1_{\mathbb{F}} \cdot a = a$

6.3   Pedersen Commitment Scheme

Recall that in the last chapter to prove that if a ballot was valid or invalid, we generated a secret permutation π and published its commitment using Pedersen commitment scheme. Later, we use this to shuffle each row and column of the ballot by this (secret) permutation. In the shuffle algorithm [Wikström, 2009]), the data structure of permutation π is matrix (a permutation matrix to be precise). In this section, we discuss that how to use generalize Pedersen commitment scheme for matrix.

A Pedersen commitment for any two give group element $g\ h \in G$, a message $m \in \mathbb{F}$, and random element $r \in \mathbb{F}$ is:

Pedersen_commitment g h m r := $g^r \cdot h^m$

We can extend the Pedersen commitment for a group element g, vector of n group element $h_1, h_2 \ldots h_n$, vector of n field element $m_1, m_2 \ldots m_n$, and a random field element r is:

Pedersen_generalized g hs ms r := $g^r \cdot \prod_{i=1}^{n} h_i^{m_i}$

The $\prod_{i=1}^{n} h_i^{m_i}$ function can be written in Coq as (this program is written using Equation library [Sozeau and Mangin, 2019]):

```
Equations prod {F G n}
  (f : G -> F -> G) (g : G -> G -> G)
  (v : Vector.t G (S n)) (w : Vector.t F (S n)) : G :=
 prod  (n := 0) f g (vcons h _) (vcons m  _) := f h m;
 prod  (n := S n) f g (vcons h hs) (vcons m ms)  :=
    g (f h m) (prod f g hs ms).
```

Finally, to commit a matrix of size $N \times N$, we need to call *Pedersen_generalized* on each column of matrix. Consequently, we would get a vector of commitments of size N.

6.4   Sigma Protocol: Efficient Zero-Knowledge-Proof

A sigma protocol is a two party protocol, a prover P and a verified V, where prover P try to convince the V that she holds a private input x for some public input w such that a binary relation R holds, i.e. $(x, w) \in R$. Sigma protocol, in general, is a three step protocol:

1. Initialisation: P generates a random message r, commits it, and send the committed message to V

2. Challenge: V generates a random message c, and sends it to V

3. Response: P sends a response z to V

Upon receiving the response z, V either accepts the proof or rejects the proof. Now we define the sigma protocol in Coq by using record data type.

```
Record SigmaProtocol (Statement : Type) (* Statement x *)
       (Witness : Type) (* witness w *)
       (Rel : Statement -> Witness -> bool) (* decidable relation *)
       (RandCoin : Type) (* random coin *)
       (Commitment : Type) (* commitments *)
       (Challenge : Type) (* challenges *)
       (Response : Type) (* response *) :=
  MkSigma {
  (* initial commitment send by the Prover *)
  initial : RandCoin -> Commitment;
  (* Randomness send by the verifier.  *)
  challenge : Challenge;
  (* response generate by prover *)
  response : Statement -> Witness ->
    RandCoin -> Challenge ->
    Response;
  (* verify the response *)
  verify : Statement * Commitment * Challenge * Response
  -> bool;
  (* Simulator *)
  simulator : Statement -> Challenge -> Response ->
    Statement * Commitment * Challenge * Response;
  (* Extractor *)
```

```
    extractor : Challenge -> Response -> Challenge ->
      Response -> Witness;


  (* Completness *)
  Completness : forall (s : Statement) (w : Witness)
    (r : RandCoin) (e : Challenge),vRel s w = true ->
    verify (s, initial r, e, response s w r e) = true;

  (* Special Soundness *)
  Special_Soundness : forall s c e1 e2 r1 r2,
    e1 <> e2 ->
    verify (s, c, e1, r1) = true ->
    verify (s, c, e2, r2) = true ->
    Rel s (extractor e1 r1 e2 r2) = true;


  (* Probablistic Reasoning could have made it nicer *)
  Special_Honest_Verifier_ZKP  (s : Statement)
    (w : Witness) (e : Challenge):
    Rel s w = true -> forall (r : RandCoin),
    verify (s, initial r, e, response s w r e) = true <->
    forall (z : Response), verify (simulator s e z) = true;


  (* simulator correct *)
  Simulator_correct : forall (s : Statement)
   (e : Challenge) (r : Response),
   verify (simulator s e r) = true;
  }.
```

The record *SigmaProtocol* is indexed by *Statement*, the public input known to *P* and *V*, *Witness*, secret input known to the *P*, *Rel* such that $(x, w) \in Rel$, *RandCoin*, the private random coin toss of *P*, *Commitment*, commitment computed by *P* based on the random coin toss of itself, *Challenge*, the random challenge of *V* to *P*, *Response*, the response of *P* send to *V*. Moreover, the body of record *SigmaProtocol* contains functions *initial*, *challenge*, *response*, and *verify* to reflect the three steps of sigma protocol with two auxiliary functions *simulator* and *extractor*. The reason for having *simulator* and *extractor* to prove the property of special honest verifier zero knowledge proof and special soundness of sigma protocol. Finally, we have four property *Completeness*, *Special_Soundness*, *Special_Honest_Verifier_ZKP* and *Simulator_correct*. *Completeness* states that if *P* and *V* follow the protocol, then verifier would ac-

cept the proof. *Special_Soundness*, expresses that if *P* is able to convince *V* with two accepting transcript for the same commitment, then *V* can extract the witness. Recall that *special honest verifier zero knowledge proof* amounts to a probabilistic polynomial time simulator *S* which would generate a proof transcript for some statement *s* with same probability distribution as if there were a real conversation between a prover *P* and a verifier *V* for the statement *s* and witness *w* such that $(s, w) \in R$. Informally, the real proof transcript depends on statement *s*, witness *w*, and challenge *e*, while the simulated proof transcript depends on statement *s* and challenge *e*. (simulator does have not access to witness *w*, so to generate a accepting proof just by using *s* and *e*, it uses a concept call rewinding.) In our definition of *Special_Honest_Verifier_ZKP*, we eschew the probabilistic reasoning by making randomness explicit, and it states that for any given fixed statement *s*, witness *w*, challenge *e* and assumption that *Rel s w* holds, then for every random coin *r* and a accepting real transcript, simulator can construct an accepting transcript from all random response drawn from response space. The *Correct_simulator* property states that the simulator is correct, i.e. any transcript created by simulator checks out.

Finally, we can use our construction, *SigmaProtocol*, as building block for composing different sigma protocols, which we are not explaining here. For example, we can define parallel composition, AND composition, EQ composition, OR composition, etc.

6.4.1   Concrete Sigma Protocol: Discrete Logarithm

One of the most basic sigma protocol is proof of knowledge of discrete logarithm, i.e. given two elements *g* and *h* of a group *G*, prover convinces the verifier that she knows the discrete logarithm ($\log_g h$) in zero knowledge. In mathematical notation of zero knowledge proof, it is represented as: $ZKPoK\{w \mid h = g^w\}$. We can show that this is a sigma protocol inside Coq by by encoding all the functions and proving all the axioms mentioned in our record type *SigmaProtocol*. For example, we can write the *initial* function as taking a input random coin *r* as input and computing g^r, *challenge* as a function which simply returns a challenge *e*, and so forth:

> initial r := g^r
>
> challenge := *e*
>
> response h w r e := $r + e \cdot w$

verify h a e z := $g^z = a \cdot h^e$

simulator s e z := $(g^z \cdot h^{-e}, e, z)$

extractor c_1 z_1 c_2 z_2 := $(z_1 - z_2) \cdot (c_2 - c_1)^{-1}$

Based on these definitions, we can easily discharge the three proofs, *Completeness*, *Special Soundness*, and *Special Honest Verifier Zero Knowledge* axiom by simple algebraic manipulation.

6.4.2   Honest Decryption: Diffie Hellman Tuple

We have sigma protocol at our arsenal, we focus on honest decryption problem. How can we convince someone that for a given group (G, g, p, h) and private key x ($h := g^x$), the message m is the honest decryption of ElGamal cipher text (c_1, c_2) with revealing our private key x? To solve this problem, we use a well known protocol for proving equality of the discrete logarithm [Cramer et al., 1997]. We first discuss the protocol, and later we will show that how we can adopt the protocol for our purpose.

Diffie Hellman Tuple: a tuple (g, h, u, v) is a *Diffie Hellman* tuple if there exists a w such that $u = g^w$ and $v = h^w$. The protocol to prove it is:

- P chooses a random r and sends $a = g^r$ and $b = h^r$.

- V sends a random e

- P sends $z = r + e \cdot w$

- V check $g^z = a \cdot u^e$ and $h^z = b \cdot v^e$

Now we come back to our original problem, i.e. proving that m is the honest decryption of (c_1, c_2). From these values, we construct a *Diffie Hellman* tuple by multiplying c_2 with g^{-m}, i.e. $(g, h, c_1, c_2 \cdot g^{-m})$. A simple algebraic simplification shows that this tuple can be written as $(g, h, g^r, g^m \cdot h^r \cdot g^{-m})$ for some random r. A further simplification leads to (g, h, g^r, h^r), and this tuple is clearly a *Diffie Hellman* tuple, where $u = g^r$ and $v = h^r$. We could not have been able to construct a *Diffie Hellman* tuple and proved the equality of discrete log if we had claimed otherwise. For example, suppose a cheating prover claims that m is the honest decryption of (c_1, c_2), while it is not. If m is not the honest decryption of (c_1, c_2), then we would end up with a tuple

$(g, h, g^r, g^{m_1} \cdot h^r \cdot g^{-m})$, where m_1 is the honest decryption of (c_1, c_2). Clearly, the tuple $(g, h, g^r, h^r \cdot g^{m_1-m})$ is not *Diffie Hellman* tuple; hence a cheating prove would not succeed.

6.5 Homomorphic Tally

Now that we have sorted out the correct decryption, the next challenge in our tally sheet is computing the final tally homomorphically. Since our encryption is additive ElGamal, and recall that our ballot is matrix of ciphertext:

$$
\begin{pmatrix}
(g^{r_{11}}, g^{m_{11}} * h^{r_{11}}) & (g^{r_{12}}, g^{m_{12}} * h^{r_{12}}) & \cdots & (g^{r_{1n}}, g^{m_{1n}} * h^{r_{1n}}) \\
(g^{r_{21}}, g^{m_{21}} * h^{r_{21}}) & (g^{r_{22}}, g^{m_{22}} * h^{r_{22}}) & \cdots & (g^{r_{2n}}, g^{m_{2n}} * h^{r_{2n}}) \\
\vdots & \vdots & \ddots & \vdots \\
(g^{r_{n1}}, g^{m_{n1}} * h^{r_{n1}}) & (g^{r_{n2}}, g^{m_{n2}} * h^{r_{n2}}) & \cdots & (g^{r_{nn}}, g^{m_{nn}} * h^{r_{nn}})
\end{pmatrix}
$$

To compute the finally tally, all we have to do is stack all the valid ballots (matrices) together and multiply the corresponding ciphertexts together to get the final tally matrix (point wise matrix multiplication). The final computed tally can be decrypted honestly by using the same principals described in the previous section. We can capture all these concepts in Coq based on the algebraic structures, group, field, vector space, and prove all the properties by simple algebraic manipulation. We can represent encryption, decryption and ciphertext multiplication for a given cyclic group (G, g, h, x) such that $h = g^x$:

elGamal_enc (g h : G) (r : F) :=$(g^r, g^m \cdot h^r)$

elGamal_dec (g h : G) $(c_1, c_2) := c_2 \cdot c_1^{-x}$

elGamal_mult $(c_1, c_2)(d_1, d_2) := (c_1 \cdot d_1, c_2 \cdot d_2)$

In fact, by simple algebraic manipulation, we can prove that decryption is left inverse of encryption.

$$
\begin{aligned}
\text{elGamal\_dec g h (elGamal\_enc g h r)} &= \text{elGamal\_dec g h } (g^r, g^m \cdot h^r)(unfolding) \\
&= g^m \cdot h^r \cdot (g^r)^{-x}(unfolding) \\
&= g^m \cdot (g^x)^r \cdot (g^r)^{-x}(substitution) \\
&= g^m \cdot g^{xr} \cdot g^{-rx}(algebraic - simplification) \\
&= g^m
\end{aligned}
$$

The final decrypted tally would be a matrix filled with values like $g^{m_1+m_2+\cdots}$, and we need to do a search to find the values of $m_1 + m_2 + \cdots$ from the final decrypted tally. As a consequence, if the number of candidates and ballots are large then this method is no longer a practical method.

6.6   IACR 2018 Election

We followed some of these techniques explained above write a formal certificate checker for IACR 2018 directors election scrutiny sheet. The 2018 IACR directors election considered seven candidates to fill three positions on the board of directors. The voting style was approval voting where all the eligible voters, IACR members, could vote for as many candidates as they liked. After the counting, the top three members were elected to fill the positions.

The Helios voting system [Helios, 2016] was used for the election, and the system was configured with four authorities, who generated an ElGamal [ElGamal, 1985] public key such that all four authorities were required to decrypt efficiently. Every eligible voter received the credentials by email which they used to cast their ballot from their person computer. During the cast process, each voter created seven ElGamal cipher text, encrypting either zero or one, for the seven participating candidates. Since the vote was exponent, the ElGamal cryptosystem became homomorphic additive. At this point, the voter is then offered the chance to audit her encrypted ballot to check that it does indeed contain the vote she intended. If she chooses to audit, she must discard this ballot and asked to cast a fresh ballot. This mechanism is called called Benaloh challenge, and the purpose of this challenge is to catch a cheating machines. Moreover, this method ensures the cast-as-intended because a cheating machine would have no idea when a voter would cast her ballot, so, in most likely scenario, a voter would end up cast her true intentions. Once she has an unaudited ballot with which she is happy, she casts it. The Helios website maintains an append-only bulletin board on which the voter's encrypted ballot appears. After the voting period is over, all the encrypted ballots corresponding to each candidate are multiplied together; so that there is now a single ciphertext for each candidate, encoding the number of votes for that candidate. The authorities then decrypt these (seven) ciphertexts, announce the results and prove, using a sigma protocol, that the announced result is the correct decryption.

Now we focus on three aspects of verifiability: cast-as-intended, collected-

as-cast, and tallied-as-collected. Well, the cast-as-intended has already been assured by Benaloh challenges, and collect-as-cast is ensured by every voter checking her ballot on the bulletin board. The more complicated step is the counted-as-collected check. In order to verify the tallied-as-cast, a scrutineer has to check only the valid ballots (those which are encryption of either zero or one) has contributed to final tally, the final tally has been calculated correctly, and the final tally has been decrypted honestly.

At this stage, there is a published list of encrypted ballots on the bulletin board and a published result. Moreover, to enable scrutiny, the election authority publishes, non-interactive, sigma protocol transcripts for correct encryption and decryption. Using these transcripts, the scrutineer can verify the election by checking the following three things. First, all the ballots included in final tally are indeed the encryption of zero or one, and any ballot containing any other value has been discarded. Second, the scrutineer reruns the (multiplication) computation and checks that the resulting ciphertexts matches the published one. Finally, she checks that the transcripts are valid for the decryption of these combined ciphertexts with respect to the announced result. These three checks suffice to ensure that the ballots were counted-as-collected.

IACR used Schnorr group to avoid attacks various attacks on solving the discrete logarithm problem. A Schnorr group is a multiplicative Abelian subgroup of prime order q of the field of integers modulo a prime p, where p = k * q + 1 for some integer. In IACR election, the primes used were:

```
Definition P : Z := 163286320849330100023840550338054573296016147711
8595538973916730908621480040646579903858363495375294167564556218249
8120750264980492381375579367675648771293800310370964745767014243638
5184425538239734829952673040443267770476629574802693913227893783846
1942859644646984694306187644767462460965622580087564339212631775817
8959584090166763989756712661796378985576873170761772188432331506951
5578810612570530191330785459289835622213963131696224755098184426610
4701843626480690102396623671836720471075593589901375030610773800236
4137917426595737403871114187750804346564731250609196846638183903982
3878845782661365036974934746820711.
```

```
Definition Q : Z := 61329566248342901292543872769978950870633559608
6693371311393755083704587789171.
```

Since theorem provers are known for proving mathematical statements, but not for being good at running computation inside their environment. Naturally, proving any mathematical statement, e.g. number theoretic proofs,

which are computational intensive would not be a ideal situation for theorem provers. However, the recent advancement in theorem provers (specifically Coq) led us to prove primality of two large prime numbers inside the Coq. To begin with we utilise the CoqPrime library[1] to prove in Coq that the numbers used to define the Schnorr group are in fact prime.

```
Lemma P_prime : prime P.
```

```
Lemma Q_prime : prime Q.
```

Finally, we extracted the Coq code as a OCaml and wrote a main file to glue the extracted code and parsing code. Upon execution, the code returned yes, which asserts that the result produced were correct.

6.7 Summary

In this chapter, we have sketched the ideas for developing a formally verified certificate checker for the certificate we produced in the last chapter. However, due to time constraint and complexity of shuffle primitive, we ended up verifying a simple certification, which did not not involve any zero-knowledge-proof of shuffle. Finally, in this chapter we closed the loop of decrease in number of scrutineers because any one can run the certificate checker. Moreover, we open sourced [2] the checker, so that it can be inspected by anyone (we would like to call it correctness by democratic process). One thing we would like to emphasize that cryptographic concepts are inherently very complex, so running a certificate checker certainly not amounts to understanding the various bits of cryptography and formal method used to develop the certificate checker.

In the next chapter, we will discuss some of the properties of Schulze method which we have formalized in Coq.

---

[1]https://github.com/thery/coqprime
[2]https://github.com/mukeshtiwari/secure-e-voting-with-coq

Machine Checked Schulze Properties

Stay Hungry. Stay Foolish.

Steve Jobs

Since the beginning of democracy, social scientist are constantly looking for methods which would aggregate the individual choice to arrive at acceptable group decisions. In general, these acceptable group decisions were based on intuition of the society at that time, but not backed by mathematical theory. The first mathematical treatment to combine the individual choices (social mathematics) can be attributed to french philosopher and mathematician Marquis de Condorcet (Condorcet method, 1785) and his contemporary and co-national mathematician Jean-Charles de Borda (Borda count, 1770). However, the first formal system, foundational cornerstone of modern social choice theory, for collective preference was given by Kenneth Arrow. In 1950, Kenneth Arrow published a paper titled *A Difficulty in the Concept of Social Welfare* [Arrow, 1950b]. In this paper, Kenneth Arrow envisioned an axiomatic system having the following properties:

- Unrestricted domain

- Non-dictatorship

- Pareto efficiency

- Independence of irrelevant alternatives

Moreover, he concluded that no preferential voting method which can combine or aggregate the individual choices into a community wide ranking would have all the properties of his axiomatic system. This result is now known as *Arrow's impossibility theorem*.

In the light of impossibility theorem, Schulze method, a preferential voting method, can not have all the properties, and it fails on Independence of irrelevant alternatives (IIA) criterion. Despite the fact that Schulze method fails on IIA, it has a plenty of other nice properties established by social choice theorist. In this chapter, we will discuss some of the properties. Moreover, we will show that our implementation adheres to these properties.

7.1   Condorcet Winner

A *Condorcet winner* is a candidate who beats every other candidate in pairwise comparison (also known as head to head competition). Recall that in Schulze method, the pairwise comparison method was margin function, denoted as *marg*, which defined as:

> Given a set of ballots P and candidate set C, we construct graph G based on the margin function $marg : C \times C \rightarrow \mathbb{Z}$. Given two candidates $c, d \in C$, the *margin* of c over d is the number of voters that prefer c over d, minus the number of voters that prefer d over c. In symbols:
>
> $$marg(c, d) = \sharp\{b \in P \mid c >_b d\} - \sharp\{b \in P \mid d >_b c\}$$
>
> where $\sharp$ denotes cardinality and $>_b$ is the strict (preference) ordering given by the ballot $b \in P$.

Now we define the *Condorcet winner* in Coq as:

```
Definition condorcet_winner (c : cand)
  (marg : cand -> cand -> Z) := forall d, marg c d >= 0.
```

Informally, the definition, *condorcet_winner*, states that if a candidate c is *condorcet winner*, then she has been ranked higher against every other candidate. Having the definition of condorcet winner, our goal is to concluded

that if there is a condorcet winner, the Schulze methods always elects it as a winner.

```
(* if candidate c is condorcet winner then it's winner of election *)
 Lemma condorcet_winner_implies_winner (c : cand)
    (marg : cand -> cand -> Z) :  condorcet_winner c marg ->
    c_wins marg c = true.
Proof.
      intros Hc.
      pose proof  condorcet_winner_genmarg.
      pose proof c_wins_true.
      apply H0. intros d.
      pose proof (H c d (length cand_all) marg Hc).
      auto.
Qed.
```

The proof of this theorem hinges on the two key facts:

1. If a canddiate beats everyone in pairwise comparison, then generalized margin between her and every other candidate would be greater than or equal to 0.

2. If a canddiate beats everyone in pairwise comparison, then generalized margin between every other candidate and her would be less than or equal to 0.

It is not very hard to see these two facts based on the definition of generalized margin. Intuitively, if a candidate c is the condorcet winner, then the strongest path between her and every other candidate, say d, would be either a direct path, *marg c d*, or a more stronger path, $M(c, d)$, via some other intermediate candidates.

A directed *path* in the graph, G, from candidate c to candidate d is a sequence $p \equiv c_0, \ldots, c_{n+1}$ of candidates with $c_0 = c$ and $c_{n+1} = d$ ($n \geq 0$), and the *strength*, st, of path, p, is the minimum margin of adjacent nodes, i.e.

$$st(c_0, \ldots, c_{n+1}) = \min\{m(c_i, c_{i+1}) \mid 0 \leq i \leq n\}.$$

For candidates c and d, let $M(c, d)$ denote the maximum strength, or generalized margin of a path from c to d i.e.

$$M(c, d) = \max\{st(p) : \text{p is path from c to d in G}\}$$

We capture these two facts in Coq:

```
Lemma gen_marg_gt0 :
  forall c d n marg,
  condorcet_winner c marg ->
  M marg n c d >= 0.
Proof.
  (* Coq terms omitted *)
Qed.

Lemma gen_marg_lt0 :
  forall c d n marg ,
  condorcet_winner c marg ->
  M marg n d c <= 0.
Proof.
  (* Coq terms omitted *)
Qed.
```

Using these two key facts, we concluded that for any condorcet winner candidate *c*, the generalized margin between her and every other opponent is greater than or equal to reverse generalized margin between every other candidate and her. Formally, in Coq we prove the following theorem on which the proof of *condorcet_winner_implies_winner* hinges.

```
Lemma condorcet_winner_genmarg :
    forall c d n marg,
    condorcet_winner c marg ->
    M marg n d c <= M marg n c d.
```

7.2 Reversal Symmetry

The *Reversal symmetry* is a voting method criterion which states that if the voting method has produced a unique winner, say *c*, based on the cast ballots,

then c should not be elected if the individual choices were reversed. In context of Schulze method, we first need to define the unique winner, and ballot reversal.

```
Definition unique_winner
(marg : cand -> cand -> Z) (c : cand) :=
 c_wins marg c = true /\
 (forall d, d <> c -> c_wins marg d = false).
```

Informally, our definition of *unique_winner* states that the candidate c is a unique winner if it wins the election with respect to computed margin function, *marg*, and every other candidate other than c loses the election.

We capture the ballot reversal in terms of margin function. For any given ballot set, if the computed margin between two candidates c and d is:

$$marg(c,d) = \sharp\{b \in P \mid c >_b d\} - \sharp\{b \in P \mid d >_b c\}$$

If we reverse each ballot from the ballot set, then the new margin function, denoted as *rev_marg*, would be:

$$rev\_marg(c,d) = -1 * marg(c,d)$$

This fact can also be demonstrated using a single ballot ABC. The interpretation is A is strictly preferred over B, and B is preferred over C (but we don not need strict preferences to have this property). The margin function constructed from this ballot is:

$$
\begin{array}{c c}
 & \begin{array}{c c c} A & B & C \end{array} \\
\begin{array}{c} A \\ B \\ C \end{array} &
\left(\begin{array}{c c c}
0 & 1 & 1 \\
-1 & 0 & 1 \\
-1 & -1 & 0
\end{array} \right)
\end{array}
$$

After reversing the original ballot, we get CBA and the margin function is:

$$
\begin{array}{c c}
 & \begin{array}{c c c} A & B & C \end{array} \\
\begin{array}{c} A \\ B \\ C \end{array} &
\left(\begin{array}{c c c}
0 & -1 & -1 \\
1 & 0 & -1 \\
1 & 1 & 0
\end{array} \right)
\end{array}
$$

We capture this notion formally in Coq as:

```
Definition rev_marg
   (marg : cand -> cand -> Z) (c d : cand) :=
   -marg c d.
```

Based on the our definition of *rev_marg*, we can formally state the reversal symmetry as:

```
Lemma winner_reversed :
      forall marg c, unique_winner marg c ->
      c_wins (rev_marg marg) c = false.
```

The lemma, *winner_reversed*, expresses that if a candidate c is a unique winner with respect to *marg* (computed from some ballot set P), then she is not a winner with respect to *rev_marg* (computed from reversing all the entries in the ballot set P).

The proof for this lemma is fairly intuitive, but it takes some efforts to prove it in Coq. In this lemma, we assume the existence of unique winner, say c with respect to *marg*, which means that the generalize margin between her and every other candidate would be greater than the reverse generalized margin, i.e. $\forall d$, M marg $(c, d) >$ M marg (d, c). If we compute the generalize margin with respect to *rev_marg*, then for the candidate c it would be the case that: $\forall d$, M rev_marg $(c, d) <$ M rev_marg (d, c). One key observation is that the graph we get after computing the generalized margin with respect to *rev_marg* is simply a mirror image, every path is reversed, of the graph we get after computing the generalize margin with respect to *marg*. In terms of Coq, it is:

```
Lemma path_with_rev_marg :
  forall k marg c d,
  Path marg k c d <->  Path (rev_marg marg) k d c.
```

Using the lemma *path_with_rev_marg*, we can conclude that if the strength of a path going from c to d with respect to *marg* is greater than or equal to k (by definition of *Path* inductive type), then the strength of a reverse path from d to c would also be greater than or equal to k with respect *rev_marg*.

Using all this fact, and some auxiliary lemma the proof of reversal symmetry is mere rewriting the facts [1].

7.3 Summary

Although, we have just proved two properties of Schulze method (more are still going on), and so far, this chapter is far from being complete. The rationale behind this chapter was to put forward the idea of not only implementing the voting method and proving its correctness, but also proving that the implementation follows the property of voting method. In the next chapter, I will conclude this thesis and some possible direction for future work.

---

[1] At the time of writing this thesis, proof of reversal symmetry hinges on a auxiliary lemma which is fairly intuitive, but demands a lot of Coq machinery (a typical situation in theorem proving). We are in the process of proving it

Conclusion and Future Work

> Education is the most powerful weapon which you can use to change the world.
>
> *Nelson Mandela*

This chapter is summarizes the key outcomes of this dissertation, followed by possible future work.

8.1 Conclusion

Recall that the journey started with the purpose to make the electronic voting process transparent and trustworthy. The premise was:

> Given the potential advantages of electronic voting, we need to address the correctness, privacy, and verifiability concerns for its widespread adoption.

The current state of art software program used by many government is mostly black-box which takes a pile of ballots and produces result. To improve the current situation, we focussed on answering the four main concerns: *(i) correctness*, *(ii) privacy*, *(iii) coercion resistance*, and *(iv) verifiability (tallied-as-cast)* by using *Schulze method* an example.

8.1.1 Correctness

With the intentions to solve the correctness in electronic voting, we approached from mathematical logic route. Rather than implementing the Schulze method, we gave a logical specification of the Schulze winning condition and losing condition with respected to already computed margin function. These specifications were simple enough that any one could inspect them, and make sure that the intent is captured correctly in these specifications. Moreover, we proved the correctness properties about our specification. Also, we put forward the idea of formalizing the properties of voting protocol, in our case Schulze method, in the framework developed by Kenneth Arrow in theorem prover itself. For example, in the last chapter, we proved the *Condorcet* winner property of Schulze method. Given that ballot counting is one of the most crucial phase of any election, the correctness of counting software should be explored from all the possible directions.

8.1.2 Verifiability

We answered the verifiability issue by producing a independently checkable scrutiny sheet. In our case, it contained the step by step computation of margin, winners and loser with the proof why they win or lose the election. This scrutiny sheet can be used any independent third to verify the outcome of election. In case of plain-text ballots, achieving verifiability was trivial, but the encrypted ballot case was very complex. The reason for complexity was the inherent nature of cryptography, so to keep the encrypted ballot election verifiable we augmented the scrutiny sheet with zero-knowledge-proofs. Finally, we developed a formally verified certificate checker to ease the auditing of election (although, we ended up developing checker of different election).

8.1.3 Privacy and Coercion Resistance

Our approach to privacy and coercion problem was homomorphic encryption. To keep the content of ballot private, we did not decrypt any individual ballot and computed the final tally homomorphically by multiplying the cipher-texts. In the final step, we decrypted the fully computed tally, which in turn did not reveal any individual ballots. Since there was no decryption

of any individual ballot, there was no way a voter could have convinced any one about her choices.

8.2 Future Work

8.2.1 Formalizing Cryptographic Entities

During the formalization, we assume the cryptographic primitives for various construction, e.g. encryption primitive, decryption primitive, zero-knowledge-proof primitive, etc. Moreover, we assumed axioms about their correctness property. Formalizing all these primitives and proving the axioms we assumed would further close the trust gap.

8.2.2 Formalizing Properties of Schulze Method

We have formalized just two properties, *Condorcet winner*, and *Reversal symmetry*, but taking it further and proving all the properties would put more trust in the implementation. Moreover, it would be a good stress testing for the specification/implementation and see how far it can go.

8.2.3 Formally Verified Checker

Another interesting avenue would be to explore the formally verified certificate checker. In our formally verified certificate checker, we extracted OCaml code for certificate checking. However, we wrote a substantial amount of unverified OCaml code for parsing the scrutiny sheet. To alleviate this kind of concerns, it would be worth exploring a verified parser, and, probably, evaluating the whole certificate inside Coq environment. There has already been verified parsers written in Coq [Jourdan et al., 2012], and given that we proved the two large primes inside a Coq, it does not seem a far fetched concept.

8.2.4   Risk Limiting Audit for Preferential Voting Scheme

One challenging potential opportunity would be developing a risk limiting audit system for Schulze method. In nutshell, risk limiting audit is a method to audit the election by randomly sampling the ballots. Risk limiting is very well understood in the first-past-the-post voting system, and some of the work has been done in the context IRV elections [Blom et al., 2018], but so far none, to the best of my knowledge, for Schulze method.

8.2.5   Formalizing Code Extraction

This is orthogonal, but very important from the perspective of electronic voting. One of the biggest concern with code extraction of Coq in OCaml/Haskell is that all the security properties proved inside Coq are no longer valid in OCaml/Haskell. Taking the mission critical importance of electronic voting into account, we need a mechanism to translate the proofs all the way up to assembly level. CertiCoq [Anand et al.] seems promising, but it is not very mature yet. The other possibility is using the CakeML [Kumar et al., 2014] to develop the electronic voting schemes.

Bibliography

e-governance. https://e-estonia.com/solutions/e-governance/i-voting/. Accessed on October 17, 2019. (cited on page 2)

Electronic Voting. https://www.e-voting.cc/en/it-elections/world-map. Accessed on October 17, 2019. (cited on page 11)

German Constitution. https://www.bundesverfassungsgericht.de/SharedDocs/Entscheidungen/EN/2009/03/cs20090303_2bvc000307en.html. Accessed on October 18, 2019. (cited on page 13)

New South Wales Election. https://www.abc.net.au/news/2015-02-04/computer-voting-may-feature-in-march-nsw-election/6068290. Accessed on October 18, 2019. (cited on page 15)

Voting Computer Tempest Attack (Information Leaking). https://www.youtube.com/watch?v=B05wPomCjEY. Accessed on October 17, 2019. (cited on page 14)

Western Australia Senate Election. https://www.theguardian.com/world/2014/feb/28/western-australia-senate-election-re-run-to-be-held-on-5-april. Accessed on October 18, 2019. (cited on page 11)

Adrian, D.; Bhargavan, K.; Durumeric, Z.; Gaudry, P.; Green, M.; Halderman, J. A.; Heninger, N.; Springall, D.; Thomé, E.; Valenta, L.; VanderSloot, B.; Wustrow, E.; Zanella-Béguelin, S.; and Zimmermann, P., 2015. Imperfect forward secrecy: How diffie-hellman fails in practice. In *Proceedings of the 22Nd ACM SIGSAC Conference on Computer and Communications Security*, CCS '15 (Denver, Colorado, USA, 2015), 5–17. ACM, New York, NY, USA. doi:10.1145/2810103.2813707. http://doi.acm.org/10.1145/2810103.2813707. (cited on page 36)

Aguillon, J. Ocaml ↔ Java Interface. https://github.com/Julow/ocaml-java. Accessed on April 29, 2019. (cited on page 94)

Alkassar, E.; Böhme, S.; Mehlhorn, K.; and Rizkallah, C., 2014. A framework for the verification of certifying computations. *Journal of Au-*

tomated Reasoning, 52, 3 (Mar 2014), 241–273. doi:10.1007/s10817-013-9289-2. https://doi.org/10.1007/s10817-013-9289-2. (cited on page 21)

ANAND, A.; APPEL, A.; MORRISETT, G.; PARASKEVOPOULOU, Z.; POLLACK, R.; BELANGER, O. S.; SOZEAU, M.; AND WEAVER, M. Certicoq: A verified compiler for coq. (cited on page 126)

ANAND, A. AND RAHLI, V., 2014. Towards a formally verified proof assistant. In *Interactive Theorem Proving*, 27–44. Springer International Publishing, Cham. (cited on page 35)

APPEL, A. W.; MICHAEL, N.; STUMP, A.; AND VIRGA, R., 2003. A trustworthy proof checker. *Journal of Automated Reasoning*, 31, 3 (Nov 2003), 231–260. doi:10.1023/B:JARS.0000021013.61329.58. https://doi.org/10.1023/B:JARS.0000021013.61329.58. (cited on page 35)

ARANHA, D. F.; BARBOSA, P. Y.; CARDOSO, T. N.; ARAÚJO, C. L.; AND MATIAS, P., 2019. The return of software vulnerabilities in the brazilian voting machine. *Computers Security*, 86 (2019), 335 – 349. doi:https://doi.org/10.1016/j.cose.2019.06.009. http://www.sciencedirect.com/science/article/pii/S0167404819301191. (cited on page 2)

ARKOUDAS, K. AND RINARD, M. C., 2005. Deductive runtime certification. *Electr. Notes Theor. Comput. Sci.*, 113 (2005), 45–63. (cited on page 19)

ARROW, K. J., 1950a. A difficulty in the concept of social welfare. *Journal of Political Economy*, 58, 4 (1950), 328–346. (cited on page 4)

ARROW, K. J., 1950b. A difficulty in the concept of social welfare. *Journal of political economy*, 58, 4 (1950), 328–346. (cited on page 115)

AUSTRALIAN ELECTORAL COMMISSION, 2013. Letter to Mr Michael Cordover, LSS4883 Outcome of Internal Review of the Decision to Refuse your FOI Request no. LS4849. avaliable via http://www.aec.gov.au/information-access/foi/2014/files/ls4912-1.pdf, retrieved October 23, 2019. (cited on pages 2 and 15)

BACKES, M.; HRITCU, C.; AND MAFFEI, M., 2008. Automated verification of remote electronic voting protocols in the applied pi-calculus. In *Proceedings of the 2008 21st IEEE Computer Security Foundations Symposium*, CSF '08, 195–209. IEEE Computer Society, Washington, DC, USA. doi:10.1109/CSF.2008.26. https://doi.org/10.1109/CSF.2008.26. (cited on page 7)

BARENDREGT, H. P., 1992. Handbook of logic in computer science (vol. 2). chap. Lambda Calculi with Types, 117–309. Oxford University Press, Inc.,

New York, NY, USA. ISBN 0-19-853761-1. http://dl.acm.org/citation.cfm?id=162552.162561. (cited on page 24)

BARRAS, B., 1996. Coq en coq. (1996). (cited on page 35)

BAYER, S. AND GROTH, J., 2012. Efficient zero-knowledge argument for correctness of a shuffle. In *Proc. EUROCRYPT 2012*, vol. 7237 of *Lecture Notes in Computer Science*, 263–280. Springer. (cited on pages 78 and 82)

BECKERT, B.; GORÉ, R.; SCHÜRMANN, C.; BORMER, T.; AND WANG, J., 2014. Verifying voting schemes. *Journal of Information Security and Applications*, 19, 2 (2014), 115 – 129. doi:https://doi.org/10.1016/j.jisa.2014.04.005. http://www.sciencedirect.com/science/article/pii/S2214212614000246. (cited on page 16)

BEN-OR, M.; GOLDREICH, O.; GOLDWASSER, S.; HÅSTAD, J.; KILIAN, J.; MICALI, S.; AND ROGAWAY, P., 1988. Everything provable is provable in zero-knowledge. In *CRYPTO*, vol. 403 of *Lecture Notes in Computer Science*, 37–56. Springer. (cited on page 78)

BENALOH, J.; MORAN, T.; NAISH, L.; RAMCHEN, K.; AND TEAGUE, V., 2009. Shuffle-sum: coercion-resistant verifiable tallying for STV voting. *IEEE Trans. Information Forensics and Security*, 4, 4 (2009), 685–698. (cited on page 9)

BENALOH, J. AND TUINSTRA, D., 1994. Receipt-free secret-ballot elections (extended abstract). In *Proceedings of the Twenty-sixth Annual ACM Symposium on Theory of Computing*, STOC '94 (Montreal, Quebec, Canada, 1994), 544–553. ACM, New York, NY, USA. doi:10.1145/195058.195407. http://doi.acm.org/10.1145/195058.195407. (cited on pages 2 and 68)

BERARDI, S., 1988. Towards a mathematical analysis of the coquand-huet calculus of constructions and the other systems in barendregt's cube. *Technical report, Carnegie-Mellon University (USA) and Universita di Torino (Italy)*, (1988). (cited on page 24)

BERNHARD, M.; BENALOH, J.; HALDERMAN, J. A.; RIVEST, R. L.; RYAN, P. Y. A.; STARK, P. B.; TEAGUE, V.; VORA, P. L.; AND WALLACH, D. S., 2017. Public evidence from secret ballots. In *Proc. E-Vote-ID 2017*, vol. 10615 of *Lecture Notes in Computer Science*, 84–109. Springer. (cited on pages 2, 19, 68, and 78)

BERTOT, Y.; CASTÉRAN, P.; HUET, G.; AND PAULIN-MOHRING, C., 2004. *Interactive theorem proving and program development : Coq'Art : the calculus of inductive constructions*. Texts in theoretical computer science. Springer. (cited on pages 4, 23, and 47)

BLOM, M.; STUCKEY, P. J.; AND TEAGUE, V. J., 2018. Ballot-polling risk limiting audits for irv elections. In *Electronic Voting*, 17–34. Springer International Publishing, Cham. (cited on page 126)

BRUNI, A.; DREWSEN, E.; AND SCHÜRMANN, C., 2017. Towards a mechanized proof of selene receipt-freeness and vote-privacy. In *Electronic Voting*, 110–126. Springer International Publishing, Cham. (cited on page 8)

CARR'E, B. A., 1971. An algebra for network routing problems. *IMA Journal of Applied Mathematics*, 7, 3 (1971), 273. (cited on page 61)

CHAUM, D., 2004. Secret-ballot receipts: True voter-verifiable elections. *IEEE Security & Privacy*, 2, 1 (2004), 38–47. (cited on page 68)

CHAUM, D. AND PEDERSEN, T. P., 1992. Wallet databases with observers. In *CRYPTO*, vol. 740 of *Lecture Notes in Computer Science*, 89–105. Springer. (cited on page 81)

CHEN, H.; ZIEGLER, D.; CHAJED, T.; CHLIPALA, A.; KAASHOEK, M. F.; AND ZELDOVICH, N., 2015. Using crash hoare logic for certifying the fscq file system. In *Proceedings of the 25th Symposium on Operating Systems Principles*, SOSP '15 (Monterey, California, 2015), 18–37. ACM, New York, NY, USA. doi:10.1145/2815400.2815402. http://doi.acm.org/10.1145/2815400.2815402. (cited on page 16)

CHLIPALA, A., 2013. *Certified Programming with Dependent Types: A Pragmatic Introduction to the Coq Proof Assistant*. The MIT Press. ISBN 0262026651, 9780262026659. (cited on page 47)

COCHRAN, D. AND KINIRY, J., 2010. Votail: A formally specified and verified ballot counting system for Irish PR-STV elections. In *Pre-proceedings of the 1st International Conference on Formal Verification of Object-Oriented Software (FoVeOOS)*. (cited on page 8)

COCHRAN, D. AND KINIRY, J. R., 2013. Formal model-based validation for tally systems. In *Proc. Vote-ID 2013*, vol. 7985, 41–60. Springer. (cited on page 8)

COHEN, E.; DAHLWEID, M.; HILLEBRAND, M.; LEINENBACH, D.; MOSKAL, M.; SANTEN, T.; SCHULTE, W.; AND TOBIES, S., 2009. Vcc: A practical system for verifying concurrent c. In *Theorem Proving in Higher Order Logics*, 23–42. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 21)

CONSTABLE, R. L.; ALLEN, S. F.; BROMLEY, H. M.; CLEAVELAND, W. R.; CREMER, J. F.; HARPER, R. W.; HOWE, D. J.; KNOBLOCK, T. B.; MENDLER, N. P.; PANANGADEN, P.; SASAKI, J. T.; AND SMITH, S. F., 1986. *Implementing Mathematics*

with the Nuprl Proof Development System. Prentice-Hall, Inc., Upper Saddle River, NJ, USA. ISBN 0-13-451832-2. (cited on page 23)

COQUAND, T. AND HUET, G., 1988. The calculus of constructions. *Inf. Comput.*, 76, 2-3 (Feb. 1988), 95–120. doi:10.1016/0890-5401(88)90005-3. http://dx.doi.org/10.1016/0890-5401(88)90005-3. (cited on pages 4 and 24)

COQUAND, T. AND PAULIN, C., 1988. Inductively defined types. In *International Conference on Computer Logic*, 50–66. Springer. (cited on page 4)

CORTIER, V. AND SMYTH, B., 2011. Attacking and fixing helios: An analysis of ballot secrecy. In *2011 IEEE 24th Computer Security Foundations Symposium*, 297–311. doi:10.1109/CSF.2011.27. (cited on page 7)

CORTIER, V. AND WIEDLING, C., 2012. A formal analysis of the norwegian e-voting protocol. In *Principles of Security and Trust*, 109–128. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 8)

CRAMER, R.; DAMGÅRD, I.; AND SCHOENMAKERS, B., 1994. Proofs of partial knowledge and simplified design of witness hiding protocols. In *Advances in Cryptology — CRYPTO '94*, 174–187. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 44)

CRAMER, R.; GENNARO, R.; AND SCHOENMAKERS, B., 1997. A secure and optimally efficient multi-authority election scheme. In *Advances in Cryptology — EUROCRYPT '97*, 103–118. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on pages 41 and 109)

DE BRUIJN, N. G., 1983. *AUTOMATH, a Language for Mathematics*, 159–200. Springer Berlin Heidelberg, Berlin, Heidelberg. ISBN 978-3-642-81955-1. doi:10.1007/978-3-642-81955-1_11. https://doi.org/10.1007/978-3-642-81955-1_11. (cited on page 23)

DE MOURA, L.; KONG, S.; AVIGAD, J.; VAN DOORN, F.; AND VON RAUMER, J., 2015. The lean theorem prover (system description). In *Automated Deduction - CADE-25*, 378–388. Springer International Publishing, Cham. (cited on page 23)

DELAHAYE, D., 2000. A tactic language for the system coq. In *Proceedings of the 7th International Conference on Logic for Programming and Automated Reasoning*, LPAR'00 (Reunion Island, France, 2000), 85–95. Springer-Verlag, Berlin, Heidelberg. (cited on page 5)

DELAUNE, S.; KREMER, S.; AND RYAN, M., 2010a. Verifying privacy-type properties of electronic voting protocols: A taster. In *Towards Trustworthy Elections, New Directions in Electronic Voting*, vol. 6000 of *Lecture Notes in Computer Science*, 289–309. Springer. (cited on pages 2 and 68)

DELAUNE, S.; KREMER, S.; AND RYAN, M., 2010b. *Verifying Privacy-Type Properties of Electronic Voting Protocols: A Taster*, 289–309. Springer Berlin Heidelberg, Berlin, Heidelberg. ISBN 978-3-642-12980-3. doi:10.1007/978-3-642-12980-3_18. https://doi.org/10.1007/978-3-642-12980-3_18. (cited on page 7)

DEYOUNG, H. AND SCHÜRMANN, C., 2012. Linear logical voting protocols. In *Proc. VoteID 2011*, vol. 7187 of *Lecture Notes in Computer Science*, 53–70. Springer. (cited on page 8)

DIFFIE, W. AND HELLMAN, M., 2006. New directions in cryptography. *IEEE Trans. Inf. Theor.*, 22, 6 (Sep. 2006), 644–654. doi:10.1109/TIT.1976.1055638. http://dx.doi.org/10.1109/TIT.1976.1055638. (cited on page 36)

DIJKSTRA, E. W., 1972. The humble programmer. *Commun. ACM*, 15, 10 (Oct. 1972), 859–866. doi:10.1145/355604.361591. http://doi.acm.org/10.1145/355604.361591. (cited on page 16)

ELGAMAL, T., 1985. A public key cryptosystem and a signature scheme based on discrete logarithms. *IEEE transactions on information theory*, 31, 4 (1985), 469–472. (cited on pages 36, 38, and 111)

ERBSEN, A.; PHILIPOOM, J.; GROSS, J.; SLOAN, R.; AND CHLIPALA, A., 2019. Simple high-level code for cryptographic arithmetic - with proofs, without compromises. In *2019 IEEE Symposium on Security and Privacy, SP 2019, San Francisco, CA, USA, May 19-23, 2019*, 1202–1219. doi:10.1109/SP.2019.00005. https://doi.org/10.1109/SP.2019.00005. (cited on page 16)

FELDMAN, A. J.; HALDERMAN, J. A.; AND FELTEN, E. W., 2007. Security analysis of the diebold accuvote-ts voting machine. In *Proceedings of the USENIX Workshop on Accurate Electronic Voting Technology*, EVT'07 (Boston, MA, 2007), 2–2. USENIX Association, Berkeley, CA, USA. http://dl.acm.org/citation.cfm?id=1323111.1323113. (cited on page 2)

FIRSOV, D. AND UUSTALU, T., 2015. Dependently typed programming with finite sets. In *Proceedings of the 11th ACM SIGPLAN Workshop on Generic Programming, WGP@ICFP 2015, Vancouver, BC, Canada, August 30, 2015*, 33–44. doi:10.1145/2808098.2808102. https://doi.org/10.1145/2808098.2808102. (cited on page 56)

FUJIOKA, A.; OKAMOTO, T.; AND OHTA, K., 1993. A practical secret voting scheme for large scale elections. In *Advances in Cryptology — AUSCRYPT '92*, 244–251. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 7)

GAMAL, T. E., 1984. A public key cryptosystem and a signature scheme based on discrete logarithms. In *CRYPTO*, vol. 196 of *Lecture Notes in Computer Science*, 10–18. Springer. (cited on page 85)

GENTRY, C., 2009. *A Fully Homomorphic Encryption Scheme*. Ph.D. thesis, Stanford, CA, USA. AAI3382729. (cited on page 40)

GEUVERS, H., 2001. Induction is not derivable in second order dependent type theory. In *Typed Lambda Calculi and Applications*, 166–181. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 28)

GEUVERS, H.; WIEDIJK, F.; AND ZWANENBURG, J., 2002. A constructive proof of the fundamental theorem of algebra without using the rationals. In *Types for Proofs and Programs*, 96–111. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 16)

GHALE, M. K.; GORÉ, R.; AND PATTINSON, D., 2017. A formally verified single transferable vote scheme with fractional values. In *Proc. E-Vote-ID 2017*, LNCS. Springer. This volume. (cited on page 8)

GIMÉNEZ, E., 1995. Codifying guarded definitions with recursive schemes. In *Types for Proofs and Programs*, 39–59. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 29)

GIRARD, J.-Y., 1987. Linear logic. *Theoretical Computer Science*, 50, 1 (1987), 1 – 101. doi:https://doi.org/10.1016/0304-3975(87)90045-4. http://www.sciencedirect.com/science/article/pii/0304397587900454. (cited on page 8)

GOLDREICH, O.; MICALI, S.; AND WIGDERSON, A., 1991. Proofs that yield nothing but their validity for all languages in NP have zero-knowledge proof systems. *J. ACM*, 38, 3 (1991), 691–729. (cited on page 78)

GOLDWASSER, S.; MICALI, S.; AND RACKOFF, C., 1985. The knowledge complexity of interactive proof-systems (extended abstract). In *STOC*, 291–304. ACM. (cited on pages 42 and 78)

GONTHIER, G., 2008. The four colour theorem: Engineering of a formal proof. In *Computer Mathematics*, 333–333. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 16)

Gu, L.; Vaynberg, A.; Ford, B.; Shao, Z.; and Costanzo, D., 2011. Certikos: a certified kernel for secure cloud computing. In *APSys '11 Asia Pacific Workshop on Systems, Shanghai, China, July 11-12, 2011*, 3. doi: 10.1145/2103799.2103803. https://doi.org/10.1145/2103799.2103803. (cited on page 16)

Halderman, J. A. and Teague, V., 2015. The new south wales ivote system: Security failures and verification flaws in a live online election. In *E-Voting and Identity*, 35–53. Springer International Publishing, Cham. (cited on pages 2 and 15)

Hales, T. C.; Adams, M.; Bauer, G.; Dang, D. T.; Harrison, J.; Hoang, T. L.; Kaliszyk, C.; Magron, V.; McLaughlin, S.; Nguyen, T. T.; Nguyen, T. Q.; Nipkow, T.; Obua, S.; Pleso, J.; Rute, J.; Solovyev, A.; Ta, A. H. T.; Tran, T. N.; Trieu, D. T.; Urban, J.; Vu, K. K.; and Zumkeller, R. A formal proof of the kepler conjecture. *arXiv*. (cited on page 16)

Harrison, J., 1996. Hol light: A tutorial introduction. In *Proceedings of the First International Conference on Formal Methods in Computer-Aided Design*, FMCAD '96, 265–269. Springer-Verlag, London, UK, UK. http://dl.acm.org/citation.cfm?id=646184.682934. (cited on page 23)

Harrison, J., 2006. Towards self-verification of HOL Light. In *Automated Reasoning*, 177–191. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 35)

Helios, 2016. The helios voting system. Http://heliosvoting.org/, accessed June 25, 2016. (cited on pages 7 and 111)

Hirt, M. and Sako, K., 2000. Efficient receipt-free voting based on homomorphic encryption. In *Proc. EUROCRYPT 2000*, vol. 1807 of *Lecture Notes in Computer Science*, 539–556. Springer. (cited on page 78)

Jackson, D., 2002. Alloy: A lightweight object modelling notation. *ACM Trans. Softw. Eng. Methodol.*, 11, 2 (Apr. 2002), 256–290. doi:10.1145/505145.505149. https://doi.org/10.1145/505145.505149. (cited on page 8)

Jacobs, B. and Pieters, W., 2009. *Electronic Voting in the Netherlands: From Early Adoption to Early Abolishment*, 121–144. Springer Berlin Heidelberg, Berlin, Heidelberg. ISBN 978-3-642-03829-7. doi:10.1007/978-3-642-03829-7_4. https://doi.org/10.1007/978-3-642-03829-7_4. (cited on page 14)

JOURDAN, J.-H.; POTTIER, F.; AND LEROY, X., 2012. Validating lr(1) parsers. In *Programming Languages and Systems*, 397–416. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 125)

KAUFMANN, M. AND STROTHER MOORE, J., 1996. Acl2: an industrial strength version of nqthm. In *Proceedings of 11th Annual Conference on Computer Assurance. COMPASS '96*, 23–34. doi:10.1109/CMPASS.1996.507872. (cited on page 23)

KLEIN, G.; ELPHINSTONE, K.; HEISER, G.; ANDRONICK, J.; COCK, D.; DERRIN, P.; ELKADUWE, D.; ENGELHARDT, K.; KOLANSKI, R.; NORRISH, M.; SEWELL, T.; TUCH, H.; AND WINWOOD, S., 2009. sel4: formal verification of an OS kernel. In *Proceedings of the 22nd ACM Symposium on Operating Systems Principles 2009, SOSP 2009, Big Sky, Montana, USA, October 11-14, 2009*, 207–220. doi:10.1145/1629575.1629596. https://doi.org/10.1145/1629575.1629596. (cited on page 16)

KOHNO, T.; STUBBLEFIELD, A.; RUBIN, A. D.; AND WALLACH, D. S., 2004. Analysis of an electronic voting system. In *IEEE Symposium on Security and Privacy, 2004. Proceedings. 2004*, 27–40. doi:10.1109/SECPRI.2004.1301313. (cited on page 2)

KREMER, S. AND RYAN, M., 2005. Analysis of an electronic voting protocol in the applied pi calculus. In *Programming Languages and Systems*, 186–200. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 7)

KUMAR, R.; MYREEN, M. O.; NORRISH, M.; AND OWENS, S., 2014. Cakeml: a verified implementation of ML. In *Proc. POPL 2014*, 179–192. ACM. (cited on pages 16, 76, and 126)

KÜSTERS, R.; TRUDERUNG, T.; AND VOGT, A., 2011. Verifiability, privacy, and coercion-resistance: New insights from a case study. In *2011 IEEE Symposium on Security and Privacy*, 538–553. doi:10.1109/SP.2011.21. (cited on pages 2 and 68)

LANDIN, P. J., 1964. The mechanical evaluation of expressions. *The Computer Journal*, 6, 4 (1964), 308. (cited on page 71)

LEROY, X., 2006. Formal certification of a compiler back-end or: programming a compiler with a proof assistant. In *Proceedings of the 33rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2006, Charleston, South Carolina, USA, January 11-13, 2006*, 42–54. doi:10.1145/1111037.1111042. https://doi.org/10.1145/1111037.1111042. (cited on page 16)

LETOUZEY, P., 2003. A new extraction for coq. In *Proc. TYPES 2002*, vol. 2646 of *Lecture Notes in Computer Science*, 200–219. Springer. (cited on page 93)

LETOUZEY, P., 2008. Extraction in Coq: An overview. In *Proc. CiE 2008*, vol. 5028 of *Lecture Notes in Computer Science*, 359–369. Springer. (cited on page 29)

LEWIS, S. J.; PEREIRA, O.; AND TEAGUE, V. Ceci n'est pas une preuve. https://people.eng.unimelb.edu.au/vjteague/UniversalVerifiabilitySwissPost.pdf. Accessed on October 17, 2019. (cited on page 2)

LOCHER, P. AND HAENNI, R., 2014. A lightweight implementation of a shuffle proof for electronic voting systems. In *44. Jahrestagung der Gesellschaft für Informatik, Informatik 2014, Big Data - Komplexität meistern, 22.-26. September 2014 in Stuttgart, Deutschland*, 1391–1400. https://dl.gi.de/20.500.12116/2747. (cited on pages 6 and 94)

MCCONNELL, R.; MEHLHORN, K.; NÄHER, S.; AND SCHWEITZER, P., 2011. Certifying algorithms. *Computer Science Review*, 5, 2 (2011), 119 – 161. doi:https://doi.org/10.1016/j.cosrev.2010.09.009. http://www.sciencedirect.com/science/article/pii/S1574013710000560. (cited on page 19)

MEHLHORN, K. AND NÄHER, S., 1995. Leda: A platform for combinatorial and geometric computing. *Commun. ACM*, 38, 1 (Jan. 1995), 96–102. doi:10.1145/204865.204889. http://doi.acm.org/10.1145/204865.204889. (cited on page 21)

MENEZES, A. J.; VANSTONE, S. A.; AND OORSCHOT, P. C. V., 1996. *Handbook of Applied Cryptography*. CRC Press, Inc., Boca Raton, FL, USA, 1st edn. ISBN 0849385237. (cited on page 47)

MILLER, B. P.; FREDRIKSEN, L.; AND SO, B., 1990. An empirical study of the reliability of unix utilities. *Commun. ACM*, 33, 12 (Dec. 1990), 32–44. doi:10.1145/96267.96279. http://doi.acm.org/10.1145/96267.96279. (cited on page 17)

MILNER, R., 1972. Implementation and applications of scott's logic for computable functions. *SIGPLAN Not.*, 7, 1 (Jan. 1972), 1–6. doi:10.1145/942578.807067. http://doi.acm.org/10.1145/942578.807067. (cited on page 23)

MOORE, J. S., 2019. Milestones from the pure lisp theorem prover to acl2. *Formal Aspects of Computing*, (Jul 2019). doi:10.1007/s00165-019-00490-3. https://doi.org/10.1007/s00165-019-00490-3. (cited on page 17)

MYREEN, M. O. AND DAVIS, J., 2014. The reflective milawa theorem prover is sound - (down to the machine code that runs it). In *Interactive Theorem Proving - 5th International Conference, ITP 2014, Held as Part of the Vienna Summer of Logic, VSL 2014, Vienna, Austria, July 14-17, 2014. Proceedings*, 421–436. doi:10.1007/978-3-319-08970-6_27. https://doi.org/10.1007/978-3-319-08970-6_27. (cited on page 16)

NIPKOW, T.; PAULSON, L. C.; AND WENZEL, M., 2002a. *Isabelle/HOL — A Proof Assistant for Higher-Order Logic*, vol. 2283 of *Lect. Notes in Comp. Sci.* Springer. (cited on page 23)

NIPKOW, T.; WENZEL, M.; AND PAULSON, L. C., 2002b. *Isabelle/HOL: A Proof Assistant for Higher-order Logic*. Springer-Verlag, Berlin, Heidelberg. ISBN 3-540-43376-7. (cited on page 21)

NORELL, U., 2009. Dependently typed programming in agda. In *Proceedings of the 6th International Conference on Advanced Functional Programming*, AFP'08 (Heijen, The Netherlands, 2009), 230–266. Springer-Verlag, Berlin, Heidelberg. http://dl.acm.org/citation.cfm?id=1813347.1813352. (cited on page 23)

OTTEN, J., 2003. Fuller disclosure than intended. (2003). http://www.votingmatters.org.uk/ISSUE17/I17P2.PDF. Accessed on October 17, 2019. (cited on pages 9, 75, and 77)

OWRE, S.; RUSHBY, J. M.; AND SHANKAR, N., 1992. Pvs: A prototype verification system. In *Proceedings of the 11th International Conference on Automated Deduction: Automated Deduction*, CADE-11, 748–752. Springer-Verlag, London, UK, UK. http://dl.acm.org/citation.cfm?id=648230.752639. (cited on page 23)

PAAR, C. AND PELZL, J., 2009. *Understanding Cryptography: A Textbook for Students and Practitioners*. Springer Publishing Company, Incorporated, 1st edn. ISBN 3642041000, 9783642041006. (cited on page 47)

PATTINSON, D. AND SCHÜRMANN, C., 2015. Vote counting as mathematical proof. In *Proc. AI 2015*, vol. 9457 of *Lecture Notes in Computer Science*, 464–475. Springer. (cited on pages 8 and 19)

PATTINSON, D. AND VERITY, F., 2016. Modular synthesis of provably correct vote counting programs. In *Proc. E-Vote-ID 2016*. To appear. (cited on page 8)

PAULIN-MOHRING, C., 1993. Inductive definitions in the system coq - rules and properties. In *Proceedings of the International Conference on Typed Lambda*

Calculi and Applications, TLCA '93, 328–345. Springer-Verlag, London, UK, UK. http://dl.acm.org/citation.cfm?id=645891.671440. (cited on pages 24 and 28)

PEDERSEN, T. P., 1992. Non-interactive and information-theoretic secure verifiable secret sharing. In *Advances in Cryptology — CRYPTO '91*, 129–140. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 82)

PFENNING, F. AND PAULIN-MOHRING, C., 1989. Inductively defined types in the calculus of constructions. In *International Conference on Mathematical Foundations of Programming Semantics*, 209–228. Springer. (cited on page 28)

PIERCE, B. C., 2004. *Advanced Topics in Types and Programming Languages*. The MIT Press. ISBN 0262162288. (cited on page 31)

POLLACK, R., 1998. How to believe a machine-checked proof. *Twenty Five Years of Constructive Type Theory*, 36 (1998), 205. (cited on page 35)

RIVEST, R. L., 2008. On the notion of software independence in voting systems. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 366, 1881 (2008), 3759–3767. https://royalsocietypublishing.org/doi/10.1098/rsta.2008.0149. (cited on page 19)

RIVEST, R. L.; ADLEMAN, L.; DERTOUZOS, M. L.; ET AL., 1978. On data banks and privacy homomorphisms. (1978). (cited on page 39)

RIVEST, R. L. AND SHEN, E., 2010. An optimal single-winner preferential voting system based on game theory. In *Proc. COMSOC 2010*. Duesseldorf University Press. (cited on page 51)

RYAN, P. Y. A.; RØNNE, P. B.; AND IOVINO, V., 2016. Selene: Voting with transparent verifiability and coercion-mitigation. In *Financial Cryptography and Data Security*, 176–192. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 8)

SCHNEIER, B., 1995. *Applied Cryptography (2Nd Ed.): Protocols, Algorithms, and Source Code in C*. John Wiley & Sons, Inc., New York, NY, USA. ISBN 0-471-11709-9. (cited on page 47)

SCHULZE, M., 2011. A new monotonic, clone-independent, reversal symmetric, and Condorcet-consistent single-winner election method. *Social Choice and Welfare*, 36, 2 (2011), 267–303. (cited on pages 4, 50, and 64)

SCHÜRMANN, C., 2009. Electronic elections: Trust through engineering. In *Proc. RE-VOTE 2009*, 38–46. IEEE Computer Society. (cited on page 19)

SLIND, K. AND NORRISH, M., 2008. A brief overview of hol4. In *Proceedings of the 21st International Conference on Theorem Proving in Higher Order Logics*, TPHOLs '08 (Montreal, P.Q., Canada, 2008), 28–32. Springer-Verlag, Berlin, Heidelberg. doi:10.1007/978-3-540-71067-7_6. http://dx.doi.org/10.1007/978-3-540-71067-7_6. (cited on page 23)

SOZEAU, M. AND MANGIN, C., 2019. Equations reloaded: High-level dependently-typed functional programming and proving in coq. *Proc. ACM Program. Lang.*, 3, ICFP (Jul. 2019), 86:1–86:29. doi:10.1145/3341690. http://doi.acm.org/10.1145/3341690. (cited on page 105)

STOLTENBERG-HANSEN, V.; LINDSTRÖM, I.; AND GRIFFOR, E., 1994. *Mathematical Theory of Domains*. No. 22 in Cambridge Tracts in Theoretical Computer Science. Cambridge University Press. (cited on page 59)

SULLIVAN, G. F. AND MASSON, G. M., 1990. Using certification trails to achieve software fault tolerance. In *[1990] Digest of Papers. Fault-Tolerant Computing: 20th International Symposium*, 423–431. doi:10.1109/FTCS.1990.89397. (cited on page 19)

TARSKI, A., 1955. A lattice-theoretical fixpoint theorem and its applications. *Pacific Journal of Mathematics*, 5, 2 (1955), 285—309. (cited on page 59)

TERELIUS, B. AND WIKSTRÖM, D., 2010. Proofs of restricted shuffles. In *AFRICACRYPT*, vol. 6055 of *Lecture Notes in Computer Science*, 100–113. Springer. (cited on page 81)

VERITY, F. AND PATTINSON, D., 2017. Formally verified invariants of vote counting schemes. In *Proceedings of the Australasian Computer Science Week Multiconference*, ACSW '17 (Geelong, Australia, 2017), 31:1–31:10. ACM, New York, NY, USA. doi:10.1145/3014812.3014845. http://doi.acm.org/10.1145/3014812.3014845. (cited on page 8)

WIKSTRÖM, D., 2009. A commitment-consistent proof of a shuffle. In *Proceedings of the 14th Australasian Conference on Information Security and Privacy*, ACISP '09 (Brisbane, Australia, 2009), 407–421. Springer-Verlag, Berlin, Heidelberg. doi:10.1007/978-3-642-02620-1_28. http://dx.doi.org/10.1007/978-3-642-02620-1_28. (cited on pages 83, 102, and 105)

WILCOX, J. R.; WOOS, D.; PANCHEKHA, P.; TATLOCK, Z.; WANG, X.; ERNST, M. D.; AND ANDERSON, T. E., 2015. Verdi: a framework for implementing and formally verifying distributed systems. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation, Portland, OR, USA, June 15-17, 2015*, 357–368. doi:10.1145/2737924.2737958. https://doi.org/10.1145/2737924.2737958. (cited on page 16)

Wolchok, S.; Wustrow, E.; Halderman, J. A.; Prasad, H. K.; Kankipati, A.; Sakhamuri, S. K.; Yagati, V.; and Gonggrijp, R., 2010. Security analysis of india's electronic voting machines. In *Proceedings of the 17th ACM Conference on Computer and Communications Security*, CCS '10 (Chicago, Illinois, USA, 2010), 1–14. ACM, New York, NY, USA. doi:10.1145/1866307.1866309. http://doi.acm.org/10.1145/1866307.1866309. (cited on pages 2 and 14)

Yang, X.; Chen, Y.; Eide, E.; and Regehr, J., 2011. Finding and understanding bugs in c compilers. *SIGPLAN Not.*, 46, 6 (Jun. 2011), 283–294. doi:10.1145/1993316.1993532. http://doi.acm.org/10.1145/1993316.1993532. (cited on page 17)

Zhao, J. and Zdancewic, S., 2012. Mechanized verification of computing dominators for formalizing compilers. In *Certified Programs and Proofs*, 27–42. Springer Berlin Heidelberg, Berlin, Heidelberg. (cited on page 16)